

# TFPT — The Prime Front

Research documentation of the prime / zeta line: seven verified modules, one localized residual, and the complete kill list

Stefan Hamann

July 28, 2026 · v5.4 (rev 438)

## The TFPT document set — what is treated where

**Plain language:** TFPT (Topological Fixed-Point Theory) is a small discrete compiler. Two inputs — the seam constant  $c_3 = \frac{1}{8\pi}$  and the carrier rank  $g_{\text{car}} = 5$  — build  $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$  and read off the Standard-Model skeleton, the constants and the scale grammar. The series is **the introduction plus five numbered papers and five companions**, best read in order:

0. **introduction** — reading guide, compiler closure, predictions, the dependency DAG and the single proof ledger (`verification/status_ledger.csv`). (*entry point*)
1. **tfpt\_1\_architecture\_e8** — the two axioms, the derivation map, the EM fixed point  $\alpha^{-1}$ , the  $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$  construction, the QBL theorem chain.
2. **tfpt\_2\_standard\_model** — the SM in one  $\varphi_0$ -ladder formula, flavor from parabolic transport (sheet diamond, dual anchor,  $Q_+$  cohomology), the worked closures.
3. **tfpt\_3\_e8\_audit\_bootstrap** — the seven  $E_8$  slices as an audit raster, the cascade bridge, the Möbius bootstrap.
4. **tfpt\_4\_frontier** — honest status of  $\eta_B$ ,  $m_p/m_e$ , Koide, dark matter, full QG.
5. **tfpt\_5\_redteam** — the adversarial audit: declared attacks, what survives, what each reduces to, and the kill tests.

**Companions:** **R. tfpt\_research\_contracts** — the open interfaces ( $v_{\text{geo}}$ ,  $G_{\text{net}}$ ,  $F_{\text{transfer}}$ ) as numbered contracts; **H. tfpt\_horizon\_readouts** — Appendix H (unit-system reframe): one seam constant as the universal horizon thermal code, the Nariai/anchor surface, the resummed seam clock; **O. origin\_theory** — why no free fundamental number remains (exact core + one honestly-typed cyclic interpretation); **S. tfpt\_safeguards** — the verification discipline: every mechanism that defends a claim against coincidence and numerology (status calculus, no-free-pattern, the over-determination map, the firewall, frozen predictions + null model, the independent Wolfram and Lean paths, the red team); **P. tfpt\_prime\_front** — research documentation of the prime / zeta line: the seven load-bearing modules on the  $E_8$  glue census (Hecke from geometry through the transport ledger), the sandbox localization program that follows them, and the kill list. Documentation, not a result: it makes no claim about the Riemann Hypothesis.

*You are reading the **Prime Front** companion (P): the complete internal record of the prime / zeta line — how a discrete compiler’s bookkeeping kept producing classical number-theoretic objects, which of those links became load-bearing modules, and how the remaining question was compressed, over a hundred exploratory probes, down to one strict-positivity margin input — then, once that margin was measured to be an artifact of the requirement, down to a two-object core — then down to a three-point list of which only one point is an inequality — and finally, once the boundary reformulation ran, down to that single inequality in a textbook shape (a Galerkin error with Aubin–Nitsche duality), with the prime comb named as the one structure that refuses compression — and, once the missing lemmas of that inequality were measured against their classical addresses, down to exactly one genuinely new analytic statement — and finally, with the whole chain assembled end to end on one ladder, down to a load-bearing spine that is 96.2% identity or completed Cholesky certificate and whose only hypothesis is a declared accounting convention. It is a **research documentation**, not a results paper. Its purpose is that a reader (or a reviewer) can see the whole route, including everything that died on the way.*

## Abstract

A two-axiom discrete compiler ( $c_3 = 1/(8\pi)$ ,  $g_{\text{car}} = 5$ ,  $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ ) has no business meeting the prime numbers. It does anyway: the  $\mu_4$  glue completion carries a quartic character, the signed glue census is a theta series, and the shell counts of that census turn out to be *Hecke* data. It goes further than an analogy: the  $E_8$  shell count is literally  $240 \zeta(s) \zeta(s-3)$  as a Dirichlet series, and the lattice's self-duality gives its theta sum an exact functional equation whose fixed axis is the critical line — a framework picture, stated as such, which the document opens with and then keeps at arm's length. This companion documents the line end to end. Seven modules are load-bearing and ledger-carried (v535–v541): Hecke from geometry on the  $E_8$  census, the Eichler trace layer, the half-integral bridge, one finite relative-trace identity of which the first three are projections, the Weil structure of the resulting family with two explicitly isolated obstructions, the amplitude route with a positive linear carrier, and the matching-lemma / transport ledger package. Everything beyond those seven is *sandbox*: an exploratory program, run as 127 preregistered probes with 3198 machine checks, which localized the residual object. That localization is the actual content of this document. The transport ledger closes exactly, leaving one coupled inequality (called **I5**) which is — *by identity, as a typing and not as progress* — equivalent to Weil positivity and hence to the Riemann Hypothesis. The subsequent arc maps **I5** geographically (Connes–Consani dictionary, tightness curves, the Sonin crossing, the one-atom band), turns the positivity question into a per-zone handover induction, and then certifies its ingredients one by one until the entire remaining hardness sits in a single Loewner statement on the  $J$ -odd half of a window, reducible by Sherman–Morrison/Woodbury to one scalar ratio  $r = \kappa/\varepsilon \leq 1$  (measured 0.0053–0.1814 on 16 zones), then, once  $\varepsilon$  is identified as an exact energy, to two scalars — and finally, with both scalars certified (the wing scalar unconditionally by a graded matrix cap, the boundary value by a residual certificate that carries a cancellation), to a chain that closes 16/16 conditional on exactly one strict-positivity margin input, itself two to six orders of magnitude weaker than the conclusion it buys. The penultimate part then closes the induction circle on the measured zones end to end — a certified base case, fifteen certified handover steps through a graded Loewner minorant, and the structural finding that the atom entry costs nothing — while leaving three sharp, named gaps: no reserve, no scalar step law, and no uniformity in  $k$ , with an extrapolated crossing warning at  $n \approx 170$ . The next part stops extrapolating and measures: on a deep ladder of 199 prime-power zones with 117 certified handovers, the crossing is located between  $n = 461$  and  $n = 463$  — the  $n \approx 170$  figure was a fit artefact, but the wall is real, and  $n^* \approx 462$  is an upper bound — and the single warning splits into three separate walls (a measured margin wall, an arithmetic ladder wall set by the twin-prime pair  $521 \rightarrow 523$ , and a vacuous-requirement wall at  $n = 727$ ), while the handover mechanism itself never fails: all 117 steps certify at retention 1.000000, and the chain tears at the ratio, not at a step. The operating variable turns out to be the window depth, not the zone index. The next part acts on that reinterpretation and rebuilds the construction in a gap-coupled scaled frame ( $D$  tied to the local prime gap): two of the three walls fall *structurally* — the ladder death becomes a payable cell cost (the killer pair  $461 \rightarrow 463$  now certifies spectrally with the reserve opening by a factor  $\sim 1000$ ) and the requirement wall loses its depth dependence — while the margin wall proves *frame-invariant* at exponent  $-0.974$ : it is the substance of the requirement, not geometry. The hardness is thereby traded — three arithmetic walls against two analytic demands, the positivity of one limit operator and a convergence rate below the step reserve — with the prime-gap dependence disclosed rather than hidden. The next part settles the currency question, with a fundamental reinterpretation: the wall carries the same exponent in *every* admissible currency — it is substance — but it measures the *discretisation*, not the spectrum. Under refinement at a fixed window both leading eigenvalues carry the same cell-width power (the continuum window form has no gap), the positive floor survives only as a cancellation of relative size  $10^{-7}$ – $10^{-4}$ , and the T109 requirement chain turns out to divide by an artifact margin; the hardness balance is restated as four parts, led by a margin-free step certificate. The next part builds that certificate and the wall dissolves: the exact Schur complement (Albert 1969, with the Moore–Penrose inverse) certifies every ladder step margin-free — eleven of them beyond the old wall, to  $n = 1331$ , and all seven zones where the requirement chain tore — the wall was an  $O(1)$  numerator divided by an artifact floor, chains now stop only at the compute cap and never at a step, and the (R) demand itself proves grid-bound and is deleted rather than transported, leaving a two-object core:  $\varepsilon$  relative to  $\kappa$  via the exact Szegő identity, and transport of an  $O(0.1)$  Schur complement between grids. The next part splits that core: the transport across a real regrid certifies only mild refinement (a clean split at ratio  $\rho^* \approx 1.8$ , and for a principled reason — on nested ladders, where the transport error is exactly zero, the

Schur floor itself falls like  $\rho^{-1.7}$ , so no bound can undo the drop), while a two-scale compression breaks the compute cap: a certified margin-free step at  $n = 155\,921$  ( $117\times$  deeper), chains of ten steps, the stopper always cost and never a step — leaving the shortest list the program has produced: three points, of which only one is an inequality (the classical Szegő–Levinson prediction-error bound for one symbol). The next part executes the reformulation point of that list, and the induction step turns out to *be* a boundary process: partial minimisation over the interior (Haynsworth) compresses the window into a state  $(\Sigma, r, \sigma)$  of  $12 \times 12 + 12 + 1$  numbers that carries the global rank-one pole *exactly* (bordered elimination, no truncation) and updates by a Riccati-type recursion whose cost is independent of the window — one unbroken chain of 169 236 prepends ran to a window of 1.35 million cells ( $903\times$  the old cap) at a flat  $76\,\mu\text{s}$  per step, declared honestly as a cost-geometry demonstration and not a Weil certificate. What refuses compression is the *prime comb* itself: the archimedean tail decays cleanly, but the reflection comb couples every incoming cell to an interior cell at  $\alpha - \log n_j$  for every prime power in the window, so the full symbol never decays in the width and no sparse faithful state exists. In exchange, the one remaining inequality acquires a textbook shape:  $\varepsilon$  is the error of a piecewise-constant Galerkin method for the  $D$ -independent function  $2 \sinh(x/2)$ , Aubin–Nitsche duality predicts the measured exponent ( $\theta = 1.79 \Rightarrow s \approx 0.90$ ), and the target inequality is  $\varepsilon(\alpha, D) \geq c_0 D^\theta \alpha^\varphi$  across the prime-power jumps. The last part then makes that inequality *theorem-shaped*: the Galerkin reading becomes four exact identities, the lower-bound direction acquires a certified two-level chain that loses no power of  $D$  ( $\theta' = 1.74$  against  $\theta = 1.76$ ), the jumps acquire exact closed forms — with the correction that prime-power entries *raise*  $\varepsilon$ , so the factor-120 falls were a sweep artefact of the earlier run — and what remains is three named analytic lemmas about one symbol, two of them constants rather than rates. The next part measures those three lemmas against their classical addresses, and two of the three stand: saturation is an *identity* here; the Schur-floor lemma is rescued by the strengthened Cauchy–Schwarz shift onto the oscillation Gram, whose arithmetic aliasing symbol suppresses the prime comb quadratically — certified positive on half the measured windows, and on a 15 680-point FFT lever the certified floor rises logarithmically and crosses zero, so the failing windows were under-resolved rather than obstructed; and the corner lemma corrects itself — a log singularity is the boundary of all powers, so no power-law corner exponent exists to cite, and the certified chain is repaired at its last step, the rate now living on the oscillation mass. The missing analytic content is thereby exactly one statement: a  $D^{1.75}$  lower bound for that mass. The closing part then closes the *arithmetic* half of the remaining list as a theorem — the aliasing symbol is positive for every cell width below an explicit  $D_0(\alpha) = \exp(-(\Xi(\alpha) + B))$ , with  $\Xi$  the prime-power atom count and  $B = -1.0474$  universal — proves that the mass bound is *not* derivable from the identities already in hand, and reduces it, through the exact identity  $\kappa_{\text{end}} = 1/(1 + R)$ , to a single discrete Harnack inequality with a classical address; the chain contains no zeta input anywhere. The final part then *explains* that inequality: the two increment families are the odd and the even half of one sequence, offset by a single fine cell, so given one sign on the corner cells  $|R - 1| \leq 0.047$  holds *unconditionally* by a two-line pairing argument — while the per-cell version of the inequality is provably false, and the honest defect count rises from three to four, because the frame’s window grows with the zone (the unconditional  $D_0$  route covers not one handover of the real ladder) and the strengthened Cauchy–Schwarz constant decays with the width. The series then closes by *assembling* what it built: on one common resolution, thirty consecutive prime-power atoms form a run in which the output of each certified handover *is* the input of the next (the incoming atom block is bit-exactly zero, the composition residual  $2.9 \times 10^{-14}$ ), all five stages complete on 52 of 52 ladder zones with 430 completed Cholesky certificates, and the load-bearing spine turns out to be 96.2% identity or Cholesky certificate with exactly one hypothesis — an accounting convention — so that the Harnack pair leaves the spine and survives only as a second, independent route. What remains is stated as such: not size but *uniformity* in the zone index, quantified in a six-point distance statement (no zeta input, RH unused in either direction, a finite run bounded by the smallest log-gap it contains, a measured ladder that is not all zones, uniformity as the open item, and the resulting object typed as what it is — a lower bound in numerical analysis, not a statement about zeros). A second phase opens past the finale: its first part dissolves the segment-seam question into a monotone free-resolution construction (re-gridding only at record-small gaps, all twelve affordable real seams certified), turns the continuum step into an exact fractional-Dirichlet identity, and leaves a sixteen-point full-proof map containing exactly *two* genuinely new inequalities — a direction lemma and a zone-uniform seam floor. **No claim of progress on the Riemann Hypothesis**

is made anywhere in this document, and none of the sandbox findings moves any ledger marker. The document is deliberately symmetric about failure: a dedicated section lists the readings and routes that were refuted — an essential-singularity reading, a  $C/g$  proxy law, a density chain, an accumulated invariant with self-propagation, the Landau–Widom anchor, the symbol route for the odd channel, three certificate candidates for the avoidance norm plus the Combes–Thomas certificate at the boundary, and the naive margin route — because in a program of this shape the kill list is the more reliable output.

**[E]** exact / machine-proven

**[C]** conditional

**[O]** open / axiom

**[X]** kill test

## 1 How to read this document

### Scope and status — read this before anything else

1. **No RH claim.** Nothing in this document claims progress towards, or evidence for, the Riemann Hypothesis. Where an equivalence to RH appears (Section 5) it is an *equivalence typing*: a statement about which object carries the difficulty, produced by an exact identity, and deliberately *not* a statement that the difficulty has been reduced.
2. **Two layers, never mixed.** Claims marked `[ledger]` are load-bearing: they have a row in `verification/status_ledger.csv`, a module in `verification/`, and a script citation here (the blue `[verification/...]` tag). Everything marked `[sandbox]` is exploratory work in `experiments/tfpt-discovery/`; it carries no ledger row, no marker, and no authority. Section 4 is the only section that is entirely `[ledger]`.
3. **Status markers are quoted, never upgraded.** The markers `[E]` `[C]` `[O]` `[X]` shown in Section 4 are read verbatim from the ledger rows of v535–v541. This document introduces no new markers and moves none.
4. **Fits are declared as fits.** Every scaling law, slope, exponent and extrapolation quoted from the sandbox is a fit to finitely many measured windows and is labelled as such at the point of use. Sixteen resolved zones are sixteen zones, not all zones.
5. **Classical results are named as classical.** Kneser neighbours, Eichler/Witt, Shimura, Waldspurger, Cohen, Weil/Guinand, Bombieri, Yoshida, Connes–Consani, Gronwall/Robin, Douglas, Slepian–Landau–Pollak, Sherman–Morrison/Woodbury, Grenander–Szegő and the rest are classical and are used as such; the compiler-native part is the wiring, not the theorems.
6. **Numbers are literal.** Every numeric value in this document is transcribed from the research diary `experiments/next.txt` or from a ledger row. Figures that could not be drawn from literal data are marked *schematic* in their caption, and the caption states exactly which quantities in the figure are literal.

The diary parts are cited as **T** $n$  (part  $n$  of the research diary, chronological,  $n = 11 \dots 128$ ). The probe file, the contract name, the verdict and the check count of every part are listed in Appendix A. Section 2 explains why a physics compiler meets the primes at all and Section 3 states the framework interpretation that motivated the whole line — a self-dual lattice, its functional equation, and a mirror that returns at the far end of the program; Section 4 is the verified layer; Sections 5, 6 and 7 are the sandbox program in three acts; Section 8 is the kill list; Section 9 is the series balance (cascade, certification degree, the honest RH distance); Section 10 is the opening arc of the second phase (T126–T128: the seam architecture, the two inequalities dissected, then the four-point list worked cheapest first); Section 11 is the status statement.

## 2 Motivation: why a compiler meets the primes

TFPT is a discrete compiler with two inputs: the seam constant  $c_3 = \frac{1}{8\pi}$  (P1) and the carrier rank  $g_{\text{car}} = 5$  (P2). They force the lattice completion  $D_5 \oplus A_3 + \mu_4 \Rightarrow E_8$ , and the readouts of the Standard-Model skeleton are taken off that completion. None of this is about number theory. The prime line begins with an accident of the *glue*.

The completion is not  $D_5 \oplus A_3$  but  $D_5 \oplus A_3$  *plus* a glue group  $\mu_4$  — a cyclic group of order four. A cyclic group of order four carries a quartic character, and the quartic character on  $\mathbb{Z}$  is  $\chi_{-4}$ , the non-trivial character mod 4. Once the census of glue vectors is counted *with that sign*, three things happen at once, and all three are elementary consequences rather than choices:

1. The signed glue count is a **theta series**. Signed counting of a quadratic form is a theta series by definition; here the compiler’s own count  $\theta_3^2$  is exactly the Gaussian-integer theta,  $r_2(n)$ , so the census counts  $\mathbb{Z}[i]$ -ideals (T85).
2. The character makes the count a **Dirichlet object**.  $\chi_{-4}$  is the character of  $\mathbb{Q}(i)$ ; the associated  $L$ -function is the first Dirichlet  $L$ -function, and the split law “ $p = a^2 + b^2$ ” becomes a compiler-visible fibre property of each prime (T21, channel B:  $\chi_4$  fibre  $\iff p = a^2 + b^2$ , verified 81/0/87/0 on  $p < 1000$ ;  $\chi_3$  fibre  $\iff p = a^2 + 3b^2$ , 80/0/86/0).
3. The shell counts of the census carry **Hecke eigenvalues**. That is the surprise, and it is the content of the first load-bearing module.




Figure 1: The chain by which the compiler acquires arithmetic. Nothing here is chosen for number theory: the glue group is forced by the completion, the quartic character is the character of a cyclic group of order four, and the signed count of a quadratic form is a theta series. The green nodes are where the line becomes load-bearing (Section 4).

The first question of the series was therefore not “does this prove anything about primes” but “which of these links are *mechanism* and which are pretty coincidences”. Parts T11–T13 established the signed glue theta, its  $L$ -series and the glue-character atlas (19/19, 16/16, 11/11); T14–T19 tested six candidate contracts in parallel and produced two self-corrections against the project’s own earlier readings (T14 DEFLATED: the measured  $2\pi$  is the universal Bisognano–Wichmann/Unruh conversion, not “the self-dual width”). The cross-finding of that batch is worth stating because it survived the entire series: the structure is exact at the census/lattice level (weight 4), and the gap is always the *same* gap — the canonical transport to the  $GL(1) / \xi(s)$  world (multiplicity one, weight 1/2, the archimedean place, the derivation of the Hecke action). One meta-contract, ZETA.CENSUS.T0.GL1 [O].

Parts T20–T25 then killed, in four stages, the slogan “the archimedean term comes from the seam”: the raw seam density of states is flat/slightly decreasing (loglog slope  $\approx -0.40$ ) while the classical digamma kernel grows logarithmically (measured 0.15917 against  $1/(2\pi) = 0.159155$ , rms  $3 \times 10^{-6}$  for  $t \geq 20$ ) — KILLED-AS-NAIVE (T20, 25/25); the non-DOS successor located the missing logarithm in the interval cut but failed the one-constant dictionary at  $2/\pi$  (T22, PARTIAL); and both dictionary closures plus the scattering-phase observable died (T23–T25, 23/23 + 25/25 + 29/29). This matters for the rest of the document: the archimedean place was never going to come for free, and when it finally arrived (T82) it arrived from *inside* the same family, not from the seam.

### 3 The big picture: a self-dual lattice and its own mirror

#### What this section is

A **framework interpretation**, not a result. It says how the pieces fit together and why the line was pursued for a hundred parts. Every factual statement it uses is either classical (and named as such), one of the seven modules of Section 4, or a sandbox probe with its part number. Nothing here is an argument about the Riemann Hypothesis, nothing here is used later as a premise, and the picture itself carries no status marker.

#### 3.1 (i) The zeta function is literally in the geometry

Count the  $E_8$  lattice points shell by shell. The even shells satisfy  $N(2n) = 240 \sigma_3(n)$  — direct lattice counting gives 240, 2160, 6720, 17 520, 30 240, 60 480 for  $n = 1 \dots 6$  (T6, one of the parts that predate the indexed range of Appendix A) — and  $\sigma_3$  is multiplicative, so the counting function of the lattice is a Dirichlet series with an Euler product:

$$\sum_{n \geq 1} \frac{N(2n)}{n^s} = 240 \sum_{n \geq 1} \frac{\sigma_3(n)}{n^s} = 240 \zeta(s) \zeta(s-3) \quad (\text{at } s = 5 : 1.705672 \text{ vs } 1.705678).$$

The Riemann zeta function stands in  $E_8$ 's own counting function, with a simple pole at  $s = 4$  of residue  $240 \zeta(4) = 8\pi^4/3$  (T12). This is classical lattice theory — Eisenstein series of weight 4 — and it is stated here only to fix where the arithmetic enters: not through a modelling choice, but through the object the compiler already had. The contrast is instructive: the same multiplicativity *fails* for the  $D_5 \oplus A_3$  sublattice alone (52, 576, 1584, 16 128 with  $576 \cdot 1584 \neq 16\,128 \cdot 52$ , T6). Without the glue there is no Euler structure in the count.

#### 3.2 (ii) The glue carries the character of the Gaussian integers

The completion is  $D_5 \oplus A_3 + \mu_4$ , and a cyclic group of order four carries a quartic character, which on  $\mathbb{Z}$  is  $\chi_{-4}$  — the character of  $\mathbb{Q}(i)$ . Two consequences make the primes a *fibre property* of the compiler rather than an application of it. First, the split  $p \equiv 1$  versus  $p \equiv 3 \pmod{4}$  is visible to the census directly: the  $\chi_4$  fibre holds exactly for  $p = a^2 + b^2$  (81/0/87/0 on  $p < 1000$ , T21). Second, the character kills a pole. Of the three character channels of the  $E_8$  counting function, the trivial one is  $240\zeta(s)\zeta(s-3)$  with its pole at  $s = 4$ , the spinor one is  $64(1 - 2^{-s})\zeta\zeta$  and keeps it, while the  $\mu_4$ -signed channel  $16 \cdot 2^{-s}\eta(s)\eta(s-3) - 8L(f_8, s)$  is *entire* ( $L_c(4.001) = -6.976$ , finite; T12). The signed channel is the only pole-free one — exactly the  $\zeta$  versus  $L(\chi)$  pattern.

#### 3.3 (iii) Hecke structure is an output of geometry, not an input

The step that made the line load-bearing is v535: the Hecke action on the census is produced by *pure lattice adjacency* — Kneser  $p$ -neighbours — with no modular input anywhere in the derivation. The marked neighbour sum is exactly an affine Hecke element,  $\nu_p = a \text{Id} + b T_p$ , and the cusp eigenvalues come out of the glue census as  $a_p = -c(p)/8$ . Prime structure is what the geometry *emits*; nothing about primes was put in. The first four modules of Section 4 make this precise and end in one finite relative-trace identity of which the first three are projections.

#### 3.4 (iv) The core point: self-duality, and a mirror that keeps coming back

$E_8$  is unimodular — self-dual — and self-duality is exactly the hypothesis of Poisson summation. The lattice theta sum therefore satisfies an exact functional equation,  $\sum_x e^{-t|x|^2} = (\pi/t)^4 \sum_x e^{-(\pi^2/t)|x|^2}$  (verified to 25 digits, T12), whose unique self-dual width is  $t = \pi$ . A functional equation is a *reflection*, and the critical line is the fixed axis of that reflection. That is the geometric content of the picture: on the axis sits the symmetry, and beside it sits everything the lattice actually carries.

Two findings, at opposite ends of the program, put the same mirror on the table twice.

- **The reflection product is the transport operator** (T83, `fe_symmetrization_probe.py`, 27/27, INVARIANT-NULL). The involution algebra was computed exactly:  $\zeta$ 's functional equation is the reflection  $J_{1/2}$  and is absorbed in evenness (on autocorrelations it acts as the identity), while the family's functional equation is  $J_1$ . Their *product*  $J_1 \circ J_{1/2} = e^{\pm u}$  is the unit-line shift — the difference of the two centres *is* the transport operator, algebraically. The two reflections generate an infinite dihedral ladder; the value side is invariant under the whole ladder, the spectral cone only under its own reflection, which is what anchors it to the critical line; and  $\text{Fix}(J_1) \cap \{\text{Weil cone}\} = \{0\}$  exactly, so no symmetrisation can relieve I5. The verdict was a null with a structural finding: the wall is *transversal* to the functional equation, not positional.
- **The same mirror sits in the last hard residue** (T106, Section 7). The mirror parity  $J$  with  $JQJ = Q$  and  $JB = -BR$  splits the Weil pole exactly into  $ss^\top - tt^\top$ : a positive rank-one lift in the  $J$ -even channel and a negative rank-one pressure in the  $J$ -odd one. The even channel closes; everything that is still hard lives in the odd half — the anti-symmetric part of the same reflection.

Stated as a picture and nothing more: the recurring question of this line is whether the self-dual object is entirely at peace with its own mirror. That sentence is an interpretation of where the difficulty sits, in the same spirit as the classical reading of the critical line as a symmetry axis. It is *not* a reformulation of the Riemann Hypothesis, it is not used as a premise anywhere below, and it produces no claim.

### 3.5 (v) What the program did with the picture

It converted it, step by step, into smaller and more explicit objects. The transport ledger closes exactly (v541), which leaves the single coupled inequality I5; I5 is by identity equivalent to Weil positivity, an equivalence *typing* and not a reduction (Section 5); the positivity question is then reorganised as a per-zone handover induction (Section 6); the induction is certified ingredient by ingredient until only one Loewner statement (R) on the  $J$ -odd half of a window remains; and (R) reduces first to one scalar ratio, then to two scalars, and finally — with both scalars certified — to one strict-positivity margin input (Section 7). The full cascade, with the type of every arrow marked, is Figure 8; the reason it is a route and not a proof sketch is stated there and in Section 11.

## 4 The verified layer: seven modules

This section is entirely [ledger]. Each subsection states the claim as the ledger states it, cites the module, and quotes the ledger's own status markers. Nothing here depends on Sections 5–7.

| The seven modules and their ledger markers                                                                                                                                                                                                                                                                                                                                           |                       |                                                  |                 |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------|--------------------------------------------------|-----------------|
| Module                                                                                                                                                                                                                                                                                                                                                                               | Claim id              | Content                                          | Markers         |
| v535                                                                                                                                                                                                                                                                                                                                                                                 | HECKE.GEOM.01         | Hecke from geometry on the $E_8$ census channels | [E] + [C] + [O] |
| v536                                                                                                                                                                                                                                                                                                                                                                                 | HECKE.GEOM.EICHLER.01 | Eichler trace layer at the $E_8$ lattice         | [E] + [C] + [O] |
| v537                                                                                                                                                                                                                                                                                                                                                                                 | HECKE.GEOM.HALFINT.01 | half-integral bridge                             | [E] + [C]       |
| v538                                                                                                                                                                                                                                                                                                                                                                                 | HECKE.GEOM.RTF.01     | compiler relative trace identity                 | [E] + [C]       |
| v539                                                                                                                                                                                                                                                                                                                                                                                 | RTF.GNS.WEIL.01       | Weil structure of the compiler family            | [E] + [C]       |
| v540                                                                                                                                                                                                                                                                                                                                                                                 | RTF.GNS.AMP.01        | amplitude route and positive linear carrier      | [E] + [C]       |
| v541                                                                                                                                                                                                                                                                                                                                                                                 | RTF.GNS.LEDGER.01     | matching lemma and transport ledger package      | [E] + [C]       |
| Marker key: [E] exact / machine-proven · [C] conditional (holds under named hypotheses) · [O] open / axiom · [X] falsifiable kill test. Per-claim fine type in <i>verification/status_ledger.csv</i> .<br>Every row carries NO marker moves in the ledger. The [O] on v535 and v536 is the same open gate throughout: $\text{weight-4} \rightarrow \text{GL}(1) \text{ transport}$ . |                       |                                                  |                 |




Figure 2: *Schematic*. The self-duality picture. *Literal*: the Poisson identity and its unique self-dual width  $t = \pi$  (25 digits, T12); the algebraic identity  $J_1 \circ J_{1/2} = e^{\pm u}$  and  $\text{Fix}(J_1) \cap \{\text{Weil cone}\} = \{0\}$  (T83); the exact pole splitting  $ss^T - tt^T$  at  $10^{-18}$  (T106). *Schematic*: the geometric arrangement, and the reading of the axis as “order” with the lattice’s arithmetic content beside it. The bottom line is an interpretation and carries no status.

#### 4.1 Hecke from geometry (v535)

The Hecke structure of the  $\mu_4$ -glue census is lattice-native: it comes out of Kneser  $p$ -neighbour geometry, not out of modular input [verification/v535\_hecke\_from\_geometry.py] (25 checks,  $\sim 11$  s; consolidated promotion of T27/T31/T32).

- **Kneser carries Hecke** [E]:  $\#_{\text{iso\_pts}} = p^7 + p^4 - p^3$  and  $\#_{\text{iso\_lines}} = \sigma_3(p) \cdot \#\mathbb{P}^3(\mathbb{F}_p)$  identically and by full enumeration at  $p = 2, 3, 5, 7$  (135/1120/19656/137600 lines). The Eisenstein eigenvalue  $\sigma_3(p) = 1 + p^3$  is an exact factor of a lattice-native correspondence count.
- **The marked neighbour sum is an affine Hecke element** [E]:  $\nu_p = a \text{Id} + b T_p$  exactly (overdetermined, residual 0), with  $b = \sigma_3 + a_p$ ; the cusp rule  $a_p = b - \sigma_3$  recovers  $a_3 = -4$  and  $a_5 = -2$ .
- **Census redundancy is 2-adic oldform structure** [E]:  $\dim V = 7 = 5+2$ ,  $\pi_{\text{cusp}} = (28 - T_3)/32$ , only 2-adic levels occur.
- Fences carried *inside* the claim: the  $p = 5$  geometric  $S[1]$  block is a T31-validated *freeze* [C]; the weight-4  $\rightarrow \text{GL}(1)$  transport is [O].

#### 4.2 The Eichler trace layer (v536)

The smooth (Eisenstein) background and the coherent (cusp) interference separate exactly [verification/v536\_eichler\_trace\_layer.py] (23 checks,  $\sim 30$  s; promotion of T33/T36/T37/T42). The Witt layer gives  $\lambda_{\text{Eis}} = \sigma_3(\#\mathbb{P}^3 - \sigma_3) = L - \sigma_3^2$  identically (anchors 336, 3780, 19264 at  $p = 3, 5, 7$ ) [E], and the Eichler identity  $\lambda_{\text{geom}} = \lambda_{\text{Eis}} + a_p^2$  holds two-sidedly (352/3784/19840). The signed cusp reading is  $a_p = -c(p)/8 \Rightarrow (-4, -2, +24)$ , and the assembler  $R(p) = a_p^2$  is exact for  $p \leq 31$ . The  $p = 5$  type-(ii) block is again a validated freeze [C]; the transport gate stays [O].

### 4.3 The half-integral bridge (v537)

In the 70-element compiler weight-5/2 theta monoid there is *exactly one* object  $g$  with  $\text{Sh}_{t=2}(g) = -8f_8$  coefficientwise [E], it is an exact  $T(p^2)$ -eigenform with the weight-4 eigenvalues  $(-4, -2, 24, -44, 22)$  [E], and  $U_4(g) = 0$  places it outside the Kohnen plus space — a scope fence, not a defect [E] [verification/v537\_halfintegral\_bridge.py] (20 checks,  $\sim 90$ s; promotion of T38/T41/T44/T45). The Waldspurger/Baruch–Mao constant  $R = 23.1873585645\dots$  is constant on ten discriminants  $d \equiv 1 \pmod 8$  (spread  $\leq 10^{-12}$ ), typed [C] because it is an AFE-numeric measurement.

### 4.4 One finite relative trace identity (v538)

The strongest structural statement of the verified layer: v535, v536 and v537 are three *projections* of one finite relative-trace identity of Jacquet type [verification/v538\_relative\_trace\_identity.py] (18 checks,  $\sim 14$ s; promotion of T53).  $\text{Tr}_V(\nu_p \circ \pi)$  is computed twice, independently — geometrically from the glue census with no modular input, and spectrally from  $\eta(2\tau)^4 \eta(4\tau)^4$  with no geometry — and the two agree exactly for all  $(p, \pi) \in \{3, 5, 7\} \times \{\text{Id}, \pi_{\text{Eis}}, \pi_{\text{cusp}}, \pi_{E_4}, \pi_{f_8}\}$  (anchors, e.g.  $p = 7$ : full trace 727680,  $f_8$  channel 19840) [E]. The Eichler split *is* the orbit decomposition: principal orbit  $a \text{Id} \rightarrow \lambda_{\text{Eis}}$ , elliptic orbit  $b T_p \rightarrow a_p^2$ , cross terms cancel exactly. And the period side is literally compiler geometry:  $g$  is the quaternary lattice form  $n = (x^2 + y^2)/2 + 2z^2 + u^2 + 2w^2$  with sign  $(-1)^{|u|+|w|}$ , giving  $[\text{signed lattice count}]^2 = R |d|^{3/2} L(f_8 \times \chi_d, 2)$  at relative  $10^{-12}$  [C]. The identity is *finite*; the infinite RTF is explicitly open.

### 4.5 The Weil structure of the family (v539)

The family carries a fully identified Weil structure *up to two explicitly isolated obstructions*, and the obstructions are promoted content, not a footnote [verification/v539\_weil\_structure\_family.py] (25 checks,  $\sim 6$ s; promotion of T55/T61/T62/T63). The canonical GNS pairing is a direct integral over the Gelfand spectrum with fibre mixing  $\leq 5.04 \times 10^{-16}$ ; the trivial Sato–Tate isotype is algebraically a  $\text{GL}(1)$  core,  $G_0 = (1 + Y)/(1 - Y)$ ; the prime side is an exact linear relation to shifted  $\zeta$ -Weil forms at relative  $\leq 9 \times 10^{-16}$ . The two obstructions are (i) a sign/doubling residue (metaplectic,  $\text{sym}^2$ ) that survives every preregistered canonical projection, and (ii) a Plancherel correction which is irreducibly non-automorphic. Finite-class positivity of  $Q_{\text{fam}}$  is measured [C] ([4.369, 11.486]); dense-class positivity is open. The ledger row says it in one line: *not almost-RH*.

### 4.6 The amplitude route and the positive linear carrier (v540)

The route out of the square plane [verification/v540\_amplitude\_linear\_carrier.py] (34 checks,  $\sim 3$ s; promotion of T67–T72). An amplitude Dirac operator  $D = \begin{pmatrix} 0 & V \\ V^\dagger & 0 \end{pmatrix}$  with  $V[d, m] = \sqrt{w_d} b(d) \chi_d(m)$  exists and is Hecke-equivariant [E]; the geometric polarisation  $b = N_+ - N_-$ ,  $\Theta = N_+ + N_-$  is lattice-exact with  $N_\pm \geq 0$  for all  $n \leq 50\,000$  [E]; the Cohen seed  $\Theta(d) = -48L(-1, \chi_d)$  is exact-rational on 159 live  $d$  [E]; the prime side is a *pure plus* combination  $Q = Q_\zeta(g_-) + Q_\zeta(g_+)$  at relative  $2 \times 10^{-15}$  [E]; the functional equation of the linear family holds at relative  $1.4 \times 10^{-40}$  [E]. The open boundary is *inside* the claim: the gap functional  $\lambda^*$  on  $n \equiv 6 \pmod 8$ , with an exact Farkas witness (0.0511), is a named limit, and positivity is asserted only in the Euler region.

### 4.7 The matching lemma and the transport ledger (v541)

The largest module of the line [verification/v541\_matching\_lemma\_ledger.py] (33 checks,  $\sim 10$ s; promotion of T78–T82 and T85; every core check recomputed, not cited).

- **Window proof** [E]: the matching lemma is proved exact-integer on  $[4, 10^6]$  —  $7S(j) < 40A(j)$  at every atom, 0 violations over 939 870 clash atoms, no-overflow proven in exact integers, 268 big-integer rechecks; exact margin  $X = 22242575693/270725400000 = 0.082159$  at  $j^* = 554400$ ;

$\rho_{\text{crit}} = 1.143989 > 21/20$ . A maximal-greedy identity makes the certificate uniform in the target set and the greedy weights.<sup>1</sup>

- **Transport ledger [E]:**  $Q_{\text{Weil}} = Q_{\text{cert}} + \Delta_{\text{arch}} + \Delta_2$  closes with residual  $4.4 \times 10^{-16}$  on a 24-row battery, and  $\Delta_{\text{pole}} \equiv \Delta_{\text{conv}} \equiv 0$  are *proven*, not measured (kernel collapse in `sympy`, exact scale invariance).
- **Character-exact envelope [E]:**  $-\psi = (\chi_{-4} + \frac{1}{4}\chi_8 + \frac{1}{4}\chi_{-8})\Theta$  coefficient-exact for all odd  $n \leq 50\,000$ ; confinement holds as a *set equality* on  $10^6$ .
- **Archimedean term internal [E]:** the Legendre duplication bridge  $(2\pi)^{-s}\Gamma(s) = \frac{1}{2}\Gamma_{\mathbb{R}}(s)\Gamma_{\mathbb{R}}(s+1)$  and the kernel identity  $k_{\zeta} = K_{\text{fam}} - K_{\text{shift}}$  at  $7.9 \times 10^{-31}$  pointwise.
- **Frontier facts [C]:** unlifted coherent chain crossing at  $k^* = 14$  ( $\log_{10} N^* = 23.1$ ); Robin tail factor 6.16 at  $J = 10^6$  and divergent — the constant route stays open.
- **Named limits as content.** Two limits are part of the claim rather than hidden by it: one classically-shaped open lemma (correlated cancellation on credit-rich non-coherent tail atoms), and the I5 equivalence typing discussed next. The row states explicitly: *NOT “almost RH”*; ZETA.HP.CARRIER untouched.

## 5 The localization program (sandbox)

### Provenance

Everything from here to Section 11 is [sandbox]: exploratory probes in `experiments/tfpt-discovery/`, no ledger row, no marker, no authority. The findings are reported because the *shape* of the residual is the useful output, not because any of them is established.

### 5.1 I5: the ledger closes and one inequality is left (T79)

The transport wall — the thing that had blocked the line since T20 — was fully decomposed in T79 (`transport_ledger_probe.py`, 22/22). The key observation is that the odd-prime side of the Weil functional agrees *exactly* with the certified plus combination (relative  $2.6 \times 10^{-16}$ ): the prime side was never the wall. What remains after the ledger closes is a minimal transport set of two exact identities (proved), two classically-provable bounds, and exactly one coupled inequality:

$$(I5) \quad Q_{\text{cert}}(h) + \Delta_2(h) + A_{\text{fam}}(h) - A_{\text{shift}}(h) \geq 0 \quad \text{for all autocorrelations } h.$$

I5 is the prime  $\leftrightarrow$  archimedean coupling. By the closed ledger it is, as an identity, equivalent to Weil positivity and hence to RH. The diary is explicit about what that means: *the built-in safeguard fired*. RH does not follow from finite bookkeeping; the hidden infinite step is exactly I5 — and the value of T79 is that the step was *found and named* rather than concealed. This is an equivalence typing, not progress.

T82 then changed I5’s *type* (`heat_arch_probe.py`, 22/22, ARCH-INTERNAL):  $\Delta_{\text{arch}}$  is exactly the internal  $\Gamma$  difference  $A_{\text{fam}} - A_{\text{shift}}$  via Legendre duplication, verified on 18/18 battery rows at max relative  $6.6 \times 10^{-15}$ . I5 stopped being a coupling between two worlds and became a *self-consistency statement about one heat family*. Again: a type change is not a proof.

T83 (`fe_symmetrization_probe.py`, 27/27, INVARIANT-NUL) then asked the obvious question of the self-duality picture (Section 3): can the functional equation itself do the work? It cannot, and the way it fails is informative. The  $\zeta$  reflection  $J_{1/2}$  acts as the identity on autocorrelations — the symmetry is already absorbed in evenness, and  $Q_{\text{Weil}}(Jh) = Q_{\text{Weil}}(h)$  holds floating-point-exactly on 15 test functions — and the family reflection  $J_1$  moves nothing in the value geography (5/24

<sup>1</sup>Prior-art note (literature pass, 2026-07-27): the coefficient laws L1–L4 that T77/T78 use to shape the lemma are specialisations of Cohen 1975 (Math. Ann. 217: the Hurwitz class-number series  $H(r, n)$  and the Cohen–Eisenstein series of weight  $r + \frac{1}{2}$ ) and of Pei–Wang 2003 — corollaries of known theory, not discoveries. The restricted divisor inequality itself remains the independent part.

overlap invariant,  $\lambda^*$  covariant-exact) while pushing  $m_+h$  out of the Weil cone in closed form, with  $\text{Fix}(J_1) \cap \{\text{Weil cone}\} = \{0\}$  exactly. The structural finding underneath the null is the one quoted in the big picture: the product of the two reflections is the unit-line shift, i.e. the transport operator, and the wall is transversal to the functional equation rather than positioned by it. Symmetrisation is therefore not a route; the same mirror returns, as an obstruction, in T106.

## 5.2 The Connes dictionary (T87) and where positivity is anchored

T87 (`i5_connes_dictionary_probe.py`, 22/22, DICTIONARY-EXACT-CORE) identified the I5 one-family form as a compiler instance of the semi-local Connes structure. Four exact matches: the two-shell orbit route equals Connes' expression at relative  $6.3 \times 10^{-16}$ ; the orbit kernel of Connes–Consani (2021) is literally the compiler's Poisson ladder at  $\rho = e^u$  (`sympy` plus  $1.7 \times 10^{-50}$ );  $h_+$  is by definition the compiler kernel  $k_\zeta = K_{\text{fam}} - K_{\text{shift}}$ ; and  $k_\zeta = 2\theta'_{\text{RS}}$  — the derivative of the *Riemann–Siegel phase*, at relative 0.0. The dihedral shift is the scaling group  $\mathbb{R}_+^*$ .

The most useful part of T87 is not a match but a *boundary*. The two positivity programs are complementary: Connes–Consani anchor positivity *before* the primes (the prime-free Sonin window, support  $(\frac{1}{2}, 2)$ , i.e.  $u \in (-\log 2, \log 2)$ ), the compiler anchors it *after* the primes (atom certificates, all supports). The window boundary is exactly the first prime atom  $u = \log 2$  (`sympy`-exact; on narrow rows the prime side,  $Q_{\text{cert}}$  and  $\Delta_2$  all vanish identically). The sectors *touch* and do not overlap; the I5 coupling lives in the crossing region between them, outside both proved sectors.

## 5.3 The tightness map (T88) and the crossing (T89–T91)

T88 (`i5_tightness_probe.py`, 27/27) mapped where I5 is actually tight: a band-limited zero plateau (classical), safe zones, and *eight* near-vertical tight curves whose spacing follows the  $\Gamma$ -side density law (ratio  $1.01 \pm 0.08$ , computed zero-free — no spectral identification is claimed or made). There are no negatives on genuine autocorrelations; the earlier T76 negatives were the shadow of missing archimedean/ $p = 2$  bookkeeping, exactly as the ledger predicted.

T89 (`sonin_crossing_probe.py`, 19/19) established that the window compression of the Weil form *is* Bombieri's classical object, and settled the sharpness question: the boundary at  $\log 2$  is *not* sharp. The margin falls smoothly from  $a \approx 0.5$  towards the floor with no structure at  $\log 2$ ; the prime-free sector still has margin 0.557 at its own boundary and keeps  $\geq 0.25$  of the reference comfort to  $a \approx 0.75$  (+0.057 beyond the proved boundary). The residual subspace grows softly  $0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 4$  (out of 32) and carries the global minimum. T90 (`residual_dissection_probe.py`, 17/17) made those residual vectors explicit —  $n$ -stable to  $\leq 2.1^\circ$ , closed Gaussian  $\times$  cos/Gaussian  $\times$  sin fits with 99+% capture — and decided the core question: the residual is *not* just pole coupling (3/10 vectors are controlled by no structure at all; beyond  $a \approx 0.85$  genuinely pole-free atom  $\leftrightarrow$  arch content remains).

T91 (`thin_band_analytic_probe.py`, 19/19) identified the band as the *one-atom zone* and found the mechanism that dominates the rest of the document: beyond an inner edge the prime atom  $u = \log 2$  becomes *load-bearing* — the prime *rescues* the form where the archimedean margin is exhausted (rescue term  $+0.22 \dots + 0.41$ , purely odd sector). Two uncertainty constants were decided in closed form:  $a \cdot t_{\text{rms}} \rightarrow \pi$  (Wirtinger, 0.3%) and  $a \cdot t_{\text{cent}} \rightarrow 2\text{Si}(\pi) - 4/\pi = 2.4306$  (two independent routes,  $10^{-25}$ ).

The band that these three parts map out is narrow and its landmarks are explicit. Reading outwards from the classical side, the map is: classical edge 0.34657, prime-free crossing 0.3723–0.3763, band top  $\alpha^* = 0.46265$ , second atom  $\log 3/2 = 0.54931$ . Everything in Sections 6 and 7 happens inside that interval, between the first and the second prime atom, and the four numbers are the coordinates used throughout. The map was re-derived by a third independent implementation (T93, `band_reconciliation_probe.py`, 41/41) before it was used further.

An attempt to convert the band into a proof was made and did not succeed. T92 (`zone_extension_proof_probe.py`, 36/36, T-SKELETON) certified something small and real —  $Q \geq 6.7 \times 10^{-12} > 0$  on an 8-dimensional sine-window subspace across the whole band, with an

error certificate and a 1057-point covering — but not the lattice zone extension it was aiming at, and it named two independent walls:  $\lambda_{\min}$  collapses geometrically (factor 11.4 per mode), and the high-mode complement would need  $N^* \gtrsim 5069$  modes. The finding, and the reason it is quoted here rather than buried, is a statement about instruments: *the finite-block route is structurally the wrong instrument for this target, not an under-resourced one*. That is what motivated the change of approach in Section 6.

## 5.4 The blind prime demonstration (T94)

A separate, deliberately modest exercise. T94 (`prime_blind_demo_probe.py`, 16/16, BLIND-100) used the exact geometric criterion “odd  $n$  is prime  $\iff r_4(n) = 8(n+1)$ ” (Jacobi 1834, on the  $\mathbb{Z}^4 = \text{rank-}|\mu_4|$  theta tower) to predict all 753 primes of the never-seen window  $[1\,000\,001, 1\,010\,000]$  — count and MD5 declared *before* any ground truth, then evaluated: exactly 753, zero false positives, zero false negatives, 100.00% on all  $10^4$  integers, with an AST-enforced prediction path containing no divisibility test (no `isprime`, no sieve, no `%`, `//` or `/`). The cost is part of the result: a sieve wins by  $\sim 820\times$  (1 ms against 0.83 s). This is structure completeness, not an algorithmic advance — and the *spectral* prediction (zeros  $\rightarrow$  primes) remains bound to I5/RH and is not claimed.

# 6 The handover induction (T95–T101)

The band work left a picture: each prime atom rescues its own zone and hands over to the next. T95–T101 turned that picture into an induction, and then measured whether the induction can run.

## 6.1 The relay is real (T95, T96)

T95 (`continuum_extension_probe.py`, 28/28) proved the C1 chain unconditionally ( $|h_f(\log 2)| \leq 1/2$  by disjoint supports with margin 0.2305; the windowed symmetrised shift has  $\|S\| = 1/2$  exactly,  $1.1 \times 10^{-16}$ ) and found that the binding mechanism is *not* the two-bump extremal (share 0.046) but the atom rescue: above  $\alpha_{\text{free}} = 0.371654$  the prime-free minimiser has  $h(\log 2) < 0$  down to  $-0.328$ , so the atom adds positivity exactly in the losing direction.

T96 (`relay_test_probe.py`, 21/21) is the model of a self-correcting run. The apparent “edge” of T95 at  $\alpha^*$  was a map artefact: the value  $+7.8 \times 10^{-6}$  reproduced exactly, but it sits on a smooth, strictly positive curve with no feature at all ( $\lambda_{\min} > 0$  on the whole of  $[0.38, 0.86]$ , decay law  $\lambda_{\infty} \sim \exp(-49.2\alpha)$ ). The *relay*, however, is real — and it lives entirely in the counterfactual: remove the  $\log 3$  atom and  $\lambda_{\min}$  crashes to  $-0.445$ , with the loser being exactly the anti-double-bump at distance  $\log 3$  (alignment  $-0.99$ ) and a rescue identity at  $5 \times 10^{-15}$ . All four handover windows are positive: each atom arrives strictly *before* it is needed ( $+0.025/ +0.009/ +0.011/ +0.007$  for  $k = 1.4$ ). Three consequences were recorded: it is a margin problem, not an edge problem; numerics as a witness is exhausted beyond  $\alpha \approx 0.55$ ; and the counterfactual is the proof target, because it works on  $O(0.1)$ -sized quantities instead of  $10^{-6}$ .

## 6.2 The induction step and its one missing lemma (T97, T98)

T97 (`relay_induction_probe.py`, 105/105) gave the induction step proof shape. Alignment is true and sharp with a constant law; the  $t = 0$  killer loss is proved by three structure identities (old atoms vanish identically on  $E_-$ ; the pole rescue flips with factor  $1 - \cosh(u/2)$ ;  $|\hat{f}|^2 = 2(1 - \cos(tu))|\hat{g}|^2$ ), giving  $k_{\text{eff}} = (1 - \cos(tu))k(t)$  with infimum  $-1.11 \dots -2.59$  instead of  $k(0) = -5.37$ .<sup>2</sup> The structural pearl:  $E_0$  is *literally* the same form at a smaller window ( $7 \times 10^{-14}$ ) — the induction hypothesis, by self-similarity. What remained was a single cross-block lemma about a single vector.

<sup>2</sup>In closed form (a T92 side result)  $k(0) = -\gamma - 3\log 2 - \pi/2 - \log \pi$ . Prior-art note (literature pass, 2026-07-27): in substance this is the classical digamma value  $\psi(1/4) = -\pi/2 - 3\log 2 - \gamma$ , shifted by  $-\log \pi$ , as it appears in Suzuki 2023 (J. London Math. Soc. 107) — a good consistency check, not a find.

T98 (`cross_lemma_probe.py`, 44/44) attacked it and produced the most instructive negative result of the series: the measured  $\sqrt{\lambda_0}$  law is Douglas (1966) range inclusion, i.e. a *tautology* of the positivity one is trying to prove — the identification *dissolves* the lemma rather than proving it. Three T97 premises were refuted in the same run. The replacement target is an exact scalar inequality with no constants and no near-null vector:

$$D_k(\alpha) := -\lambda_{\min}(Q|_{E_-} - Q_{-0}(Q|_{E_0})^{-1}Q_{0-}) \leq \frac{\mu_k}{2}.$$

T99 (`dk_recursion_probe.py`, 23/23) then found the a-priori structure: an exact mirror-parity selection rule ( $J_-Q_{-0}J_0 = -Q_{-0}$  at  $2 \times 10^{-15}$ ) giving  $D_k = \max(D_{\text{even}}, D_{\text{odd}})$ , with the binding even channel coupling only to the *odd* spectrum of the smaller window — so the fragile near-null mode is identically absent from the binding channel. T100 (`bulk_remainder_probe.py`, 27/27) closed the bulk remainder from 11/24 to 24/24 samples and showed the 10.8% drift was a grid artefact.

### 6.3 The race (T101)

T101 (`k_asymptotics_probe.py`, 31/31, CROWDING-TRENDS) resolved 16 zones ( $n = 2 \dots 29$ ) and separated two questions that had been fused. The *mathematics* trends self-supportingly: the primitive quantities are flat over 16 zones ( $w/g$  slope  $+0.002 \pm 0.178$ ; the relative  $D_k$  margin does not trend;  $D_k \leq \mu_k/2$  holds 64/64 and never fails). Atom *strengths* do not fall ( $\mu \sim n^{-0.087}$ ), only the *spacings* do ( $n^{-0.615}$ , exact) — pure arithmetic. What loses the race is the *closure instrument*: the demand ratio  $r_k$  decays like  $\exp(-0.1622k)$  (fit,  $\pm 0.0562$ ). The collapsed one-parameter law  $w_k = 0.0838 \cdot (\text{atom spacing})/\mu_k$  reduced the scatter from  $8.85\times$  to  $2.66\times$  with zero residual trend.




Figure 3: *Schematic*. The instrument race over the 16 resolved zones. *Literal*: the two decay slopes ( $-0.1622 \pm 0.0562$  measured in T101;  $-0.0748 \pm 0.0116$  measured in T103, a factor 2.2 flatter), and the T103 endpoints  $r_1 = 9.33 \rightarrow r_{16} = 2.70$  with the statement that the improved instrument never leaves the spectrum. *Schematic*: the straight-line interpolation in  $\log_{10}$ , the common anchor of the two lines, and the position of the “spectrum floor”. Both slopes are declared fits in the diary. What loses the race is the instrument, not the inequality: the true inequality  $D_k \leq \mu_k/2$  holds 64/64 and its relative margin does not trend.

T101 also stated what an asymptotic theorem would need, and that list has not changed since: (A) the collapsed lower bound  $w_k \geq c g_k/\mu_k$  with a fixed  $c > 0$  — a sharp, non-perturbative, *arithmetic* lower bound at a Weil-positivity edge, with Taylor arguments forbidden by the  $\exp(-c/\delta')$  onset; (B) a uniform relative margin; (C) a replacement for the bulk remainder; (D) a finite check of the low zones. If crowding ever wins, the hardness sits entirely in (A).

## 7 The certification arc (T102–T125)

The last twenty-four parts stopped improving measurements and started converting them into certificates — and, where that failed, into explicit refutations. The arc closes with T125, which

assembles the certificates into one chain and states what that chain is and is not (Section 9).

### 7.1 The mechanism: a crossing, not a singularity (T102)

T102 (`arithmetic_bound_probe.py`, 42/42, MECHANISM-IDENTIFIED) computed the mechanism exactly. The  $k$ -th atom acts on  $E_-/E_0/E_+$  as  $\text{diag}(-\frac{1}{2}, 0, +\frac{1}{2})$ , so  $Q_{k-1} = Q_{\text{full}} - \frac{\mu_k}{2}P_- + \frac{\mu_k}{2}P_+$ : the atom strength  $\mu_k$  enters exactly once, linearly. That gives a two-sided sandwich through the Schur profile  $\sigma_k(\delta)$ , with  $\mu_k/2 \leq \sigma_k$  implying that the handover holds, and  $2w_k$  lying between the crossings in 16/16 zones.

Two corrections came out of the same run. First, the onset is *anchored* at  $\delta_c = 2w_k$  — a crossing of two finite quantities ( $R^2 = 0.968$ ) — and exponential-singular versus power behaviour is not separable; T96's essential-singularity reading is compatible but no longer distinguished. Second, and more consequentially, the binding constraint *flips*: no concentration condition binds. The bare  $E_-$  form is strongly positive ( $2.65 \dots 3.52$ , i.e.  $4\text{--}14 \times \mu_k/2$ ), the pole flip and band mass saturate their classical ceilings (96.9–97.1% of Cauchy–Schwarz and Landau–Pollak/prolate respectively), and the onset is manufactured *entirely* by the Schur dressing against  $E_0 \oplus E_+$ , which eats 35.7%–97.3% of the bare eigenvalue. The loser does not live in  $E_-$ .




Figure 4: *Schematic*. The crossing mechanism of T102. *Literal*:  $\mu_k$  enters exactly once and linearly, so the condition is a two-sided sandwich through the Schur profile  $\sigma_k(\delta)$  against the flat level  $\mu_k/2$ ; the window width  $2w_k$  lies between the crossings in 16/16 zones; the onset is anchored at  $\delta_c = 2w_k$  with  $R^2 = 0.968$ ; the candidate bound  $w_k \geq A_k \mu_k^{1/q_k}$  holds 16/16 with ratio 0.749–0.940. *Schematic*: the curve shapes and all absolute scales.

The same run performed the most useful kill of the arc: a triple decomposition test showed that the T101 law in  $C/g$  was a *proxy*.  $g_k$  is causally impossible (the form  $Q_{k-1}$  is blind to the next atom), statistically dispensable (a causal law  $C \cdot u$  has CV 23.3% against 27.6%), and arithmetically only a ceiling ( $C_k \leq g_k \mu_k$  exactly; the 16-zone extrapolation violates the ceiling from  $k = 69$ ; checked over 18 120 prime-power atoms to  $n = 200\,000$ ).

### 7.2 The instrument, rebuilt (T103), and the two doors

T103 (`instrument_probe.py`, 29/29, INSTRUMENT-IMPROVED) rebuilt the closure instrument with certified weights  $w_j \geq 1/\lambda_j$  and a  $\theta$ -weighted band sum, reproducing both anchors exactly. With *one* fixed,  $k$ -uniform instrument ( $\Lambda_0 = 3$ ,  $r = 2$ , no per-zone tuning) the closure map jumps from 7/64 to 44/64 and the race slope halves (Figure 3). The price is explicit: the number of explicit modes grows  $2 \rightarrow 232$ . And the verdict relocated the remainder: it is no longer the bulk (the bulk is measurably not low-rank — effective rank up to  $0.579 \dim E_-$ ) but the *wing slack*  $S = 1 - \rho$ , which falls  $0.2091 \rightarrow 0.0392$ .

The convergence recorded at the end of T102 is the pivot of the whole arc: T103's wing slack  $S = 1 - \rho$  and T102's Schur dressing are *the same object seen from two sides*. Both doors lead to the wing near-null direction.




Figure 5: *Schematic*. The closure map. *Literal*: the grid is 16 zones  $\times$  4 wing fractions = 64 cells; the old instrument closes 7 of them, the rebuilt one closes 44 with a single fixed  $k$ -uniform setting ( $\Lambda_0 = 3, r = 2$ ) and no per-zone tuning. *Schematic*: which cells are shaded — the diary records the totals, not the pattern; the shading is drawn as a monotone wedge purely to visualise the counts.




Figure 6: Two doors, one object. Both entries are literal diary findings; the convergence statement (“the same object from two sides”) is quoted from the T102 closing note.

### 7.3 Inductive chains, and the death of the margin route (T104)

T104 ran as *two independent arms* after a file collision between two parallel workers was resolved by keeping both implementations of the same contract — `schur_profile_bound_probe.py` (21/21) and `schur_profile_chain_probe.py` (47/47). All core findings converged, which makes it a de-facto independent replication.

- **The naive margin route is dead:**  $M \geq m \cdot \text{Id}$  holds in 0/16 zones; the coupling mass to margin ratio is  $10^3$ – $10^6$ ;  $m$  vanishes in the continuum like  $M^{-1.7}$  (fit).
- **The mechanism is avoidance:** the coupling  $B_-$  decouples from the soft bulk modes — the softest mode carries 0.00% of the coupling mass for  $n \geq 3$ , and the mass sits in mid/high bands (lowest band only 2–12%). Positivity is not saved by a margin but by structural avoidance.
- **Exact spectral-split chains close 16/16 with finite data:** arm A via the exact block-inverse identity with matrix caps in PSD order (headroom 1.15–1.44, reach  $\gamma_{\max} = 0.50$ – $0.90$ ); arm B via a threshold split with an  $O(1)$ , resolution-stable threshold  $w = 2.47$ – $5.90$  — but requiring  $r_{\min} = 64 \dots 1024$  soft modes  $\propto M$ , which makes the continuum statement a spectral-density

condition.

- A recommendation from T96/T103 was tested and *rejected*: the prolate wing basis loses against the raw basis in both arms.

Consequence: the hard core migrated from  $\sigma_k$  as a whole to four named ingredients — a lower bound on  $\text{bare}_k = \lambda_{\min}(Q_{\text{full}}|_{E_-})$ , the soft dressing scalar  $L$ , finite induction data, and a one-sided edge estimate instead of a fit. The chain around them is classical (Schur complement, exact block inverse, Parseval/Bessel, Weyl) and fit-free.

#### 7.4 The certified bare bound and the avoidance theorem (T105)

T105 (`bare_wing_bound_probe.py`, 28/28, ONE-OF-TWO) delivered one of the two missing ingredients in closed form. First it resolved the T104 arm discrepancy: *bare* depends only on  $(u_k, \delta)$  — the window centre cancels ( $\leq 0.74\%$  drift under  $8\times$  refinement) — with the exact split  $\text{bare} = \mu_k/2 + b_0$ . Then three classical steps (Bessel; a Legendre/level-set optimisation at  $k(\pi/\delta)$ ; Cauchy–Schwarz at the pole pair) collapse to a closed lower bound:

$$b_0 \geq \frac{\delta}{\pi} \int_0^{\pi/\delta} k(t) dt - \delta K_{\max} - 4 \left( \cosh \frac{u}{2} - 1 \right) \sinh \frac{\delta}{2}.$$

This is the band average of the archimedean kernel over the band conjugate to the wing, of size  $\sim \log(1/2\delta) - 1$ . It is positive in 16/16 zones (1.27–3.54) at 81.1–92.7% of the measured value (median 92%), it sharpens as the wing flattens, and it needs no eigenvalue or induction data.

Second, the *avoidance law* became a theorem rather than an observation:  $Q_{\text{full}} \geq 0 \Rightarrow \|B^T v_i\|^2 \leq w_i \Lambda_-$  (a semi-inner-product argument; coupling proportional to the mode eigenvalue), verified on every mode of every zone with worst ratio 0.377. That explains the measured 0.00% avoidance exactly. Alongside it, an exact parity superselection:  $JQJ = Q$  ( $5.7 \times 10^{-15}$ ) and  $JB = -BR$  ( $3.4 \times 10^{-15}$ ) — the two channels do not mix, and the softest mode is  $J$ -even in 16/16 zones. Notably the *Euclidean* angles are not small (cosine up to 1.0): the avoidance lives in the form metric, as a Friedrichs angle.

What T105 left is one Loewner statement:  $Q_{\text{full}}(\alpha) \succeq (\mu_k/2) P_-^{(k)} \iff \rho \leq 1 - (\mu_k/2)/\text{bare}_k$  — a Friedrichs-angle / spectral-density condition, explicitly *not* a margin condition. The free cap  $L \leq \Lambda_- = 5.63\text{--}7.52$  diverges like  $\log(1/D)$ , so the certificate has to come from the spectral density near the soft end.

#### 7.5 Parity superselection: the pole splits (T106)

T106 (`friedrichs_angle_probe.py`, 32/32, DENSITY-MAPPED) killed two routes and found the structural result of the arc. The kills are in Section 8. The finding:

$$ab^T + ba^T = ss^T - tt^T \quad (10^{-18}; Ja = b)$$

Here  $J$  is the mirror parity of Section 3, arriving for the second time and now as the obstruction itself. The rank-2 Weil pole splits *exactly* into a positive rank-one lift in the  $J$ -even channel and a negative rank-one pressure in the  $J$ -odd channel. The Toeplitz part has exactly one negative eigendirection; it is  $J$ -even, and  $ss^T$  removes it precisely there. The odd channel never needed the pole at all.

The channel-wise angle chain closes 16/16 in the even channel and 9/16 in the odd channel, against only 3/16 unsplit; all 16 handovers close at operating depth in both channels under the exact Schur criterion, with certified  $b_0$  at 89–93%. Certified after T106: bare, the parity split, the pole splitting, the channel inertia, the avoidance law, the growth inequality, the accumulation bounds. Measured:  $\rho_-$ .




Figure 7: Parity superselection (T106). All numbers are literal; the layout is illustrative. The practical consequence: it is not the density that decides the angle but the *avoidance* — the dangerous channel is the one with the better density. T105’s certified bare bound sits 16/16 exactly in the hard channel, and  $\rho_-$  is resolution-stable ( $\leq 17.6\%$ ).

## 7.6 One scalar ratio (T107)

T107 (`odd_channel_closure_probe.py`, 30/30, SCALAR-TRACTABLE) reduced the matrix statement (R) to a scalar one. Sherman–Morrison gives exactly (R)  $\iff s^* \leq 1$  (identity at  $3.4 \times 10^{-11}$ ; equivalence confirmed 16/16 against Rayleigh–Ritz; precondition  $G \succ 0$  via Cholesky 16/16) — but  $s^*$  is *saturated* (0.99885...0.99999, with only  $n = 2$  at 0.2569), so no uniform  $s^*$  chain can work. The Woodbury decomposition rescues it:  $s^* = \tau + \kappa$  and

$$(R) \iff r := \frac{\kappa}{\varepsilon} \leq 1, \quad \varepsilon := 1 - \tau, \quad \text{measured } r = 0.0053 \dots 0.1814.$$

Two orders of magnitude of headroom instead of five decimal places, and  $s^*$  is resolution-stable (spread  $\leq 2.60\%$  over  $M = 600/900/1200$ ). The synthesis is stated with the usual asymmetry: certified-closed 0/16, measured 16/16 (with  $q = 1.31$ –4.16 at  $M = 1200$ , replacing T106’s 0.747–1.156 — a factor 1.7–2.5 gain). Exactly one object is missing: a certified positive lower bound for  $\varepsilon = 1 - \tilde{t}^T T_{\text{odd}}^{-1} \tilde{t}$  — how far the pole vector is from exhausting the odd channel — or, since  $\varepsilon$  and  $\kappa$  both fall with  $M$  ( $\varepsilon \sim M^{-1.99 \dots -1.69}$ , fit), a certificate for the *ratio*  $r$  itself. The named route is via a wing statement of the T105 type plus an avoidance norm, explicitly not via Szegő.

## 7.7 The $\varepsilon$ identity, and two scalars (T108)

T108 (`ratio_certificate_probe.py`, 44/44, EPSILON-IDENTITY) turned the missing positivity of  $\varepsilon$  from a bound to be won into an *identity*. Everything hangs on the pole response  $x := T_{\text{odd}}^{-1} \tilde{t}$ , and three exact identities (verified at  $10^{-11}$ – $10^{-13}$ ) pin it down:  $\varepsilon = x^T Q|_{\text{odd}} x / \tau$  — so  $\varepsilon$  is a  $Q|_{\text{odd}}$ -energy, not an accident of subtraction —  $1/\varepsilon = 1 + \tilde{t}^T Q|_{\text{odd}}^{-1} \tilde{t}$ , and the overlap of  $\varepsilon$  and  $\kappa$  is exactly one rank-one term  $hh^T/\varepsilon$ , so that (R)  $\iff \hat{r} = (\mu/2) \lambda_{\max}(\Gamma + hh^T/\varepsilon) \leq 1$ . The third route supplies the classical reading:  $\tilde{t}$  is exactly two-term geometry,  $\tau$  a  $2 \times 2$  Christoffel–Darboux form, and  $\varepsilon$  exactly the square of the last Cholesky pivot — the Szegő/Levinson prediction error. *The positivity of  $\varepsilon$  coincides with the induction positivity in the odd channel*; it is not a second, independent thing to certify.

Two consequences, and one correction of the previous picture. The consequence: (R) reduces from an  $M/2$ -dimensional Loewner statement to **two scalars** —  $\omega < 1$  (measured 0.2366–0.7531, 16/16; the T105 structure survives the parity split verbatim,  $V^T S|_{\text{odd}} V = -\frac{1}{2}I$  at  $4.9 \times 10^{-14}$ ) and one statement at the vector  $x$ . The chain (R)  $\Leftarrow [\omega < 1] \wedge [\lambda_{\min}(Q|_{\text{odd}}) \geq (\mu/2)\tau\theta/(1-\omega)]$  has C1 closing 16/16 (margin 5.5–178;  $\geq 1$  at all 48 ladder points, 2.7–343), C2 15/16, C3 0/16 — the

first time a margin chain in this arc is *not vacuous*: the induction margin it needs is explicitly  $9.3 \times 10^{-8}$ – $3.7 \times 10^{-3}$ , i.e. two to six orders of magnitude *below*  $\mu/2$ . The correction: the soft-edge picture of T107 was wrong in its mechanism —  $\tau$  sits to 99.66% on modes  $\lambda \geq 1$ , and  $\|h\|^2$  is small by sign cancellation ( $\sum |k| = 401$  versus  $\sum k = 1$ ), not by softness. T107's Cauchy–Schwarz chain turns out to be exactly the Weyl bound (slack 1.0000).

What is exactly open after T108 is one boundary value. The remaining object is a certified upper bound for the avoidance norm  $\|V^\top T_{\text{odd}}^{-1} \tilde{t}\|^2$ , and on 8 zones that is literally the single edge entry  $x_0$  of an explicit vector (suppressed at  $|x_0|/\max |x| = 2.9 \times 10^{-3}$ – $5.7 \times 10^{-3}$ , with a linear edge profile  $r^{1.01}$  quoted as a fit) — a boundary-decay statement about an explicit vector, the same object class as the T105 carrier separation.  $\hat{r}$  is flat in  $M$  ( $6.1\times$  more stable than  $\varepsilon$ ) and  $r, \hat{r} < 1$  at all 92 (zone,  $M$ ) points up to  $M = 3000$ , tightest at 0.9648: measured, on resolved zones, as always. That boundary value, and the still-uncertified  $\omega$ , are exactly what the next part attacked.

## 7.8 Boundary certificates: cancellation, not decay (T109)

T109 (`boundary_decay_probe.py`, 29/29, BOUNDARY-CERTIFIED) closed the arc: both remaining scalars are now certified, and the whole chain for (R) closes 16/16 conditional on exactly one explicit input. Three findings and one honest remainder.

First, the mechanism. The edge value  $x_0$  is not small because anything *decays* — it is small because a sum *cancels*. Carrier separation is excluded geometrically:  $\tilde{t}$  has its maximum at  $r = 0$ ; the source sits on the boundary itself. The Combes–Thomas hypothesis is satisfied ( $T_{\text{odd}}$  is of Jaffard class, with local rate  $\rho(s) = D(\frac{1}{2} + 2/(e^{2s} - 1))$ , matched at a factor 0.85–1.00, quoted as a fit) — but its conclusion is empty: the Green sum for  $x_0$  cancels by a factor 44– $4.6 \times 10^4$  on 15/16 zones (the known outlier  $n = 2$  does not cancel), and the edge profile is a nearly linear zero,  $r^{0.77\dots 1.28}$  as a fit. A boundary condition, not a decay length — and no absolute-value bound can survive a cancellation.

Second, the  $x_0$  certificate that works. Four candidates were run certified-against-measured: Combes–Thomas, evaluated with the *exact* Green row so that no constant could rescue it, is *refuted* (1/16; an upper bound on its whole class); a  $T$ -metric Cauchy–Schwarz at the edge entry closes 1/16; the exact Szegő/Levinson two-point evaluation exhibits the cancellation but bounds nothing. The fourth is the new object: a residual / goal-oriented certificate (Prager–Synge; second order in the two residuals) which *carries* the cancellation by construction instead of bounding it away. Sharpness 1.0000–2.4937 against the measured norm, 16/16 within budget.

Third,  $\omega$  is cracked — unconditionally. The certificate is a graded matrix cap in PSD order: one trial level per soft direction, validity by a single Cholesky factorisation (Sylvester inertia) plus Loewner antitonicity of the inverse,  $n_{\text{top}} = 17$ –512. It gives  $\omega_{\text{cert}} = 0.2561$ – $0.7569$  against the measured 0.2366–0.7531 — a factor 1.004–1.124 — with *no* hypothesis input. The compression Schur distance that blocked T108 shrinks from 0.18–0.52 to 0.0024–0.039, and the ungraded cap (one level for the whole soft block, which is what T108's global route amounted to) is vacuous everywhere: the grading is the whole difference. The trial data are numerically generated; their *validity* is trial-independent (identity + Cholesky + Loewner), only their sharpness is not.

The synthesis feeds both certificates from one object ( $\Gamma_{\text{cert}} = S_{\text{cert}}^{-1}$ ): the chain closes 16/16 with margin 1.3–10.4 and is  $M$ -stable (44/48 ladder points to  $M = 3000$ ; the certification price 1.5–28.6 does not grow). What remains open is exactly **one input**: the *strict* positivity margin of the odd channel,  $\lambda_{\min}(Q|_{\text{odd}}) \geq 1.7 \times 10^{-6}$ – $7.5 \times 10^{-3}$  against a measured  $1.4 \times 10^{-5}$ – $5.1 \times 10^{-2}$  — two to six orders of magnitude *below*  $\mu/2$ , i.e. strictly weaker than the conclusion it buys; the induction hypothesis as it stands supplies only the non-strict  $Q|_{\text{odd}} \succeq 0$ . Declared caveats: floating point is not audited, and  $\lambda_{\min}$  is read through Rayleigh–Ritz. Whether that one number propagates itself through the induction is exactly what the next part measured.




Figure 8: The reduction cascade, with its arrow types made explicit. The top line is an *equivalence typing* produced by an exact identity (Section 5); the arrows below are the architecture of a would-be sufficient route, verified on finitely many windows only. Since T109 the cascade ends at one strict-positivity margin input; T110 shows that this input propagates itself along the measured ladder (certified base case, graded Loewner minorant), leaving three gaps — which T111 then drives deep and measures as three separate walls (the two subsections below). The dashed box is the reason this figure is not a proof sketch.

## 7.9 The margin propagates: the circle closes on the measured zones (T110)

T110 (`margin_propagation_probe.py`, contract `MARGIN.PROPROPAGATION`, 28/28, `MARGIN-PROPAGATES` — with a precision statement) asked the only question T109 left open: does the strict scalar margin  $m_k = \lambda_{\min}(Q|_{\text{odd}})$  propagate itself along the zone ladder, against the explicit demand  $\text{need109}(k)$  of the T109 chain? Three findings close the circle on the measured zones; three gaps state precisely what a theorem would still need.

*The margin map.* On a common grid (16 zones,  $\delta_0 = 0.00284$ ) the margin falls monotonically,  $6.95 \times 10^{-2}$  ( $n = 2$ ) to  $1.66 \times 10^{-5}$  ( $n = 29$ ), against  $\text{need109} = 3.09 \times 10^{-7}$ – $3.87 \times 10^{-4}$ :  $m_k \geq \text{need109}$  on 16/16 zones, with ratio 2.09–179.6. Dilution during pure window growth is weak — a factor 1.02–2.05 per atom gap, i.e. 0.993–0.997 per cell, a  $1 - O(1/M)$  loss; the scaling fits are honestly poor (the power law wins 7/16). The structural finding of the part is the atom entry: it costs *nothing*. The naive expectation was a  $-(\mu/2)$  pressure on the odd channel; measured, the entry *raises*  $\lambda_{\min}$  on 15/15 handovers (the first-order Kato term has helping sign), and at  $n = 29$  the atom-free window is not even positive ( $-9.2 \times 10^{-3}$ ): the atom makes the window co-positive. The T106 killer argument — transport destroys a demand that *points* somewhere — has no analogue for the scalar floor.

*The step law.* The growth step splits exactly (embedding error  $\leq 2.6 \times 10^{-14}$ ), and  $\max |\mu N| = 0.0$  on 15/15: the new atom lies outside the old lag reach (gap 0.0606 against  $\delta_0 = 0.00284$ ), so its restriction to the old window is the exact zero matrix — everything is decided by the bordering. Scalar routes fail decisively: bordered Weyl closes 0/15 and the Schur/Friedrichs scalar cap is vacuous 15/15 (both recorded in the kill list). What passes certified is the **graded Loewner minorant** — its validity condition is exactly the induction hypothesis, and the level is then found by bisection through a Cholesky/Sylvester factorisation: with one soft direction ( $n_{\text{soft}} = 1$ ) it holds 15/15 with retention 1.0000; the strictly scalar variant ( $n_{\text{soft}} = \text{dim}$ ) closes 0/15.

*Base case and circle test.* The base case  $n = 2$  is *certified* by an explicit Cholesky of  $Q|_{\text{odd}} - \lambda I$  at three grid resolutions ( $\lambda_{\min} \geq 6.93 \times 10^{-2}$ , a factor 88–180 over need). The circle test then runs the whole chain from that certified base value: the strictly scalar induction breaks at  $k = 3$  (its Schur factor compounds to  $2.7 \times 10^{-36}$ ), but the graded chain runs through completely — 16/16 zones above  $\text{need109}$ , each step one Cholesky, with the propagated floor as the only input. On the measured zones, the induction circle closes end to end.

### The three gaps — the asymptotic mountain, made precise

1. **No reserve.**  $f_{\text{crit}}$  at the first handover is 1.00; the chain lives on a retention of  $1 - 2 \times 10^{-7}$ , and a factor-2 loss anywhere would break it within two steps. What a theorem needs is a step law *without* loss.
2. **No scalar step law.** Structurally excluded, not merely missed: the new cells attach at the boundary layer where the pole source sits, so any one-number summary of the bordering is blind to exactly the direction that decides it.
3. **No uniformity in  $k$ .** Every certificate above is one finite Cholesky; flat-in- $k$  is a measurement, not a theorem. And the measured trend warns rather than reassures:  $m_k \sim n^{-1.93}$  against  $\text{need109} \sim n^{-0.98}$ , ratio  $\sim n^{-0.96}$  (fit, rms 0.735) — extrapolated crossing at  $n \approx 170$ . Propagation gets *harder* with  $k$ , not easier. (T111 subsequently *measured* that crossing: it sits at  $n^* \approx 462$ , not 170 — the extrapolation was a fit artefact of the 16-zone window, but the wall itself is real. See the next subsection.)

## 7.10 The three walls (T111)

T111 (`deep_zone_stress_probe.py`, contract `DEEP.ZONE.STRESS`, 23/23, `CROSSING-CONFIRMED`) stopped extrapolating and measured. On exactly the T110 cell grid — the first 16 zones reproduce T110 bit-exactly — the ladder was extended to *every* prime-power zone  $n \leq 1000$  plus a thin tail to 3001: 199 zones in the margin map, with the chain ladder running  $n = 2 \dots 521$  through 117 handovers.

*The margin wall, measured.* The ratio  $m_k/\text{need109}$  falls  $179.59 \rightarrow 0.0174$  across the deep ladder, and the crossing is now a measurement, not a fit: the last zone with ratio  $\geq 1$  is  $n = 461$ , the first below is  $n = 463$  (ratio 0.935). The  $n \approx 170$  figure of T110 was a fit artefact of the 16-zone window — the wall sits much further out, but it is real. The failure follows the *prime branch*: prime-power zones (512, 529, 625, 729), whose atoms are weaker, survive longer than the primes around them. The fit families were again compared honestly (best by AIC: a broken-scale law with  $n^* = 501$ ; the jackknife band of the pure power law is [298, 352]), and a depth-preserving resolution study shows that refinement pushes the ratio *down*:  $n^* \approx 462$  is an *upper* bound, not an estimate to be rounded up.

*The ladder wall, arithmetic.* The uniformity object is no longer flat:  $n_{\text{soft}}^*$  runs  $1 \dots 8$  across the 117 handovers. But retention stays 1.000000 at every step (largest single loss  $5.96 \times 10^{-8}$ ), and the atom entry stays constructively free on 117/117 — the restriction of the new atom to the old window is the exact zero matrix, deep into the ladder. The new and harder finding is the *ladder wall*: the nested step needs a lag gap exceeding twice the window depth, and the twin-prime pair  $521 \rightarrow 523$  (gap  $0.003831 < 2D = 0.005617$ ) kills the ladder *arithmetically* — at frozen cell width,  $n = 521$  is the deepest reachable zone, and the bound is set by twin primes, not by spectra. An arithmetic scan to 40 000 finds 4052/4284 consecutive prime-power pairs endangered; letting  $D$  shrink like  $1/n$  instead forces the cell count to grow like  $n \log n/4$  (horizon  $n \approx 690$  at  $h \leq 1500$ ).

*The circle tears at the ratio, not at a step.* The no-reserve finding of T110 softens:  $f_{\text{crit}} \sim n^{-0.39}$  (fit) — the reserve *opens* with depth, and the bottleneck stays at the first step  $2 \rightarrow 3$ . The circle test runs from the certified base case (factor 179.6) and breaks first at  $k = 109$ ,  $n = 463$  — exactly where the margin map measures the crossing. All 117 handovers certify at retention 1.000000, including beyond the break: what fails is the demand ratio, never the handover mechanism.

#### The three walls — and the reinterpretation

The hardness balance names the driver: the growth of  $1/\kappa$  (exponent +0.86) against the margin (exponent −0.68). Three separate walls, in order of appearance:

1. **The margin wall**,  $n^* \approx 462$ : measured, an upper bound (refinement pushes it down); the T109-chain demand ratio crosses 1 between  $n = 461$  and  $n = 463$ .
2. **The ladder wall**,  $521 \rightarrow 523$ : purely arithmetic — at frozen cell width the nested handover dies at the first twin-prime gap below  $2D$ , independently of any spectral quantity.
3. **The requirement wall**,  $n = 727$ :  $\omega_{\text{cert}} \geq 1$  makes need109 vacuous on 46 deep zones — beyond it the T109 chain no longer states a demand at all.

The reinterpretation that comes with the measurement: the operating variable is the window *depth*, not the zone index  $n$  — doubling the cheeks moves the margin wall from 463 to 47. A proof along this route must beat the exponent gap  $-0.974 = -0.681 - 0.293$  and couple the depth  $\delta_0$  to the local prime gap.

T112 subsequently ran the construction in exactly that coupled frame (next subsection): walls (2) and (3) fall *structurally*, wall (1) proves frame-invariant.

### 7.11 The scaled frame: two walls fall, the margin wall is frame-invariant (T112)

T112 (`adaptive_scaling_probe.py`, contract ADAPTIVE.SCALING, 20/20, SCALING-PARTIAL) stopped freezing the cell width and coupled it to the local scale. Two couplings were built: frame A, gap-coupled,  $D_k = g_k/(2\nu)$  with  $g_k$  the local prime-power gap and  $\nu$  the resolution constant of the frame; and frame B, mean-field,  $D_k = \log n/(2\nu n)$  — the smooth PNT version of the same idea. The first finding is a structural asymmetry: *frame B has no ladder* (only 38% of consecutive pairs nest; the first failure is already at  $n = 3$ ) — the nested handover lives exclusively in the gap coupling, not in any smooth  $1/n$  law. T105 admissibility, re-measured in the scaled frame, sets a resolution floor  $\nu \geq 4$  (at  $\nu = 1$ , 3/10 test zones are inadmissible). The flatness verdict is honest and double-edged: the ratio is *not* flat, but the exponent is *frame-invariant* — the components shift massively ( $m_k$ :  $n^{-1.93} \rightarrow n^{-0.95}$ ; need109:  $n^{-0.98} \rightarrow n^{+0.03}$ , i.e. practically  $n$ -independent in the scaled frame), yet the difference stays  $-0.974$ .

*The three walls, scaled.* The **ladder wall is structurally gone**: with  $D_k = g_k/(2\nu)$  the nested step has exactly  $\nu$  cells per end *by construction* — twin pairs included — verified as an exact arithmetic lemma over all zones; the atom entry is the exact zero matrix by construction rather than by arithmetic accident; and the T111 killer pair  $461 \rightarrow 463$  now certifies *spectrally* ( $h = 1418$ , retention 1.000000), with the reserve *opening*:  $f_{\text{crit}} = 8.1 \times 10^{-4} - 4.2 \times 10^{-3}$  instead of 1.00 — roughly a factor 1000. (The pair  $521 \rightarrow 523$  would need  $h = 1634$ , above the dimension cap: arithmetically co-certified, spectrally declared honestly as not run.) The  $\omega$  **wall is no longer a depth wall**: the depth coefficient is  $n^{-0.079 \pm 0.121}$ , compatible with zero —  $\omega_{\text{cert}}$  is driven by the resolution and the local gap, not by the depth of the ladder; a residual requirement-vacuum tail of 5 zones on the deepest ladder remains and is not booked as “gone”. The **margin wall stays**, practically unmoved: the first sub-1 zone sits at  $n = 449$  against the frozen frame’s 461–463 (a factor 1.07) — it is not geometry but the substance of the T109 requirement itself.

*The limit object.* Overlaying the scaled windows, the archimedean part (0.657) and the pole part (0.558) contract Cauchy-like with amplitude laws  $D^{-0.26}$  and  $D^{+0.37}$ ; the atom part does *not* contract (0.935) — it is the prime-gap statistics itself, written in scaled coordinates. The spectral shape does not drift:  $\lambda_2/\lambda_1 \sim n^{-0.029 \pm 0.052}$ . The resulting formulation is a split object: *[a deterministic limit shape: arch + rank-one pole] + [the atoms as a controlled, non-convergent perturbation]*. New and uncomfortable is the **regrid object**: the zone frames are not nested, so Rayleigh–Ritz transfers nothing between them; the regrid cost is measured at rate  $(D'/D)^{+2.93}$  — an object that simply does not exist in the frozen frame, disclosed here because any limit argument has to pay it.

#### The trade — and the disclosed prime-gap dependence

Three arithmetic walls are exchanged for two analytic demands:

- [P1] **Positivity of the limit operator** — *one* operator instead of a sequence of matrices; the deterministic part (arch + rank-one pole) contracts, the atom part is the perturbation to be controlled.
- [P2] **A convergence rate below the step reserve** — the Cauchy rates and the regrid cost against the opening reserve  $f_{\text{crit}}$ .

The prime-gap inputs are disclosed rather than hidden. Upward, Bertrand–Chebyshev 1852 ( $g_k \leq \log 2$ ; used and verified in-run) and the trivial even-gap bound suffice. The only remaining arithmetic parameter is the normalised gap  $\theta_k = g_k n / \log n = 0.18 - 4.12$  — and downward,  $\theta$  is *provably* unboundedly small: by Zhang 2014 / Maynard 2015 (infinitely many bounded gaps),  $\theta \rightarrow 0$  on subsequences, so any gap-coupled construction must be uniform in  $\theta \rightarrow 0$ , at the price of the  $1/\theta$  cost explosion.

T113 subsequently answered the currency question (next subsection): the wall is currency-invariant — and what it measures is the discretisation, not the spectrum.

### 7.12 The currency question: the wall is substance, and it measures the discretisation (T113)

T113 (limit\_operator\_probe.py, contract LIMIT.OPERATOR, 27/27, SUBSTANCE-CONFIRMED) answered the question T112 left open — is the margin fall in the scaled frame a normalisation artefact? — and the answer reinterprets the program’s hardest number.

*The currency table.* The floor/requirement ratio carries the *same* exponent,  $n^{-1.168 \pm 0.259}$ , in all five currencies tested (raw, per  $\lambda_{\text{max}}$ , per trace density, per  $D$ , per  $D^2$ ) — and the invariance is structural, not accidental: the piecewise-constant basis is  $L^2$ -orthonormal (Gram matrix exactly  $I$ , so there is no free  $D$ -power to spend), and the zone matrices are restrictions of *one* quadratic form to subspaces of the same  $L^2$  space (cross-grid form identity at relative  $1.2 \times 10^{-14}$ ). The floor’s homogeneity degree in the cell width is exactly 1.0; the requirement’s is 1.0 under every legitimate rescaling. No flat currency exists. The wall is currency-invariant: substance, not bookkeeping.

*Two surprises about what that substance is.* First, the limit object is *not* “arch + pole with the atoms as a perturbation”: the atom-free window form is deeply negative ( $\lambda_{\text{min}} = -3.08$  to  $-19.36$ , growing like  $n^{+0.60}$ ), the atom block is equally large (norm 3.21–19.44), and the positive

floor survives only as a *cancellation* of relative size  $1.3 \times 10^{-7}$ – $9.7 \times 10^{-5}$  — on the floor vector,  $-11.64$  against  $+11.64$ . The Weyl bound is saturated: norm perturbation theory is five orders of magnitude too coarse for this object. Second, there is *no plateau under refinement*: at a fixed window, on nested grids over 8 levels (a factor 128 in resolution),  $\lambda_1 \sim D^{1.83}$  and  $\lambda_2 \sim D^{1.76}$  — the *same* power, with  $\lambda_2/\lambda_1$  flat. **The continuum window form has no gap;  $m_k$  measures the gap of the discretisation.** This resolves the T112 tension exactly:  $\lambda_2/\lambda_1$  constant while the floor falls is what it looks like when both eigenvalues carry the same  $D$ -power.

*The atom part is almost invisible at the floor.* The entering (deepest) atom contributes almost nothing to the coupling ( $8.3 \times 10^{-11}$ – $1.3 \times 10^{-6}$ ) — which is why handover steps certify at retention 1.000000 while the margin tears. The floor is almost pure geometry:  $\log m = \text{const} + 1.87 \log D - 2.07 \log \alpha$ , with an arithmetic residual of a factor 1.13 and no  $\Lambda(n)$  dependence. The  $\theta \rightarrow 0$  stress test separates cost from substance: the operator norm does *not* diverge ( $\theta^{-0.005}$ ); an extra atom at the window edge shifts the floor by only 0.2–0.4%, while the same atom in the interior shifts it  $10^7$  times more — the  $1/\theta$  explosion of T112 is purely a *cost* statement ( $h \sim \nu n/\theta$  cells), not an operator instability.

*Regrid, re-read.* The raw regrid rate  $(D'/D)^{3.85}$  is mostly the floor's own  $D$ -power — a normalisation jump, not an instability; the reserve certifies  $f_{\text{crit}} = 9.5 \times 10^{-2}$  (better than the T112 estimate), and after subtracting the  $D$ -power the reserve buys 15 steps (against 0 on the raw jump).

#### The new hardness balance — replacing [P1]/[P2]

The T112 trade is restated by measurement; this balance supersedes the earlier two-demand form.

**[P1] Semidefiniteness of the balanced form without a gap** — the continuum window form is semidefinite at best; strict limit-floor formulations demand what refinement demonstrably refuses.

**[P2] Certify the  $D$ -power itself** — a Grenander–Szegő-type consistency statement for  $\lambda_{\min}$  of the scaled form; the power, not an abstract rate, is the object.

**[P3] New, and now leading: the requirement homogeneity** — need109 carries  $D^{0.63}$  against the floor's  $D^{2.38}$ : the T109 chain divides by an *artifact* margin. The repair the measurements point at is a *margin-free step certificate* — a Schur/Loewner formulation certified from semidefiniteness alone, never dividing by a margin.

**[P4]  $\theta \rightarrow 0$  is a cost-uniformity demand only** — the frame needs  $h \sim \nu n/\theta$  cells, but the operator itself is  $\theta$ -stable.

*Disposition after T114 (next subsection):* [P3] is **discharged** — the margin-free step certificate exists and runs past the wall; [P4] stands as a pure cost statement; a new **[P5]** appears (the (R) demand is grid-bound and is deleted, not transported); and [P1]/[P2] are restated as the two-object core —  $\varepsilon$  relative to  $\kappa$ , and Schur transport between grids.

T114 (`margin_free_step_probe.py`, contract `MARGIN.FREE.STEP`) then rebuilt the chain without the margin division — and the wall dissolved.

### 7.13 The wall dissolves: the margin-free step (T114)

T114 (`margin_free_step_probe.py`, contract `MARGIN.FREE.STEP`, 22/22, WALL-DISSOLVES) built the certificate that [P3] demanded, and the answer to the preregistered question — does the margin wall disappear as a requirement artefact? — is yes.

*The margin-free step.* The structural fact that makes it possible was already on the table (T110/T111): the restriction of the grown window's form to the old window is bit-exactly  $Q_{\text{old}}$  — the entering atom's block is the exact zero matrix — so margin-freeness *is* that property, once the step stops asking for more. Route (i) is Albert 1969: the generalised Schur criterion with the Moore–Penrose inverse,

$$Q' = \begin{pmatrix} A & C \\ C^\top & X \end{pmatrix} \succeq 0 \iff X \succeq 0, \quad \text{ran}(C^\top) \subseteq \text{ran}(X), \quad A - CX^+C^\top \succeq 0$$

(the range condition is Douglas 1966) — no margin appears anywhere. It certifies 27/27 ladder

steps, *eleven of them beyond the old wall* (467, 479, 512, 529, 547, 577, 661, 773, 887, 1129, 1331; deepest  $h = 1495$ ). The exact  $4 \times 4$  Schur complement is an  $O(0.1)$  object with *no* cancellation:  $\lambda_{\min}(S) = 0.068\text{--}0.154$ , i.e. 42–67% of the block scale. Route (ii), the T110 graded minorant in its degenerate limit  $x_{\text{in}} = 0$ , dies structurally on 27/27 — which retypes T110’s margin: it was not paying for positivity but for an *abandoned direction*. Route (iii), unshifted Cholesky, passes 27/27 against the declared backward-error floor. Negative controls are sharp (5/5), and the certificate is frame-invariant at  $\nu = 8$  (3/3). All seven zones at which the T109 chain tears — including the wall zone  $n = 449$  itself, where  $\text{need109}/m = 1.241$  — are certified by the margin-free step.

*The wall in one line.* The same quantity, bounded through norms (Weyl:  $\lambda_{\min}(A) - \|C\|^2/\lambda_{\min}(X)$ ), is *negative* by a factor of  $2.4 \times 10^5\text{--}9.6 \times 10^7$ : an  $O(1)$  numerator divided by the  $10^{-6}$  artifact floor. Every norm route had to fail — and the exact Schur complement passes through none of it.

*The favourable power direction.* On nested refinement chains at a fixed window,  $\varepsilon \sim D^{1.77\text{--}2.01}$  and  $\kappa \sim D^{3.64\text{--}4.84}$  (fits) — the preregistered “same power” expectation is refuted *in the favourable direction*:  $\kappa$  falls faster, so the scale-free ratio  $r = \kappa/\varepsilon \sim D^{1.6\text{--}3.1}$  shrinks under refinement. The ratio closure  $r = 2 \times 10^{-4}\text{--}0.103$  holds on 27/27 steps (11/11 beyond 463). The circle [base semidefiniteness (a declared hypothesis input) + margin-free step + ratio closure] closes 27/27; ten multi-step chains run (21 certified steps, the longest of length 4); and *every* chain ends at the cap  $h \leq 1500$  — a window *cost* — never at a step. Regrid honesty: 0/464 gap ratios are dyadic, so the non-nested transport object of T112/T113 remains real.

[P5]: *the (R) demand is grid-bound.* Holding the physical demand  $(\mu/2)P_-$  fixed and refining the grid,  $\omega$  grows monotonically (5/5) and crosses 1 (4/5): the (R) formulation is *false in the continuum* — it is to be *deleted*, not transported. The exact Schur complement needs no demand surrogate in its place.

*Cancellation robustness.* The  $10^{-7}$ -relative cancellation of T113 has 5.3–8.8 decimal digits of headroom above the Cholesky backward-error floor (Wilkinson/Higham), and it sits almost entirely in the rank-one pole — whose positivity is exactly the T108 Szegő identity  $\varepsilon$ .

#### The new core — two objects

The four-part balance collapses. [P3] is discharged (the margin-free step certificate exists and runs past the wall), [P4] was only ever a cost statement, and [P5] deletes (R). What is left:

[P1]  $\varepsilon > 0$  **with controlled size relative to  $\kappa$**  — the exact Szegő/Levinson identity of T108, now the single positivity object of the program.

[P2] **Transport of the exact Schur complement between non-nested grids** — an  $O(0.1)$ -sized object with no cancellation, sharpened from the earlier abstract  $D$ -power demand.

*Disposition after T115 (next subsection):* [P2] **splits** — the transport across a real regrid certifies only mild refinement ( $\rho \lesssim 1.8$ ), for a principled reason (the Schur floor itself falls under refinement), while a two-scale compression breaks the compute cap; [P1] **sharpens** to one scalar inequality, the classical Szegő–Levinson prediction-error bound. The core is restated as the three-point list at the end of the next subsection — only one point of which is an inequality.

T115 (`schur_transport_probe.py`, contract `SCHUR.TRANSPORT`) then ran the two last bricks — and split them.

### 7.14 Transport blocked, cap broken: the two-scale compression (T115)

T115 (`schur_transport_probe.py`, contract `SCHUR.TRANSPORT`, 26/26, `TRANSPORT-BLOCKED` — with a large compression gain) took the two-object core apart: transport of the exact Schur complement across the ladder’s non-nested grids, and a multi-resolution compression of the already-certified interior.

*The transport splits at  $\rho^* \approx 1.8$  — for a principled reason.* Haynsworth’s partial minimisation  $y^\top S y = \min_z [y; z]^\top Q' [y; z]$  makes  $\lambda_{\min}(S)$  available without any inverse and without a margin; the two-sided transport bracket built from it is an *identity* (bookkeeping), with the evidence

sitting in the consistency defect  $\eta$  and in the lower end. On the eleven real regrid pairs of the ladder ( $\rho = 1.023\text{--}3.425$ ) the lower end is positive on 7/11, with a clean split: every pair with  $\rho \leq 2.061$  certifies and none with  $\rho \geq 2.291$  does; a synthetic scan puts the threshold at  $\rho^* = 1.83$ , against a median real refinement ratio of 2.001 — 39% of the ladder’s 7686 refinement pairs are covered. The blocker is *not* the transport machinery: the exact cell-overlap projection beats three deliberately wrong transfer maps on 9/9 (by factors up to 270), and on *nested* ladders — where the transport error is exactly zero —  $\lambda_{\min}(S)$  itself falls like  $\rho^{-1.73}$  to  $\rho^{-1.65}$ . The drop is real, so no bound can undo it. Two sharpenings came free: non-nestability is a matter of *integrality*, not dyadicity (0/14807 gap ratios are integer, the closest miss  $3.4 \times 10^{-5}$ ), and the measured defect  $\eta = 2.7 \times 10^{-2}\text{--}8.1 \times 10^{-2}$  is  $10^3\text{--}10^6$  times smaller than its a-priori C  a/Strang norm surrogate — which is exactly where the margin division would have re-entered.

*The compression breaks the cap.* The two-scale window (interior merged into blocks, the boundary strip kept fine, the merge anchored at the centre) preserves  $X_{\text{mixed}} = Q_{\text{old,mixed}}$  *exactly* (rel 0.0) — the compressed step is still margin-free — and Albert certifies 66/66 (zone,  $q$ ) combinations up to  $q = 64$ , with the compression error certified *one-sided and second order* in the projection defect (Rayleigh–Ritz plus stationarity of the fine minimiser; C  a/Strang; 66/66) and savings down to  $m/h = 0.043$ . The deep run then breaks the T114 cap: the margin-free step certifies at  $n = 155\,921$  ( $117\times$  T114’s 1331), on a fine lattice of  $h = 93\,470$  cells ( $62\times$  the hard cap  $h \leq 1500$ ) compressed to  $m = 1490$ ; re-coarsening inside a chain is free (Rayleigh–Ritz, rel  $\leq 2.1 \times 10^{-12}$ ); the longest chain is now *ten* steps (T114: four), with 33 certified steps over four chains; the stopper is the cost cap on 3/4 runs and a failing step on 0; and every certified quantity sits  $10^5\text{--}10^{11}$  times above the declared Cholesky backward-error floor. Stated honestly:  $\lambda_{\min}(S_{\text{mixed}}) \approx 5.1$  is  $52\times$  the uniform value because the coarser Galerkin space sees fewer soft directions — a weaker statement about the continuum, declared as such, and flat across the whole deep band. The contiguous reach an induction needs moves from  $n \leq 125$  to  $n \leq 5437$  (factor 43.5): the  $n \log n$  cost barrier is *divided*, not broken.

#### The remaining list — three points, one inequality

The shortest list the program has produced. From T114’s balance, the margin (dissolved), the (R) demand (grid-bound, deleted), the floor/ $\varepsilon$  transport (the object to transport is  $\lambda_{\min}(S)$ , the slowest-falling of the three) and the compression risk (it does not break the step) are all *struck*. What is left:

- (1) **[P1] as one scalar inequality** — given  $T = Q + \tilde{t}\tilde{t}^T \succ 0$ , positivity of  $Q|_{\text{odd}}$  is *equivalent* to  $\varepsilon = 1 - \tilde{t}^T T^{-1} \tilde{t} \geq c(D) > 0$ : the classical Szeg  –Levinson prediction-error bound for one explicit symbol (the Haynsworth double-Albert device:  $T \succ 0 \wedge \varepsilon \geq 0 \iff Q \succeq 0$ ). This is the target-sentence candidate. Measured:  $\varepsilon \sim \rho^{-1.71}$ , falling slightly faster than  $\lambda_{\min}(S) \sim \rho^{-1.69}$  (fits).
- (2) **An a-priori bound on the transport defect  $\eta$**  — a regularity statement about the minimiser  $z^*$ ; it would turn the measured transport into a certificate for  $\rho \leq 1.8$ .
- (3) **A boundary formulation** (suggested by the exact-zero lemma) — never represent the old interior at all; it would remove the compute cap *and* the regrid at once.

Only (1) is an inequality; (2) is a regularity statement, (3) is a reformulation.

*Disposition after T116 (next subsection):* (3) is **run and splits** — the boundary recursion itself is exact (the pole rides in the state with no truncation, and the march cost is flat), but the state cannot be *bounded*: the prime comb refuses compression. (2) is **devalued** — a single boundary march needs no regrid, so the transport defect no longer sits on the critical path. (1) **sharpens** into a two-variable law with a named classical proof technique:  $\varepsilon$  is a Galerkin error, Aubin–Nitsche duality predicts the measured  $D$ -exponent, and the jumps at prime-power entries become an explicit side-condition.

T116 (`boundary_formulation_probe.py`, contract BOUNDARY.FORMULATION) then ran the reformulation — and it held exactly where the contract feared it would break, and broke exactly where it seemed free.

### 7.15 The boundary recursion: exact pole, flat cost, and the prime comb (T116)

T116 (`boundary_formulation_probe.py`, contract `BOUNDARY.FORMULATION`, 33/33, `RICCATI-PARTIAL`) executed point (3) of the list: never represent the old interior at all — run the induction as a pure boundary process.

*The recursion stands, and the pole is free.* The state is the triple  $(\Sigma, r, \sigma)$  obtained by partially minimising the pole-free form over the interior (Haynsworth 1968): the identity  $t^T T^{-1} t = r^T \Sigma^{-1} r + \sigma$  holds on six zones at rel  $6.7 \times 10^{-10}$ , the Haynsworth double-Albert equivalence certifies 6/6 with sharp negative controls, and the Riccati-type chaining of two eliminations is *exact* (synthetic  $1.4 \times 10^{-16}$ ; bit-exact on real windows). The first feared obstruction — pole bookkeeping — dissolves for free: the rank-one pole is genuinely *global* (its boundary value at half depth is 38.6%, against 0.2% for the archimedean part), but the pair  $(r, \sigma)$  carries it *exactly* through bordered elimination, with no truncation anywhere.

*The tail bookkeeping breaks — on the prime comb.* The second feared obstruction was the archimedean kernel tail, and it is innocent: the arch part decays cleanly,  $\exp(-1.73\omega)$ . What does not decay is the *full* symbol — it stays at  $\geq 93\%$  of its maximum beyond every tested width, and 100% of the off-band mass sits in the stripes of the two *prime combs*, carried by the reflection comb: every incoming cell couples through a Hankel term to the interior cell at  $\alpha - \log n_j$ , for *every* prime power  $\sqrt{n} < n_j < n$ . A truncation of the state would have to discard  $1.3 \times 10^2$ – $1.1 \times 10^6$  times the margin  $\varepsilon$  it must respect, a budget that grows like  $n^{2.17}$  ( $\mu_j \sim n^{-1/4}$  against  $\varepsilon \sim n^{-1.8}$ ; fits); 11/18 real truncations break outright; a comb-aware state that keeps the stripes still spans 83.6% of the window and reaches only rel 20; and the ceiling of the whole idea is a  $\Theta(\log^2 n)$  gain. There is no sparse faithful state — the stripes drift with fill-in.

*The deep boundary run — a cost-geometry demonstration, declared as such.* One unbroken chain of 169 236 prepends ran to a window of  $h = 1\,354\,088$  cells —  $903\times$  the old cap — with the entire state  $12 \times 12 + 12 + 1$  numbers and the cost *flat*:  $76\,\mu\text{s}$  per step, first-to-last-decile ratio 1.00. Stated honestly, this is *not* a Weil certificate: at the fixed cell width  $D = 1.9 \times 10^{-6}$  the band contains no atom, the stopper is  $\varepsilon < 0$  at  $h = 1.35 \times 10^6$  — the *band model* loses positivity there, not the machine — and the atom reach stays at  $n = 173$ . What the march certifies is the cost geometry of the boundary formulation.

*The one inequality, sharpened into a law with a proof technique.* The measured law is  $\varepsilon \approx 8.34 D^{1.790} \alpha^{-6.04}$  (jackknife errors  $\leq 0.041$ , rms  $\leq 0.097$ ;  $\varphi$  from a fixed- $D$  sweep,  $\pm 0.08$ ; all fits).  $\varepsilon$  is *not* smooth in  $\alpha$ : it jumps by factors up to 120 at prime-power entries, so any smooth ansatz is wrong from the start. (T117 subsequently resolved the same sweep cell by cell and *corrected* this reading: the entries themselves move  $\varepsilon$  *up*, and the factor-120 falls are the accumulated smooth bordering decay between sweep points — see the T117 subsection. The criticality and  $D$ -law statements of this paragraph are unaffected.) The continuum form is *exactly critical* —  $\tilde{t}^T T^{-1} \tilde{t}$  reaches 0.999975745, and Szegő–Kolmogorov says the pole is representable in the limit — with every measurement  $\geq 1.1 \times 10^5$  above the declared Cholesky floor. The classical hit: in the Galerkin reading,  $\tilde{t}$  represents the  $D$ -independent function  $2 \sinh(x/2)$ ,  $\varepsilon$  is the error of a piecewise-constant Galerkin method for it, and Aubin–Nitsche duality predicts  $\varepsilon \sim D^{2s}$ : the measured  $\theta = 1.79$  gives  $s \approx 0.90$ , exactly the regularity a log singularity plus window edges would produce. The target inequality is now formulated:  $\varepsilon(\alpha, D) \geq c_0 D^\theta \alpha^\varphi$ , uniformly across the jumps.

#### The remaining list after T116 — one inequality, one obstruction

Won: the exact pole (no truncation), the exact Riccati chaining, and flat cost per step. Broken: [SUPPORT] — the prime comb; there is no sparse faithful state, and the stripes drift with fill-in. Devalued: [P2] — a single boundary march needs no regrid, so the transport defect leaves the critical path. The basis is uncritical. What remains on the critical path:

- (1) [P1] as a two-variable law —  $\varepsilon(\alpha, D) \geq c_0 D^\theta \alpha^\varphi$  with  $\theta = 1.790$  and  $\varphi = -6.04$  (fits), to be proved *across* the factor-120 jumps at prime-power entries; the proof technique is named (piecewise-constant Galerkin error for  $2 \sinh(x/2)$ , Aubin–Nitsche duality,  $s \approx 0.90$ ).

**(2) The comb as the honest support obstruction** — any bounded-state formulation must either carry the reflection stripes ( $\Theta(\log^2 n)$  ceiling) or find a representation in which the comb is diagonal.

*Disposition after T117 (next subsection):* (1) is **theorem-shaped** — the Galerkin reading becomes four exact identities, the lower-bound direction acquires a certified two-level chain that loses no power of  $D$ , and the jumps acquire exact closed forms with the *opposite* sign: entries *raise*  $\varepsilon$ , the factor-120 falls were a sweep artefact, and the jump side-condition is *removed* from the ansatz. What is left of (1) is three named analytic lemmas about one symbol, two of them constants. (2) is untouched.

T117 (`epsilon_theorem_probe.py`, contract `EPSILON.THEOREM`) then took the inequality apart — keeping the one honest asymmetry visible everywhere: the classical Galerkin machinery produces *upper* bounds for approximation errors, and the target needs a *lower* bound.

## 7.16 The one inequality becomes theorem-shaped (T117)

T117 (`epsilon_theorem_probe.py`, contract `EPSILON.THEOREM`, 23/23, `THEOREM-SHAPED`) converted the T116 reading into identities and a certified chain. The direction problem is stated once and kept visible: C ea 1964, Aubin 1967/Nitsche 1968 and Bramble–Hilbert 1970 all bound a Galerkin error from *above* by an interpolation error; the missing lower-bound ingredient carries a classical name — the *saturation assumption* (Bank–Smith 1993; D orfler–Nochetto 2002) — and is known to be an assumption, not a theorem. Every classical anchor below is therefore cited with the direction it actually gives.

*The Galerkin identity, exact.* Four identities, verified on six windows to  $10^{-10}$  relative or better, none of them asymptotic: (i)  $\tilde{t}$  is the cell functional of the  $D$ -independent function  $f(x) = 2\sinh(x/2)$  (independent quadrature, rel  $5.2 \times 10^{-16}$ ); (ii) the family is *exactly nested* — with  $P$  the block-averaging prolongation,  $T_c = P^T T_f P$  and  $\tilde{t}_c = P^T \tilde{t}_f$  (rel  $2 \times 10^{-14}$ ) — so  $\varepsilon$  is a Galerkin best-approximation error of *one* bilinear form on a nested family, and its monotonicity under refinement is a *theorem*, not a fit; (iii)  $\varepsilon = \min_v [1 - 2\tilde{t}^T v + v^T T v]$  with the two-level Pythagoras  $\varepsilon_c - \varepsilon_f = \|u_f - Pu_c\|_{T_f}^2$  (rel  $6.2 \times 10^{-11}$ ); (iv) the dual-norm residual form  $\varepsilon_c - \varepsilon_f = \|\tilde{t}_f - T_f Pu_c\|_{T_f^{-1}}^2$ , whence  $\varepsilon > 0 \iff \tilde{t} \notin T(V_c)$ : positivity is a *non-membership* statement — which is exactly why the quantitative version is hard, since a non-membership statement carries no rate by itself.

*The lower bound, honestly.* The crude psd-minorant route ( $\varepsilon \geq 1 - \tilde{t}^T G^{-1} \tilde{t}$  for  $0 < G \leq T$ ) is measured dead: criticality forces any useful minorant to be sharp to relative  $8.3 \times 10^{-4}$  in the one direction  $u = T^{-1} \tilde{t}$  — which is the same as knowing  $u$ . What survives is the *two-level chain*

$$\varepsilon_c \geq \varepsilon_c - \varepsilon_f = y^T S y \geq \lambda_{\min}(S) \|y\|^2 \geq \lambda_{\min}(S) D_f^3 \max_j g_j^2 / 2 > 0,$$

with  $y$  the oscillation of the fine dual solution,  $S$  the  $T$ -Schur complement onto the oscillation space (Haynsworth 1968), and the last step the one-cell Payne–Weinberger 1960 cell-variance inequality used in its *lower* form — the only place where a classical approximation statement points the right way. Every link is an identity or a completed Cholesky: certified on all 19 nested pairs,  $\text{bound}/\varepsilon \in [0.111, 0.185]$ , everything at least  $6.7 \times 10^5$  above the declared backward-error floor. The rate:  $\theta'(\text{sum}) = 1.74$  against  $\theta = 1.76$  (jackknifed fits) — a loss of 0.04 powers of  $D$ , i.e. *none*, because  $\lambda_{\min}(S)$  *grows* under refinement (+0.20 per halving of  $D$ ; log-versus-small-power undecidable on an  $8\times$  lever, and nothing in the chain rests on deciding it): the chain costs a constant, not a power ( $\text{bound}/\varepsilon$  spread 1.195 over the whole lever). The one-cell version costs a further 0.93 powers — the classical price of replacing a sum over cells by its largest term.

*The jumps, exactly — and a correction to T116.* Growing the window by one cell per end at fixed  $D$  is a pure prepend in frame  $A$ , so the bordered inverse gives the *exact* drop  $\varepsilon(h+1) = \varepsilon(h) - r_0^2/s_0$  (Levinson 1947; rel  $3.7 \times 10^{-12}$  over consecutive steps). A prime-power entry at fixed window is a *corner* update supported on the two outermost Hankel anti-diagonals, of rank exactly  $k_0 + 1$  ( $k_0$  the corner distance), and Woodbury/Sherman–Morrison reproduces its effect on  $\varepsilon$  in closed form on all 23 measured entries (rel  $2.1 \times 10^{-11}$ ). The direction *corrects* T116: all 23 entries move  $\varepsilon$  *up* (factors 1.0003–1.0506; for  $k_0 = 0$  the update is  $+\beta e_0 e_0^T$  with  $\beta > 0$ , so an entry can never by

itself lower  $\varepsilon$  at that corner distance). The T116 factor-120 falls were a *sweep artefact*: resolved cell by cell over 400 steps, the largest single-cell drop is a factor 0.9555, the entry share of the total log-drop is  $-0.3\%$  *with the opposite sign*, and the falls are the accumulated smooth Levinson product between sweep points spaced 150–300 cells apart. The jump side-condition disappears from the ansatz; what a bound must survive is a smooth  $\alpha^{-6.07 \pm 0.03}$  decay (fit; the  $\alpha$  lever is short and bound by the compute cap  $h \leq 1400$ , declared as such).

#### The theorem candidate — and the three-lemma missing list

*Candidate.* Let  $V_c \subset V_f$  be the reflection-odd piecewise-constant spaces on  $(-\alpha, \alpha)$  with  $M$  and  $2M$  cells,  $T$  the pole-free odd Weil form,  $\tilde{t}$  the pole functional of  $f(x) = 2 \sinh(x/2)$ . Assume (H1)  $T_f \succ 0$  [certified per window]; (H2)  $\lambda_{\min}(S) \geq \kappa > 0$  on the oscillation space [certified per window; measured to grow under refinement]; (H3) one coarse cell carries a non-vanishing discrete slope of the dual solution [measured]. Then the four-line chain above gives  $\varepsilon_c \geq \kappa D_f^3 G^2 / 2 > 0$ ; at fixed  $D$  the  $\alpha$ -decrease per added cell pair is exactly  $r_0^2/s_0$ ; and prime-power entries only raise  $\varepsilon$  — so the certified bound carries the full measured rate. Line-by-line attribution: the identities are exact algebra, (H1)/(H2) are completed Choleskys, the cell step is classical, the exponents are jackknifed fits.

*What is still missing — three named analytic lemmas about one symbol, each with a classical address:*

- (1) a uniform lower bound on *one* local slope of the dual solution  $u_f = T_f^{-1} \tilde{t}_f$  (H3 as analysis): the corner asymptotics of  $T^{-1}$  for a symbol with a log singularity — the Kac–Murdock–Szegő 1953 / Widom 1974 circle of results;
- (2) a lower bound for  $\lambda_{\min}(S)$  at *one* level (H2 as analysis):  $S$  is the form seen by oscillations, i.e. the symbol at the Nyquist frequency  $\pi/D$ , where the archimedean weight is largest — the measured monotonicity then carries it;
- (3) the saturation constant: dropping  $\varepsilon_f$  costs the factor  $1 - \varepsilon_f/\varepsilon_c$ , measured in  $[0.675, 0.744]$  — literally the Bank–Smith / Dörfler–Nochetto saturation assumption, bounded away from 0 by a constant, so it costs no rate.

Two of the three are *constants*, not rates. None of this touches the T116 [SUPPORT] comb obstruction — that is a separate question about *state compression*, not about  $\varepsilon$ .

*Status after T118 (next subsection):* (3) is **classically done** — saturation is an *identity* here,  $\varepsilon_c - \varepsilon_f = y^T S y$ , plus the one hypothesis Bank–Smith themselves assume; (2) is **certified on windows and arithmetically reduced** — not through the symbol of  $S$  itself, which is a weighted *harmonic* mean of the aliasing pair and vacuous on the comb, but through the strengthened Cauchy–Schwarz shift onto the oscillation Gram; (1) is **corrected and open** — the corner exponent does not settle, because a log singularity is the boundary of all powers, and (H3) must be upgraded to a mean-square statement, which also repairs the chain’s last step.

T118 (`symbol_lemmas_probe.py`, contract `SYMBOL.LEMMAS`) then measured the three lemmas against their classical addresses, one by one — and two of the three stand.

### 7.17 Two of the three lemmas stand (T118)

T118 (`symbol_lemmas_probe.py`, contract `SYMBOL.LEMMAS`, 36/36, TWO-OF-THREE) put each missing lemma against its classical address. The outcome, stated before the detail: the saturation lemma (3) is an *identity* here; the Schur-floor lemma (2) acquires a fully certified classical route on half the measured windows and an arithmetic reduction on all of them; the corner lemma (1) does *not* contradict its address — the address simply does not cover this symbol — and its failure *corrects* both an old exponent estimate and the last step of the T117 chain. What the theorem still needs is exactly one genuinely new analytic statement.

*The harmonic versus the arithmetic aliasing mean.* The two-level oscillation Schur complement of a Toeplitz form has an *exact* two-grid symbol (re-derived in closed form; local Fourier analysis, Brandt 1977/Hackbusch 1985; multigrid for Toeplitz, Fiorentino–Serra 1991), and it is the weighted *harmonic* mean of the symbol at a low frequency and at its Nyquist partner:  $1/\sigma_S = \cos^2(\theta/2)/f(\theta + \pi) + \sin^2(\theta/2)/f(\theta)$ . A harmonic mean needs *both* entries positive, and  $f$  is negative on a comb of dips at the window’s own resolution scale:  $\inf \sigma_S = 0$  on 14/14 (zone, level) rows — every pointwise

route through  $\sigma_S$  is *vacuous* on the real symbol, not merely weak (a perturbative split recovers at most  $2.6 \times 10^{-5}$  of  $\lambda_{\min}(S)$ ; the comb is non-perturbative). The classical reason is computed rather than asserted:  $\lambda_{\min}$  at finite  $M$  is a Fejér–Riesz *moment* problem, not a pointwise infimum — the ground state of  $S$  *concentrates* on the negative set ( $3.3\text{--}43\times$  the uniform share), and positivity is a cancellation of relative size  $10^2\text{--}10^4$  between the positive and negative parts of the symbol integral. Coarse-graining shows where the mass sits: the one-cell-averaged symbol is positive on 14/14 windows (a heuristic, declared as such). The growth is a *logarithm*, not a power: on an FFT-only lever of up to  $128\times$  in  $D$  the log model beats the power model (at  $b = 0.44\text{--}0.93$  per  $\log(1/D)$ ), exactly what the log-singular kernel demands (Kac–Murdock–Szegő 1953; Widom 1974). And the pole mass in the oscillation space is  $\leq 2.1 \times 10^{-5}$  — an identity via the DTFT factor  $|\sin(\theta/2)|$  — which is why the finite-section effect is absent there.

*Saturation is an identity here — and the CBS shift rescues Lemma B.* For a nested piecewise-constant pair, saturation is not an assumption but the *identity*  $\varepsilon_c - \varepsilon_f = y^T S y$  (rel  $4.9 \times 10^{-11}$ ), so  $\beta$  is *computed*:  $\beta \in [0.252, 0.336]$  over 20 triples, with the monotonicity in the favourable direction; what remains of Lemma C is the one hypothesis Bank–Smith 1993 themselves assume. Their strengthened Cauchy–Schwarz constant  $\gamma$  (Yserentant 1986) is then the lever that rescues Lemma B: the identity  $\lambda_{\min}(L_z^{-1} S L_z^{-T}) = 1 - \gamma^2$  (rel  $2.1 \times 10^{-15}$ ) moves the symbol question off  $S$  and onto the oscillation Gram, whose symbol  $\sigma_z(\phi) = \sin^2(\theta/2)f(\theta) + \cos^2(\theta/2)f(\theta + \pi)$  is the *arithmetic* mean of the same aliasing pair — the low-frequency negativity that poisoned the harmonic mean is suppressed *quadratically* by  $\sin^2(\theta/2)$ . The route is then certified:  $\inf \sigma_z > 0$  on 7/14 windows (systematically: every window with  $\alpha \leq 1.30$  and  $h \geq 200$ ), with the certified bound recovering up to 52% of the measured  $\lambda_{\min}(S)$ ; on *all* 14 windows the reduction  $\lambda_{\min}(S) \geq (1 - \gamma^2) \lambda_{\min}(T_h(\sigma_z))$  holds at 32–76% sharpness — Lemma B is reduced to the positivity of a plain Toeplitz section of an explicit symbol; and on the long FFT lever ( $M$  up to 15 680, no matrix ever formed) the certified floor  $\inf \sigma_z$  *rises* logarithmically and *crosses zero* on 3/5 zones — the failing windows were *under-resolved*, not obstructed.

*The corner lemma corrects itself — and repairs the chain.* The classical statement is an edge exponent: for a symbol vanishing like  $|\theta|^{2s}$  the finite-section solution vanishes like  $\text{dist}(\cdot, \text{edge})^s$ . The estimator is calibrated on a model class where this is exact ( $|2 \sin(\theta/2)|^{2s}$ , lags in closed form; the Kac–Murdock–Szegő 1953 closed-form inverse as the  $s = 0$  control). On the *real* symbol the corner exponent *drifts*:  $+0.238$  per halving of  $D$  ( $3\text{--}7.4\sigma$ ) —  $\log(1/|\theta|)$  is the boundary of all powers, so there is no asymptotic  $(s, \kappa)$  to cite. (The old T108/T109 24-cell exponent was doubly mis-estimated; its apparent stability was the cancellation of two errors.) The (H3) correction: the single-cell form is mis-dimensioned — the concentration factor grows like  $h^{0.99}$ , the slope mass sits on a fixed *fraction* of cells — and the repair is at the T117 chain’s last step: the rate lives on  $\|y\|^2$  (the stage-two bound carries  $D^{1.751}$  against  $\varepsilon \sim D^{1.770}$  — lossless; the single-cell stage costs  $D^{2.74}$ ), so (H3) must be a mean-square statement.

#### The theorem after T118 — five unconditional links, two hypotheses

*Candidate (V2).* Five unconditional links — the saturation identity  $\varepsilon_c - \varepsilon_f = y^T S y$ , the CBS identity  $\lambda_{\min}(L_z^{-1} S L_z^{-T}) = 1 - \gamma^2$ , the isometry ordering onto  $T_h(\sigma_z)$ , Grenander–Szegő in the lower direction, and the certified Cholesky floors — give, for every window the probe touched,

$$\varepsilon(\alpha, D) \geq (1 - \gamma^2) \lambda_{\min}(T_h(\sigma_z)) \|y\|^2.$$

Two hypotheses close it: **(H1)**  $\gamma$  bounded away from 1 uniformly in  $(\alpha, D)$  — the standard CBS hypothesis of the two-level literature, measured  $1 - \gamma^2 \geq 0.181$  with the favourable trend; **(H2)**  $\|y\|^2 \geq c D^{1.75}$  uniformly — (H3) upgraded to mean square, and *the* one genuinely new analytic statement the theorem needs, not importable from the Kac–Murdock–Szegő/Widom power-law corner theory because this symbol is log-singular with atoms.

*The per-lemma missing lists.* **L-B** (two points, both *arithmetic*):  $\inf \sigma_z > 0$  for all  $D < D_0(\alpha)$  —  $D_0$  set by the comb depth against the Nyquist-band mass, an atom-counting question — and narrow-dip positivity of finite sections. **L-C** (one point):  $\gamma$ -uniformity, i.e. (H1). **L-A** (two points): the mean-square form (H3’), and the Fisher–Hartwig extension of the corner theory to log-singular symbols with atoms.

*Disposition after T119 (next subsection):* the first L-B point is now a **theorem** with an explicit  $D_0(\alpha)$  and a universal constant; the second is diagnosed as a moment problem, with one window newly certified by the large sieve; (H2) is proven *not derivable* from the links above — genuinely new content, as suspected — and reduced, through the mean-square form (H3'), to a single scalar that is itself an exact identity,  $\kappa_{\text{end}} = 1/(1+R)$ ; what remains of the whole list is one discrete Harnack inequality plus two constants.

T119 (`oscillation_mass_probe.py`, contract `OSCILLATION.MASS`) then took the two halves that remained — the arithmetic half of Lemma B and the oscillation-mass half (H2)/(H3') — and closed the first as a theorem.

## 7.18 The arithmetic half closes; one Harnack inequality remains (T119)

T119 (`oscillation_mass_probe.py`, contract `OSCILLATION.MASS`, 27/27, ARITHMETIC-DONE) is the part where the comb, the standing obstruction since T116, becomes an explicit and *bounded* object. The outcome, stated before the detail: the arithmetic half of Lemma B is a theorem with an explicit  $D_0(\alpha)$ ; the narrow-dip point is diagnosed as a moment problem, with one window newly certified by the large sieve; (H2) is proven *not derivable* from the existing identities — and the sharpest identity of the run reduces the whole oscillation-mass half to one discrete Harnack inequality with a classical address.

*The  $D_0$  theorem, with a universal constant.* The lag sequence splits exactly as  $c = c_{\text{arch}} + c_{\text{comb}}$ , and the comb half has a *closed-form* symbol: every prime-power atom  $(u_j, \mu_j)$  contributes exactly one cosine at the frequency of its own lag, split over the two nearest integer lags by linear interpolation — and under the aliasing average  $\sigma_z$  the *parity* of the lag index decides everything (an even lag survives at full strength, an odd one is multiplied by  $-\cos \theta$ ). Two consequences, both elementary and both new: the whole comb is bounded *uniformly in  $D$*  by the atom count,  $|\sigma_z^{\text{comb}}| \leq \Xi(\alpha) = \sum_{n \leq e^{2\alpha}} 2\Lambda(n)/\sqrt{n} \sim 4e^\alpha$  (Chebyshev); and the archimedean growth is an *identity*,  $\inf \sigma_z^{\text{arch}} = \log(1/D) + B$  with slope exactly one and a *universal*  $B = -1.0474$  (drift  $\leq 3.1 \times 10^{-4}$  over  $\alpha = 0.98\text{--}4.70$ ; measured slope  $1.0030 \pm 0.0012$ ). Their conjunction is the theorem:

$$\inf \sigma_z > 0 \quad \text{for all } D < D_0(\alpha) = \exp(-(\Xi(\alpha) + B)),$$

verified deep on a  $1024 \times$  FFT-only lever ( $M$  up to 65 536, sixteen zones to  $n = 11\,981$ , no matrix ever formed).

*The two-regime correction, and the large sieve's one win.* The deep run corrects T118's reading: there are two regimes. Narrow zones ( $\alpha \leq 2.00$ , 5/16) cross zero *certified* at the frame's own resolution; wide zones ( $\alpha \geq 2.22$ , 11/16) also cross — the theorem is not in doubt — but at  $M \sim 10^{13}\text{--}10^{187}$ : a defect of the *constant*, not of the theorem. The narrow-dip point (L-B.2) becomes a certifiable criterion through the large sieve (Montgomery–Vaughan 1974, used in the honest direction: an upper bound on the  $L^2$  mass a polynomial of degree  $< h$  can place on the dips, which is exactly what a lower bound on  $\lambda_{\min}$  needs). It wins *once* — a certified positive section floor  $+0.396$  on a window where  $\inf \sigma_z = -0.349$ , the series' first certificate of section positivity *without* pointwise positivity — and it dies structurally for  $\alpha \geq 2.5$ , where a single Fejér bump fits inside one dip: the obstruction is a moment problem, not a localisation problem. A by-product is proved exactly: coarse-graining can only raise the section floor,  $\lambda_{\min}(T_h(\sigma * K)) \geq \lambda_{\min}(T_h(\sigma))$  by unitary averaging — the direction *opposite* to T118's heuristic reading, which is thereby retired.

(H2) is *genuinely new* — proven, not suspected.  $\|y\|^2$  is the squared  $L^2$  distance of the fine piecewise-constant solution from the coarse space, so (H2) is an *inverse* approximation statement. Its structure is dissected: the mass is distributed (corner share 0.18–0.42, middle 0.19–0.38, inner half 0.36–0.48 — no single-cell carrier), the exponent bookkeeping closes ( $\varepsilon \sim D^{1.761}$ ,  $\|y\|^2 \sim D^{1.961}$ ,  $\lambda_{\min}(S)\|y\|^2 \sim D^{1.757}$  — the second stage loses no power), and the Levinson structure of the solution is verified in full (the corner identity  $u_0 = -\sqrt{2}\rho/E$  to  $2.8 \times 10^{-8}$ , the mass budget  $1 - \varepsilon = \sum_n \rho_n^2/E_n$  to  $2.5 \times 10^{-13}$ , exactly one negative pivot). The derivability question is answered

*no*: the energy route is valid and empty — it returns  $\varepsilon \geq r \varepsilon$  with  $r < 1$ , a tautology — so (H2) is genuinely new analytic content, exactly as T118 suspected.

*The constant becomes an identity.* The sharpest result of the run removes the constant from (H3'): Cauchy–Schwarz is nearly sharp, the corner increments are monotone on 100% of the cells, and the end-cell share  $\kappa_{\text{end}} = 0.4906\text{--}0.5093$  is flat over 50 rows and 15 ladder pairs and *exponent-blind* — it drifts by  $\leq 0.0019$  per level while the corner exponent itself creeps by  $+0.046\text{--}0.242$  per level, a factor 128 less. And it is not a measured constant at all: for a monotone corner profile,

$$\kappa_{\text{end}} = \frac{1}{1 + R}, \quad R = \frac{\sum_j |a_j|}{\sum_j |w_j|},$$

*exactly* (verified to  $1.1 \times 10^{-16}$ ; a telescoping sum), with  $w_j$  the within-cell and  $a_j$  the across-cell increment of  $u$  and  $R$  measured  $0.9634\text{--}1.0349$ . The one hypothesis left inside the chain is therefore a *discrete Harnack inequality* for the increments of a log-singular Toeplitz inverse — the profile does not oscillate at the mesh scale — and not a bare constant.

#### The theorem after T119 — fourteen links, three defects

*Candidate (V3).* Fourteen links close the chain  $\varepsilon_c \geq (1 - \gamma^2) \lambda_{\min}(T_h(\sigma_z)) \|y\|^2 > 0$  for every window the probe touched: eight identities (among them, new in T119, the closed-form comb symbol, the parity rule, the Levinson corner value with its mass budget, and  $\kappa_{\text{end}} = 1/(1 + R)$ ), one arithmetic bound (the atom count  $\Xi(\alpha)$ , uniform in  $D$ ), one certified growth law (the archimedean identity with universal  $B$ ), one theorem (the explicit  $D_0$ ), the isometry, Cauchy–Schwarz, and the chain itself — five of the fourteen new in T119.

*The defect list, shortest form — three items:*

- (D1) the discrete Harnack inequality  $\sum_j |a_j| \leq R_{\max} \sum_j |w_j|$  on the corner cells — the one genuinely new analytic statement, with a classical address (Widom 1974; Trench 1964 for the explicit inverse; Böttcher–Silbermann); *any* finite  $R_{\max}$  gives  $\kappa_{\text{end}} \geq 1/(1 + R_{\max}) \approx 0.49$ ;
- (D2) the  $D_0$  constant for wide  $\alpha$  — the theorem is unconditional, but at the frame’s own resolution the pointwise floor is negative for  $\alpha \gtrsim 2$ , and the section route that would cover it is a Fejér–Riesz moment problem;
- (D3)  $\gamma$ -uniformity — (H1), inherited from T118 unchanged.

*Distance to a paper, honestly:* as a *conditional* lemma paper — under (D1) and (H1),  $\varepsilon \geq c D^\theta$  with an explicit  $c$  — the material is essentially complete, and the missing work is writing, not discovery; an unconditional version needs exactly the Harnack statement. The chain is classical numerical analysis end to end and contains no zeta input anywhere.

*Disposition after T120 (next subsection):* (D1) is *explained* —  $R \approx 1$  is parity symmetry, and the summed inequality follows unconditionally from one sign of the corner increments by a two-line pairing argument, while the per-cell version is provably false — so the one hard statement is replaced by a strictly weaker and far more classical hypothesis; but (D2) is enlarged (the frame consumes broad windows by construction, and the unconditional  $D_0$  route covers not one handover of the real ladder) and (D3) loses ground as a uniformity claim ( $1 - \gamma^2$  drifts down with  $\alpha$ ): the defect count goes three  $\rightarrow$  four.

T120 (`harnack_probe.py`, contract HARNACK) then attacked that defect list head on — and explained the Harnack core while sharpening the side defects.

### 7.19 The Harnack core is proof-shaped; the wide-window defects sharpen (T120)

T120 (`harnack_probe.py`, contract HARNACK, 21/21, HARNACK-EXPLAINED) is the part where the one hard statement of V3 stops being an inequality about two families and becomes a two-line consequence of a symmetry. The outcome, stated before the detail:  $R \approx 1$  is *parity symmetry* — the two increment families are one sequence offset by a single fine cell — with an unconditional  $|R - 1| \leq 0.047$  certificate; two structural routes to the remaining hypothesis are refuted, and the per-cell version of the inequality is provably *false*; and the honest defect count rises from three to four, because the  $\alpha$ -audit shows the unconditional  $D_0$  route covers not one handover of the real ladder and the CBS constant decays with the width.

*The parity symmetry, and the certificate.* The two families of (D1) are not two objects: with  $v_i = u[i+1] - u[i]$  the single increment sequence of the *same* solution  $u = T^{-1}\tilde{t}$ , the within-cell and across-cell increments are the even and the odd half,  $w_j = v_{2j}$  and  $a_j = v_{2j+1}$  — offset by exactly one fine cell. If  $v$  has one sign on the corner cells (measured at 100% of them on every one of 3915 rows), the difference of the two parity sums is a sum over *disjoint* adjacent pairs, so

$$|R - 1| \leq \frac{\sum_j |v_{2j+1} - v_{2j}|}{\sum_j |v_{2j}|} \quad (\text{C1})$$

*unconditionally* — no monotonicity, no smoothness. Measured over the whole sweep, (C1) gives  $|R - 1| \leq 0.04745$  against a measured 0.03485, and it dominates the measurement on all 3112 sign-pure rows, as it must. If  $v$  is moreover monotone the left side telescopes to a boundary term over a bulk sum,  $|R - 1| \leq |v_{2n_C} - v_0| / \sum_j |v_{2j}| \leq (C - 1)/n_C$  (C2) — recorded as *conditional* (the monotone share is only 0.50–0.96) and kept as the explanation of the rate:  $R$  sits at 1 because a shift by one fine cell is a boundary perturbation, and  $R$  is resolution-blind because refining raises  $n_C$  and shrinks the boundary term at the same rate. The certificate itself confirms the rate: (C1) decays like  $n_C^{-0.58 \pm 0.05}$  (a fit). The brute-force sweep (155 windows  $\times$  resolutions  $\times$  corner fractions, 3915 rows) keeps  $R \in [0.9098, 1.1011]$  in the corner regime with sign purity exactly 1.0000; the single excursion above 1.1 sits at the clamped minimal corner  $n_C = 4$ , exactly where the boundary-term mechanism predicts it, and for  $n_C \geq 8$  the window certificate is  $\kappa_{\text{end}} \geq 0.480827$ . One negative result belongs to this positive one: the *per-cell* Harnack inequality is *false* —  $|a_j/w_j|$  reaches  $4.8 \times 10^3$  — so no cellwise comparison can be true; only the summed form holds, which is exactly the form (C1) proves.

*Two structural routes refuted.* Both obvious routes to the one remaining sign hypothesis are now closed, each by a measurement. First, the  $40 \times 40$  corner block of  $T^{-1}$  is entrywise positive on 18/18 (zone, level) rows without exception — and the checkerboard alternative is refuted exactly (0.5000, a coin flip) — but it is *not* TP2 (the share of non-negative  $2 \times 2$  minors is only 0.51–0.88), so the Gantmacher–Krein oscillation-matrix route is closed; nor is  $T$  an M-matrix. Second, differencing  $Tu = \tilde{t}$  once gives an *exact* discrete boundary-value problem for the increments alone,  $Kv = s - u_0G\mathbf{1}$  with  $K = D_1TL$  (relative residual  $2.4 \times 10^{-12}$ ), with a smooth, positive right-hand side — the correct classical shape for a maximum-principle argument — but  $K^{-1}$  is *not* entrywise positive (positive share 0.44–0.51, minimum entry  $-333$ ): the discrete maximum principle is not unproved, it is *false*. The reason is visible in the solution itself:  $u$  is monotone at the corner and oscillatory in the bulk — that is the prime-power comb. The corner block of  $K^{-1}$  is entrywise positive (0.9961–1.0000), so a localised statement is not excluded — but the corner components of  $v$  contract full rows of  $K^{-1}$  carrying a genuine cancellation (no-cancellation ratio  $10^{-7}$ – $10^{-3}$ ), so what the sign hypothesis needs is a corner-localised Fisher–Hartwig decay estimate (Widom 1974; Böttcher–Silbermann), not a sign pattern.

*The  $\alpha$ -audit: (D2) is not vacuous but total.* The frame’s window is not a free parameter: in frame A the handover at the atom  $u_k = \log n_k$  uses  $D = g_k/(2\nu)$  and  $M_o = \lceil u_k/D \rceil + 1$ , hence  $\alpha_o = u_k/2 + O(D)$  *exactly* (maximum deviation  $1.5 \times 10^{-2}$  over all 732 admissible handovers) — the consumed half-width is half the log of the zone, and only 1.4% of the handovers have  $\alpha \leq 2$ . Because  $B = -1.0474$  is a property of the archimedean kernel alone and hence universal (spread  $1.2 \times 10^{-6}$ ), the criterion  $\log(1/D_{\text{frame}}) \geq \Xi(\alpha) + B$  can be evaluated on the whole zone table: it holds on exactly 3 of 1492 zones ( $n = 2, 3, 4$ ), i.e. for  $\alpha \leq \alpha^* = 0.693$  — *below* the first admissible handover  $n = 7$ . Not one step of the induction is covered by the unconditional  $D_0$  route (the resolution it would need explodes to  $M \sim 10^2$ – $10^{188}$ ). The fallbacks, measured and labelled: the comb supremum is only 0.038–0.603 of its certified cap  $\Xi$  — an incoherence gain of up to  $1.7\times$ , but a grid supremum, hence a measurement and never a certificate (a certified version would be an exponential-sum estimate over prime powers) — and the large-sieve narrow-dip floor (Montgomery–Vaughan 1974) is positive on 2/6 surveyed zones, at  $n = 32$  where  $\inf \sigma_z = -0.45$  — the dip route genuinely buys something — and dies beyond  $\alpha \approx 2$ .

*The  $\gamma$  formula, and the wrong direction.* The CBS constant has an exact local symbol identity:

with  $f_1 = f(\theta)$ ,  $f_2 = f(\theta + \pi)$ ,  $\sigma_c \sigma_z - \tau^2 = f_1 f_2$  (relative residual  $3.9 \times 10^{-16}$ ), hence  $1 - \gamma(\theta)^2 = f_1 f_2 / (\sigma_c \sigma_z) - (D3)$  is no longer a bare measured constant but a symbol infimum, the same class of object as  $\inf \sigma_z$ . Two honest caveats, both measured: the naive symbol infimum is  $25\text{--}210 \times$  below the matrix constant (the finite odd section never sees the worst aliasing pair — a finite-section transfer is needed, not the infimum), and the measurement itself moved:  $1 - \gamma^2$  falls with  $\alpha$  (fitted slope  $-0.297 \pm 0.013$ , range  $0.079\text{--}0.283$ , flat in the resolution at a fixed zone) — T119's  $\geq 0.181$  is reproduced on its own  $\alpha$ -range but is *not* uniform.

#### The theorem after T120 — eighteen links, four defects

*Candidate (V4).* Eighteen links (7 identities — among them, new in T120, the increment system  $Kv = s - u_0 G1$  and the symbol identity  $\sigma_c \sigma_z - \tau^2 = f_1 f_2 - 4$  certified, 2 classical, 2 measured, and the chain), and the one hard statement of V3 is *replaced*: (D1) is no longer a bare inequality but a proved two-line implication — the pairing certificate (C1), unconditional given one sign of the corner increments. What a lemma paper could claim today: *let  $T$  be the finite odd Weil form on a window,  $u = T^{-1}\tilde{t}$ , and suppose the increments of  $u$  have one sign on the outer  $n_C$  coarse cells; then  $|R - 1| \leq \delta(n_C)$  with the explicit pairing ratio  $\delta$ , hence  $\kappa_{\text{end}} \geq 1/(2 + \delta)$  and  $\|y\|^2 \geq (u_{\text{end}} - u_0)^2 / (2n_C(2 + \delta)^2)$  — everything after “suppose” is proved; measured,  $\delta \leq 0.0964$  for  $n_C \geq 8$ , falling like  $n_C^{-0.58}$ .*

*The defect list — four items (was three):*

- (E1) *one sign of the corner increments* — the single hypothesis of (C1)/(C2), measured on 100% of the corner cells on all 3915 rows; the two obvious structural routes are refuted (the corner of  $T^{-1}$  is positive but not TP2; the global discrete maximum principle is false), so what remains is a corner-localised Fisher–Hartwig estimate (Widom 1974; Böttcher–Silbermann);
- (E2) *a uniform bound on the pairing ratio  $\delta$*  — a discrete gradient estimate for the increments of  $T^{-1}\tilde{t}$ , strictly weaker than the old Harnack statement (any finite  $\delta$  suffices,  $\kappa_{\text{end}} \geq 1/(2 + \delta)$ );
- (E3)  *$D_0$  for the whole ladder* — (D2), enlarged: the frame consumes  $\alpha = u_k/2$  by construction and the criterion fails at  $\alpha^* = 0.693$ , below the first admissible handover; the open routes are a certified comb-incoherence bound or the large-sieve dip geometry, which works on the first zones and dies beyond  $\alpha \approx 2$ ;
- (E4) *the  $\alpha$ -dependence of  $\gamma$*  — (D3) with the direction now known:  $1 - \gamma^2$  shrinks like  $e^{-0.30\alpha}$  (a fit), so uniformity is not merely unproved but questionable in the direction that matters.

*Honest movement:* the one hard statement got much easier — a two-line implication with one classical hypothesis — and the side defects got sharper and both got worse; the count went three  $\rightarrow$  four, and that is the honest direction.

*Disposition after T121:* (E3) is retired *as stated* — the chain link is a section eigenvalue, not a symbol infimum, and the section stayed positive on the whole real ladder — and is replaced by the section lemma (E3\*) with the Hankel term; (E4) is replaced by one named poly-log, (E4\*); (E1) and (E2) keep their kind, with the corner certificates restated at  $n_C = h/16$  (T121 found the  $h/8$  convention used here too wide); and the symbol-to-section transfer this list implicitly relied on is refuted (the odd sector's Hankel term). The next subsection reports the measurements.

## 7.20 Section positivity survives the real ladder; the net balance names one poly-log (T121)

T121 (`wide_window_probe.py`, contract WIDE.WINDOW, 21/21, WIDE-RESTRUCTURED) asks the question the T120 defect list left sharpest: does the  $\sigma_z$  route survive the *real* ladder — the one whose window grows with the zone,  $\alpha_o = u_k/2$ ? It does, but not as the statement T120 wrote down. The route survives as a *section* statement and dies as a *symbol* statement; the net  $\alpha$ -balance of the whole chain loses only a single poly-log to a demand that itself falls like  $\alpha^{-6}$ , and that poly-log splits without residue into two repairable steps; one link of the chain, filed by T120 as classical, is honestly refuted on the way.

*Section positivity instead of symbol positivity.* The chain never needed  $\inf \sigma_z > 0$ ; it needs  $\lambda_{\min}$  of the *section*  $T_h(\sigma_z)$  at the frame's own resolution. Over 16 frame windows of the real ladder ( $n = 7 \dots 283\,303$ ,  $\alpha = 0.985 \dots 6.277$ ,  $M_{\text{frame}} = 118 \dots 431\,162$ ,  $h = 29 \dots 107\,790$ ) the section is *positive* on 16/16 while the symbol infimum is negative on 8 of them — and where it is negative,

the section eigenvalue sits at  $0.64\text{--}3.23 \times |\inf \sigma_z|$  on the *positive* side. Direction kept honest: the dense rows ( $h \leq 1700$ ) carry a shifted-Cholesky certificate with the Wilkinson–Higham floor; the large rows are matrix-free FFT-Lanczos *measurements* (a Ritz value is an upper bound for  $\lambda_{\min}$ , so it can refute positivity but never prove it). The mechanism is quantified at its classical address — the Christoffel function (Szegő 1939; Grenander–Szegő 1958, Ch. 5): a degree- $h$  polynomial cannot concentrate below the Fejér bandwidth  $2\pi/h$ , and every comb dip is only 0.156–0.720 of one Fejér width. The certified per-dip moment budget  $\lambda_{\min} \geq \text{thr} - \sum_j (\text{thr} + \delta_j) w_j (h+1)$  is positive on 8/16 rows (up to  $\alpha = 3.446$ , sharpness 0.20–0.76) and tears above, for a computed reason: the dip *count* grows like  $e^{2.41\alpha}$  (a fit) against the budget’s factor  $h$ . So (E3) is reduced from a symbol-infimum requirement that no ladder step satisfies to a *section lemma* the measurement satisfies everywhere: a Christoffel lower bound surviving  $O(e^\alpha)$  Fejér-subresolved dips — a Szegő/Fisher–Hartwig question about the comb, no longer a symbol question.

*The net  $\alpha$ -balance — the core number of the part.* Every factor of the chain is measured on one two-dimensional  $(D, \alpha)$  lever (14 zones  $\times$  3 resolutions, 42 rows). The individual links do fall with  $\alpha$ , exactly as T120’s (E4) recorded:  $1 - \gamma^2 \sim \alpha^{-0.79}$ ,  $\lambda_{\min}(A_z) \sim \alpha^{-1.07}$ ,  $\|y\|^2 \sim \alpha^{-5.63}$ , so the assembled supply falls like  $\alpha^{-7.49}$ . But the *demand* falls almost as fast ( $\alpha^{-5.93}$  on this lever; T116 quoted  $-6.04$ ): 79% of the decay (E4) complained about is the demand’s own. What is left is the only  $\alpha$  deficit of the chain,

$$\text{SUPPLY}/\varepsilon_c \sim D^{-0.177 \pm 0.096} \alpha^{-1.565 \pm 0.135}$$

(fits; the measured ratio is  $4.9 \times 10^{-3}$ –0.15 over  $\alpha = 0.98$ –6.31) — net-negative in  $\alpha$  by one to two powers (the refit per resolution lands in  $-1.88 \dots -1.62$  — a range, not a decimal) but *uniform in  $D$* : the  $D$  exponent is zero within its own band, so refining the resolution neither helps nor hurts. The deficit is then decomposed *exactly*: the identity  $y^\top S y / \text{SUPPLY} = r_{\text{ray}} r_{\text{cbs}} / (1 - \gamma^2)$  holds to  $2.8 \times 10^{-14}$ , so  $\alpha^{-1.57}$  splits without residue into  $\alpha^{-1.02}$  from the *Rayleigh step* (using  $\lambda_{\min}(A_z)$  on a  $y$  that is nowhere near the minimal eigenvector; measured slack 1.29–13.19),  $\alpha^{-0.51}$  from the CBS step (measured slack 3.42–16.67) and  $\alpha^{-0.04}$  from discarding  $\varepsilon_f$  — *nothing* of it sits in the value of  $\gamma$ . The named rebuild is therefore structural, not a new constant: bounding  $y^\top A_z y$  directly (legitimate, because  $y$  is the oscillation part of the solution, not an arbitrary vector) recovers 65% of the deficit by itself, and the CBS step’s own slack covers the rest.

*Link (5) refuted: the Hankel term.* The section comparison  $\lambda_{\min}(A_z) \geq \lambda_{\min}(T_{h_c}(\sigma_z))$  — which T120’s assembly used as a classical Grenander–Szegő transfer — is *false* on 42/42 rows (ratio 0.357–0.916, drifting with the window). The reason is structural and was hiding in the odd coordinates:  $A_z = Z^\top (\text{Toeplitz} - \text{Hankel}) Z$ , and Grenander–Szegő carries the even Toeplitz part only; the reflection Hankel term is not sign-definite. A defect made visible rather than created: the section lemma (E3\*) must be about  $A_z$ , not about  $T_{h_c}(\sigma_z)$ .

*The corner sign, and the inconclusive Fisher–Hartwig comparison.* The one-sign run of the increments covers the corner region the (C1)/(C2) certificates need on 24/24 (zone, resolution) rows at  $n_C = h/16$  (run share 0.21–0.42 of the section against the required 0.125 — a safety factor of 1.68–3.38, the sign uniformly negative); at the  $h/8$  convention T119/T120 used it holds on only 21/24, so the certificates are restated at  $h/16$  — which costs nothing, because  $\delta$  grows like  $n_C^{+0.29}$  (a fit): halving  $n_C$  *improves*  $\kappa_{\text{end}}$ . The  $\alpha$ -slope of the run share is  $-0.0258 \pm 0.0042$  (a fit); extrapolated linearly, the  $h/16$  margin would be exhausted only near  $\alpha \approx 11.7$  (an extrapolation, not a prediction). The Fisher–Hartwig comparison itself is *inconclusive*, and that is the result: the local exponent of the corner row of  $T^{-1}$  drifts by  $-0.008 \dots + 2.934$  between dyadic windows where the log-corrected FH shape would allow only  $-0.120$  — the corner is in no asymptotic power regime at any resolution a factorisation reaches, and an FH exponent quoted from these numbers would be a fit masquerading as a structure. What the block does deliver: the no-cancellation ratio  $4.6 \times 10^{-8}$ – $7.1 \times 10^{-5}$  settles that (E1) can *never* be a sign-pattern argument, and the minimal sufficient statement (E1’) is written out as an explicit head-versus-tail inequality (head share  $1/(1 + 1.63)$  to  $1/(1 + 11.73)$  of the mass at  $J = h/16$ ), with the tail exponent as the one unknown a Widom-type estimate must supply.

$\delta$  over the full ladder. The pairing quotient stays in 0.0126–0.1331 over 360 rows spanning  $n = 7 \dots 299\,983$  and three corner fractions — always below 1, no outlier — so  $\kappa_{\text{end}} \geq 1/(2 + \delta) \geq 0.4688$  everywhere measured, and  $\delta \sim n_C^{+0.29}$  confirms the  $h/16$  restatement is free.

#### The theorem after T121 — the count stays four, every entry changes kind

*Candidate (V5).* Nineteen links: the identities (Schur/Haynsworth, CBS/Yserentant, the T120 symbol identity, the aliasing form of  $\sigma_z$ ,  $\kappa_{\text{end}} = 1/(1 + R)$ ), the certified pieces (the pairing certificate (C1); the shifted-Cholesky section floors; the Christoffel/Fejér per-dip budget up to  $\alpha = 3.446$ ), the classical steps (CBS + Rayleigh, lower; Cauchy–Schwarz, lower), the measured/fitted entries (section positivity on the ten Lanczos rows; the one-sign run; the net balance) — and one link *refuted*: (5), the symbol-to-section transfer, killed by the odd sector’s Hankel term.

*The defect list — four items, each with a changed kind:*

- (E1) *one sign of the corner increments* — unchanged in kind, sharpened twice: the certificates move to  $n_C = h/16$  (clean on all rows there, only 21/24 at  $h/8$ ; the move is free since  $\delta \sim n_C^{+0.29}$ ), and (E1') is now an explicit head-versus-tail inequality whose tail exponent a corner-localised Fisher–Hartwig/Widom estimate must supply — the FH comparison is inconclusive at reachable  $h$ , and no sign-pattern argument can ever do the job (no-cancellation ratio  $4.6 \times 10^{-8}$ – $7.1 \times 10^{-5}$ );
- (E2) *a uniform bound on the pairing ratio  $\delta$*  — unchanged: no outlier over 360 rows to  $n = 299\,983$ ,  $\delta = 0.0126$ – $0.1331$ , any finite  $\delta$  suffices;
- (E3\*) *the section lemma, with the Hankel term* (replaces (E3)) —  $\inf \sigma_z > 0$  at frame resolution is *retired as a requirement*: the chain link is a section eigenvalue, positive on 16/16 frame windows up to  $\alpha = 6.277$ ,  $h = 118\,115$ , including 8 windows with a negative symbol infimum. What replaces it is strictly smaller but has two parts: (a) a Christoffel-function lower bound surviving  $O(e^\alpha)$  Fejér-subresolved dips (the certified budget delivers it for  $\alpha \leq 3.446$  and tears above); (b) the odd sector’s Hankel term, which refutes link (5) as stated — the lemma must be about  $A_z = Z^\top(\text{Toeplitz} - \text{Hankel})Z$ , not about  $T_{h_c}(\sigma_z)$ ;
- (E4\*) *one residual poly-log in  $\alpha$*  (replaces (E4)) — the net balance carries  $\alpha^{-1.565 \pm 0.135}$  and  $D^{-0.177 \pm 0.096}$ : net-negative by one to two powers of  $\alpha$  (per-resolution refits  $-1.88 \dots -1.62$ ), *uniform in  $D$* . Strictly better than (E4) as T120 wrote it: 79% of the supply’s  $\alpha$ -decay is the demand’s own, and the deficit sits in the Rayleigh step ( $\alpha^{1.02}$ ) and the CBS step ( $\alpha^{0.51}$ ), none of it in  $\gamma$  — the address is a structural lower bound for  $y^\top A_z y$ , which alone recovers 65%.

*Honest movement:* the counter stays at four — the target of getting back below three was not met, because the residual  $\alpha$  deficit is real — but every entry changed kind, and one link recorded as classical was found refuted, a defect made visible rather than created.

*Disposition after T122:* both named repairs landed. Link (5) is retired by an *identity* —  $A_z = W^\top \text{Toeplitz}_M(c) W$  with  $W = BZ$  isometric, so the Hankel term is the reflection half of an isometry and not an error term — and the certified band bound that replaces it is non-vacuous to  $\alpha = 6.31$ , where the Christoffel budget tore at 3.45; the Rayleigh step is made structural and is *sharp* (residual slack 1.00–1.03, flat in  $\alpha$ ); the certified deficit halves to  $\alpha^{-0.73}$ , exactly uniform in  $D$ , and with the measured CBS coupling the balance closes to within the chain’s own ceiling. (E3\*)/(E4\*) are thereby absorbed into one band-margin estimate (F3) plus one structural CBS estimate (F4). The next subsection reports the measurements.

## 7.21 The Hankel term is an isometry half; the certified deficit halves (T122)

T122 (`net_balance_repair_probe.py`, contract NET.BALANCE.REPAIR, 20/20, NET-IMPROVED) attacks the two worst-case steps the T121 decomposition named — the Rayleigh step and the link-(5) transfer — and both repairs land. The headline is structural: the odd sector’s Hankel term, the object that refuted link (5), was never an error term to be estimated.

*The Hankel repair is an identity, not an estimate.* The compressed Hankel term has a closed form,  $(Z^\top H Z)_{jk} = b_{j+k}$  with  $b_s = (c_{K-2s} - 2c_{K-2s-1} + c_{K-2s-2})/2$  — half a second difference of the lag sequence, exactly like the Toeplitz-side lags — verified to  $7.0 \times 10^{-16}$  on all 36 (zone, resolution) rows. So the oscillation injection  $Z$  suppresses *both* halves by the same DTFT factor  $|2 \sin(\theta/2)|^2$ : the measured norm ratio is 0.123–0.781, i.e.  $O(1)$  — the T118 mechanism does not make the Hankel term small. The honest anatomy (the archimedean half numerically of rank 10–14, the

Hartman/Peller smoothing; the comb half corner-localised on few anti-diagonals) is structure, not a bound: the valid Weyl floor  $\lambda_{\min}(T_{h_c}) - \|Z^\top H Z\|$  is positive on only 12/36 rows (up to  $\alpha = 2.44$ ) and nowhere near sharp. The valid replacement for link (5) is an *identity*: with  $W = BZ$ ,

$$A_z = W^\top \text{Toeplitz}_M(c) W, \quad W^\top W = I \quad (\text{exact to } 0.0)$$

— the Hankel term *is* the reflection half of an isometry. From this and a certified cell envelope ( $g_{\text{thr}} = (\text{thr} - \ell)_+ \geq 0$ ; Parseval is exact because  $|Wv|^2$  has degree  $< M$ ) follows the certified band bound

$$\lambda_{\min}(A_z) \geq \text{thr} - \lambda_{\max}(G), \quad G = W^\top \text{Toeplitz}_M(g_{\text{thr}}) W \succeq 0,$$

non-vacuous on 36/36 rows up to  $\alpha = 6.31$  — where the T121 Christoffel budget tore at  $\alpha = 3.45$  — recovering 0.429–0.995 of the true  $\lambda_{\min}(A_z)$ ; the  $e^{2.41\alpha}$  dip count no longer enters anywhere (measured gain over the per-dip Nikolskii budget:  $8.4\text{--}1.1 \times 10^2$ ). The honest caveat is where the threshold sweep puts its optimum:  $\text{supp}(g) = 0.95\text{--}1.00$  of the circle — one certified end of a one-parameter family whose other end (small  $\text{thr}$ ) is the dip budget that tore.

*The structural Rayleigh step is sharp.* “ $y$  is high-band” is a statement about  $Wy$  at the window resolution (99.3–100.0% of its mass above  $|\theta| = \pi/2$ , with  $\|Wy\|^2 = \|y\|^2$  exact), not about  $y$  itself (0.1–1.7% high-band in the coarse coordinates — the oscillation lives in the injection, not in the coefficients). The same certified band identity, evaluated on  $y$  instead of over the whole oscillation space —  $y^\top A_z y \geq \text{thr} \|y\|^2 - y^\top G y$  — is positive on 36/36 rows, never exceeds the truth, and beats the worst-case product  $\lambda_{\min}(A_z) \|y\|^2$  by a factor 1.28–18.90: the slack it removes is exactly T121’s  $r_{\text{ray}}$ , growing like  $\alpha^{+0.83}$ , and the residual slack of the structural bound is  $q_{\text{str}} = 1.00\text{--}1.03$  with drift  $\alpha^{-0.002 \pm 0.001}$  — flat, i.e. *sharp*: no  $\alpha$ -dependence of the chain passes through the Rayleigh step any more. The CBS step is equally pessimistic: the actual coupling is  $\kappa_y = 0.0067\text{--}0.5311$  against the worst case  $\gamma^2 = 0.7173\text{--}0.9647$  (a factor of up to 100+, growing like  $\alpha^{+1.52}$ ), and  $1 - \kappa_y = r_{\text{cbs}}$  exactly — T121’s slack, now with a name. One lever limit declared: at  $M \leq 2048$  the coarse envelope of  $\sigma_z$  is nonnegative on 13/36 rows — the deep dips T118/T121 chased appear only at the frame’s own resolution.

*The new net balance.* Five supplies on one  $(D, \alpha)$  lever (12 zones  $\times$  3 resolutions, 36 rows), all sound (36/36 below the truth  $S_4 = y^\top S y$ ). The T121 baseline reproduces exactly ( $S_0/\varepsilon_c \sim \alpha^{-1.562 \pm 0.163}$ ); repairing link (5) alone does not change the exponent at all ( $\alpha^{-1.565}$  — a uniform floor cannot beat  $\lambda_{\min}(A_z)$ ); V1’s contribution is that the floor is certified over the whole ladder instead of tearing at  $\alpha \approx 3.45$ ); with both repairs the *certified* balance is

$$S_2/\varepsilon_c \sim D^{-0.003 \pm 0.075} \alpha^{-0.729 \pm 0.090}$$

(fits; measured 0.033–0.207): the deficit *halves*, and it is exactly uniform in  $D$ . Adding the *measured* CBS coupling gives  $\alpha^{-0.113 \pm 0.063}$  — statistically indistinguishable from the chain’s own ceiling  $S_4/\varepsilon_c \sim \alpha^{-0.116 \pm 0.062}$  (the cost of discarding  $\varepsilon_f$ , untouched by both repairs). The residual decomposition closes to  $2.9 \times 10^{-16}$ :  $-0.729 - 0.002 + 0.615 = -0.116$ . In one sentence: V1 buys the ladder, V2 buys the exponent, and the chain closes *up to its own ceiling* as soon as the CBS step is no longer a worst case.

#### The theorem after T122 — four defects, consolidated

*Candidate (V6).* The chain renumbered: the Haynsworth identity; the exact CBS split  $y^\top S y = (1 - \kappa_y) y^\top A_z y$  with  $\kappa_y \leq \gamma^2$  classical (Yserentant); the new two-grid identity  $A_z = W^\top \text{Toeplitz}_M(c) W$  ( $W = BZ$ ,  $W^\top W = I$ ) — the line that retires link (5); the certified cell envelope; the certified band bound (5R) with its uniform,  $y$ -free variant (Gershgorin); and the quoted demand fit.

*The defect list — four items (T119 = 3, T120 = T121 = V6 = 4; the target  $< 4$  not met, but the kinds changed twice over):*

- (F1) *one sign of the corner increments* — unchanged from T120/T121 and untouched by this part: 24/24 clean at  $n_C = h/16$ , no sign-pattern argument possible, the minimal statement is the head-versus-tail inequality (E1’) with an open tail exponent;

- (F2) *a uniform bound on the pairing ratio  $\delta$*  — unchanged and untouched: 0.0126–0.1331 over 360 rows, no outlier, a discrete gradient estimate still missing;
- (F3) *an analytic  $(D, \alpha)$  bound for the band margin  $1 - y^\top Gy / (\text{thr} \|y\|^2)$* , measured 0.024–0.363 over  $\alpha = 0.98$ –6.31 with fitted drift  $D^{-0.042} \alpha^{-0.525}$  — the one *new* item, and it absorbs link (5), the Nikolskii budget and both readings of (E3)/(E4): a Fejér/Christoffel question about how the comb dips of the window symbol overlap the high-band mass of the two-level residual, not an eigenvector question. An  $\alpha$ -flat band margin would carry the certified chain to within the CBS step of the truth;
- (F4) *a structural CBS estimate* —  $\kappa_y \leq \gamma^2$  is the last worst-case step in the chain (worth  $\alpha^{+0.615}$  in the balance), and the actual coupling is only 0.009–0.551 of it.

*Honest movement:* link (5) is not repaired but *retired* — it was a false comparison between an isometric compression and one of its two halves — and the per-dip Nikolskii budget went with it; what is left of the old (E3)/(E4) is a single scalar estimate about dips versus high-band mass, plus the CBS item.

*Disposition after T123:* the two scalar items were run against the same band machinery and they *split*. (F3) closes almost for free — the reduction to the band profile of  $Wy$  costs under 1%, and the oscillation block gains uniform certified positivity, (5R\*) — while (F4) resists structurally: every certified route to  $\kappa_y$  needs the *coarse* form from below, whose near-null direction is smooth and lives on a cancellation that no pointwise symbol minorant can see. The remaining  $\alpha^{+0.500}$  gap to the ceiling is thereby identified with the  $\varepsilon_f$  lid itself — the recursion base case. The next subsection reports the measurements.

## 7.22 The two-isometry completion: the margin closes, the coupling resists (T123)

T123 (`structural_cbs_probe.py`, contract STRUCTURAL.CBS, 20/20, CBS-RESISTS) attacks the two scalar estimates the V6 list left — the structural CBS step (F4) and the band margin (F3) — with the same band machinery that made the Rayleigh step sharp in T122. The margin closes almost for free; the coupling does not, and the reason *why* is the content of the part.

*The coupling anatomy, in closed form.* The coarse block is the *other* half of the T122 construction: with  $V = BP$  ( $P$  the piecewise-constant prolongation),  $V$  is an isometry too —  $V^\top V = I$  and  $V^\top W = 0$  exact — so the coarse, oscillation and coupling blocks are three compressions of *one* window Toeplitz form,

$$A_c = V^\top \text{Toeplitz}_M(c) V, \quad B_x = V^\top \text{Toeplitz}_M(c) W, \quad A_z = W^\top \text{Toeplitz}_M(c) W$$

(identities to  $3.4 \times 10^{-15}$  on all 24 rows). Two new closed forms complete the T122 anatomy: the coupling is half a *first* difference of the lag sequence,  $(B_x)_{jk} = \beta_{j-k} - (c_{K-2(j+k)} - c_{K-2(j+k)-2})/2$  with  $\beta_m = (c_{|2m+1|} - c_{|2m-1|})/2$ , verified to  $1.8 \times 10^{-15}$  ( $A_z$  was half a *second* difference), and in symbol form

$$b(\varphi) = -\frac{i}{2} \sin(\varphi/2) [\sigma(\varphi/2) - \sigma(\varphi/2 + \pi)]$$

— the *aliasing difference* of the window symbol times  $\sin(\varphi/2)$ , verified to  $5.7 \times 10^{-16}$ : the coupling weight vanishes exactly where the coarse space lives. That explains  $\kappa_y$  quantitatively for the first time:  $Wy$  carries 99.64–100.00% of its mass at  $d = |\theta - \pi| < \pi/2$  and only  $8.9 \times 10^{-6}$ – $1.5 \times 10^{-3}$  in the cell containing  $\theta = 0$ , while the worst-case  $y^*$  that attains  $\gamma^2$  is a *middle-band* vector (81.22–100.00% of its mass exactly there) — the factor 2–109 between  $\kappa_y = 0.0067$ –0.4622 and  $\gamma^2 = 0.7248$ –0.9467 *is* that mismatch.

*The certified  $\kappa$  bound: both routes fail for one reason in two guises.* Route 1, band-split Cauchy–Schwarz (per frequency band separately, the strengthened-CBS mechanism of Axelsson–Vassilevski, including the exact two-grid factorisation that puts the  $\sin(\varphi/2)$  zero into the certificate), is correct on 24/24 rows and non-vacuous on 0/24: it costs exactly one coarse condition number ( $\text{cond}(A_c) = 2.1 \times 10^4$ – $3.1 \times 10^7$ ). Route 2 substitutes the certified minorant  $T_M(\ell) \preceq T_M(\sigma^{(M)})$  into the Schur complement’s minimum characterisation and delivers two new links:

$$(5R^*) \quad y^\top A_z y \geq y^\top E_z y, \quad (7R^*) \quad y^\top S y \geq y^\top E_z y - (E_{xy})^\top E_c^{-1} (E_{xy}),$$

with  $E_c = V^\top T_M(\ell) V$ ,  $E_x = V^\top T_M(\ell) W$ ,  $E_z = W^\top T_M(\ell) W$ . (5R\*) holds with  $\lambda_{\min}(E_z) = +0.303 \dots + 5.017 > 0$  on *every* row of the ladder — uniform certified positivity of the oscillation block,

at slack 1.0007–1.007, subsuming the whole T122 threshold sweep for free (since  $\min(\ell, \text{thr}) \leq \ell$  for every threshold) — while (7R\*) is certified as soon as  $E_c \succ 0$ , which holds on 0/24 rows ( $\lambda_{\min}(E_c) = -3.44 \dots - 0.043 < 0$  everywhere). The asymmetry is the result of the part:  $\lambda_{\min}(A_z) = 0.32\text{--}5.08$  is well conditioned, while  $\lambda_{\min}(A_c) = 1.7 \times 10^{-5}\text{--}2.9 \times 10^{-4}$  sits three to five orders below the certified envelope margin ( $8.4 \times 10^2\text{--}4.0 \times 10^5 \times \lambda_{\min}(A_c)$ ), rising like  $\alpha^{+4.202}$ ; against  $\lambda_{\min}(A_z)$  the same margin is only 0.037–17). The near-null direction of the window form is *smooth*, so it lives in the coarse space, and its eigenvalue comes from a cancellation *inside* the form that no pointwise symbol minorant can see. The classical worst-case step  $\kappa_y \leq \gamma^2$  (Yserentant) therefore *stays* in the chain — and V7 now states why it must.

*The band margin closes; the certified balance; the  $\varepsilon_f$  lid.* (F3) reduces by the same machinery:  $y^\top G y$  splits *exactly* over dyadic depth layers of the certified gap (closure  $1.1 \times 10^{-15}$ ), and the two-factor bound [certified symbol depth]  $\times$  [Parseval mass of  $Wy$  per band] overshoots by only 1.20–3.51 — the reduction to the band profile of  $Wy$ , the same object the coupling analysis needs, costs under 1%. The certified margin is positive on 5/24 rows ( $+0.0024\text{--}+0.3290$  against the exact margin  $0.070\text{--}0.560 \sim \alpha^{-0.449}$ ); the honest lever limit:  $\min \ell = -2187 \dots - 9.7$  at  $M \leq 2944$  — the deep dips of the frame’s own resolution stay out of reach. The certified balance — the core number: the best fully certified supply is

$$S_6/\varepsilon_c \sim D^{-0.069 \pm 0.099} \alpha^{-0.544 \pm 0.113}$$

(fits;  $\alpha = 0.98\text{--}6.31$ , uniform in  $D$  and stable across the three resolutions) against the ceiling  $S_4/\varepsilon_c \sim \alpha^{-0.044 \pm 0.076}$ , so the gap is  $\alpha^{+0.500}$  — and the residual split against  $S_6$  has exactly *two* terms, closing to  $2.8 \times 10^{-17}$ : band profile  $\alpha^{+0.000}$  (F3, now negligible) plus coupling  $\alpha^{+0.499}$  (F4, all of it). The  $\varepsilon_f$  question then identifies the lid:  $\varepsilon_c = \varepsilon_f + y^\top S y$  is an identity ( $\varepsilon_f/\varepsilon_c = 0.177\text{--}0.546 \sim \alpha^{+0.030}$ , flat), and an additive chain that carries  $\varepsilon_f$  instead of discarding it would have ceiling exactly 1 — but  $\varepsilon_f$  needs its bound one level down, 4–8 recursion levels to an exact base case, and that base case is precisely (F4). The lid is not a modelling step; it is the *same object* as coarse-level positivity.

#### The theorem after T123 — four defects, restructured

*Candidate (V7).* The chain rewritten with the two-isometry completion: the Haynsworth identity; the exact CBS split  $y^\top S y = (1 - \kappa_y) y^\top A_z y$  with  $\kappa_y \leq \gamma^2$  classical (Yserentant; Axelsson–Vassilevski) — the *only* worst-case step left in the chain, and V7 states why it stays: every certified route to  $\kappa_y$  needs coarse-level positivity from below; the two-grid identities  $A_c = V^\top T V$ ,  $B_x = V^\top T W$ ,  $A_z = W^\top T W$  with  $V = B P$ ,  $W = B Z$  two orthogonal isometries ( $V^\top W = 0$ ); the certified cell envelope; the uniform oscillation positivity (5R\*); the certified band-margin bound; and the quoted demand fit.

*The defect list — four items ( $T120 = T121 = V6 = V7 = 4$ ), restructured:*

- (F1) *one sign of the corner increments* — unchanged, quoted from T120/T121; together with (F2) it is the *Harnack pair*, untouched by the band machinery;
- (F2) *a uniform bound on the pairing ratio  $\delta$*  — unchanged: 0.0126–0.1331 over 360 rows, any finite value suffices;
- (F3) *one band-profile lemma* — the band margin reduces to the band profile of  $Wy$  at a cost of 0.7%, and (5R\*) supplies the uniform certified positivity of the oscillation block for free; what remains is one analytic statement about how the mass of  $Wy$  distributes over frequency bands;
- (F4) *coarse-level positivity — identified*, not estimated: every certified route to the coupling needs  $\lambda_{\min}$  of the coarse form from below, that eigenvalue is a smooth-direction cancellation invisible to symbol minorants, and it is the same object as the  $\varepsilon_f$  lid that defines the chain’s own ceiling.

*Honest movement:* the counter stays at four, but the terrain changed: the whole  $\alpha^{+0.500}$  gap to the ceiling now sits in one named obstruction whose resolution cannot be a sharper two-level estimate. The next step is not sharper banding but the *recursion* — the level telescope  $\varepsilon_\ell = \varepsilon_{\ell+1} + y_\ell^\top S_\ell y_\ell$  with certified per-level contributions.

*Disposition after T124:* the recursion was built, and the direction flipped. Written as a telescope, the rung contribution is a residual in the *inverse* norm — a maximum, not a minimum — so it is certifiable from *above*, where the certified envelope works: (F4) collapses from a quantitative estimate to the sign

a coarse-to-fine induction already carries, the base case becomes pure semidefiniteness, and  $\lambda_{\min}(A_c)$ ,  $\text{cond}(A_c)$  and  $\gamma^2$  leave the chain entirely. The next subsection reports the measurements.

### 7.23 The direction flip: the telescope carries, and (F4) collapses to a sign (T124)

T124 (`multilevel_telescope_probe.py`, contract `MULTILEVEL.TELESCOPE`, 28/28, `TELESCOPE-CARRIES`) builds the recursion T123 demanded, and the exit turns out to be not the recursion's *length* but its *direction*: written as a telescope, the quantity each rung must supply is not a Rayleigh quotient (a minimum, needing the form from below) but a residual in the inverse norm (a maximum, needing the form from above) — and from above, the certified envelope works, because Parseval is two-sided.

*The telescope, exact and geometric.* On the level chain  $D_0 > D_1 > \dots > D_L$  (halvings,  $\alpha$  fixed) the ladder is *one window form* on nested spaces:

$$P^\top A^{(\ell+1)} P = A^{(\ell)}, \quad P^\top b^{(\ell+1)} = b^{(\ell)}$$

exactly ( $2.4 \times 10^{-14}$ ; the pole-vector consistency is the identity  $\sinh(a-d) + \sinh(a+d) = 2 \cosh d \sinh a$ ) — the coarse block of the window form at  $D_{\ell+1}$  is the window form at  $D_\ell = 2D_{\ell+1}$ , so the T122/T123 two-level system is literally one rung of the ladder. Every rung is the T118 saturation identity,  $\varepsilon_\ell = \varepsilon_{\ell+1} + \delta_\ell$  with  $\delta_\ell = y_\ell^\top S_\ell y_\ell$ , and Galerkin orthogonality (Céa) makes it the energy of the level correction,  $\delta_\ell = \|u_{\ell+1} - Pu_\ell\|_{A^{(\ell+1)}}^2$  (verified to  $3.7 \times 10^{-8}$ ) — nonnegative, hence the exact telescope  $\sum_\ell \delta_\ell = \varepsilon_0 - \varepsilon_L$  ( $9.2 \times 10^{-10}$ ). The level distribution is geometric on average: the top rung carries 0.2722–0.8806 of  $\varepsilon_0$ , consecutive rungs fall by a median factor 0.316 (quartiles 0.229–0.473, against  $2^{-\theta} = 0.289$  from the demand law), and at  $L = 5$  the ceiling of the additive chain is 0.9911–0.9987 — against T123's two-level ceiling  $\alpha^{-0.044}$ . Rung law (a fit):  $\delta \sim D^{+1.598} \alpha^{-5.900}$ .

*The certified rung contribution — the direction flip.* The rung has two dual descriptions. Primal,  $\delta_\ell = \min_{w \in V_\ell} \|u_{\ell+1} - w\|_A^2$ : a *minimum*, every test vector an upper bound — the wrong direction for a supply. Dual: the level residual  $r_\ell = t_{\ell+1} - A^{(\ell+1)} P u_\ell$  satisfies  $V^\top r_\ell = 0$  (Galerkin), so  $r_\ell = W \hat{s}_\ell$  lives purely in the oscillation space, and

$$\delta_\ell = \|r_\ell\|_{A^{-1}}^2 = \hat{s}^\top S^{-1} \hat{s} = \max_v \frac{(\hat{s}^\top v)^2}{v^\top S v}$$

— a *maximum*: every test vector gives a lower bound, and the denominator wants  $S$  from *above*. That is where the certified machinery works:  $S \preceq A_z$  needs only  $A_c \succeq 0$  (Haynsworth monotonicity),  $A_z \preceq U := Z^\top T_M(\text{up}) Z$  is Parseval against the certified majorant up  $\geq \sigma^{(M)}$  (the *mirror* of T123's  $E_z = W^\top T_M(\ell) W$ , on the block where it works), and Loewner inversion turns the upper bound on  $S$  into the lower bound on the rung,

$$(8R) \quad \delta_\ell \geq \hat{s}^\top U^{-1} \hat{s},$$

with the optimal test vector  $v^* = U^{-1} \hat{s}$  in closed form — no coarse inverse, no  $\lambda_{\min}(A_c)$ , no CBS constant. (8R) is valid on 400/400 rungs (80 chains  $\times$   $L = 5$ , 20 zones), recovering  $\text{cb}/\delta = 0.4739$ – $0.9938$  at drift  $\alpha^{-0.080 \pm 0.007}$  — seven times weaker than T123's  $\alpha^{-0.544}$ ; of the loss, the coupling accounts for 0.5036–0.9967 and the envelope only 0.7895–0.9994. T123's obstruction remains and is *irrelevant*:  $\lambda_{\min}(E_c) < 0$  on 399/400 rungs, and (8R) never forms  $E_c$ . The Levinson bordering route honestly *falls*: the innovation ledger  $q = \sum_n \rho_n^2 / E_n$  is exact per level ( $9.6 \times 10^{-12}$ ; Levinson–Durbin–Delsarte–Genin), but refinement is not a bordering — the two filtrations are transversal ( $\text{head}/q_\ell = 0.61$ – $1.00$ ,  $\text{tail}/\delta_\ell = 13$ – $863$ ) — so the rung must be read in the (8R) form.

*Self-similarity, and why the recursion closes.* The near-null direction of the window form stays *smooth* at every level (oscillation fraction  $4.3 \times 10^{-6}$ – $1.2 \times 10^{-1}$ ) — T123's diagnosis, confirmed level by level — and it is self-similar along the ladder: prolonged,  $w_\ell$  has Rayleigh quotient exactly

$\lambda_{\min}(A^{(\ell)})$  at the finer level (an identity, from nesting), i.e. within a factor 1.30–10.11 of that level’s own minimum, at subspace overlap  $\|Q_K^T P w_\ell\| = 0.944\text{--}1.000$ . But (8R) asks only for the *sign* of  $A_c$ , and by the nesting identity that sign is the previous rung’s own conclusion: positivity propagates fine  $\rightarrow$  coarse for free (so the opposite direction would be circular — stated explicitly), and the non-circular induction runs coarse  $\rightarrow$  fine, with base case  $\varepsilon_L \geq 0$  — pure semidefiniteness of the smallest window form.

*The balance; the honest negative.* Telescope reading:  $\sum_\ell \text{cb}_\ell / \varepsilon_0 = 0.6404\text{--}0.9729$  against the ceiling  $0.9911\text{--}0.9987$  (fit  $D^{-0.066} \alpha^{-0.084 \pm 0.033}$ ). Two-level comparison, like-for-like against T123’s certified balance:  $\text{cb} / \varepsilon_c \sim D^{-0.090} \alpha^{-0.100}$  against T123’s  $D^{-0.069} \alpha^{-0.544}$  — the  $\alpha$  exponent moves by  $+0.444$  of the  $\alpha^{+0.500}$  gap. The honest negative, recorded as (F5): the *solution-free* version of the rung fails — the coupling term cancels the data high-pass rather than perturbing it ( $|r_{\text{corr}}| = 0.88\text{--}2.61$  of the explicit part), so any norm bound kills the certificate to exactly 0; what would be needed is a bound on a *direction*, not a size. (8R) therefore carries the coarse solution along — on the same bookkeeping standard as the whole chain:  $\hat{s}$  is a function of the carried coarse solution exactly as  $y$  is in T122/T123’s own certified supplies.

#### The theorem after T124 — (F4) collapses to a sign

*Candidate (V8).* The telescope chain: the nesting identity  $P^T A^{(\ell+1)} P = A^{(\ell)}$  (the ladder is one window form on nested spaces); the saturation telescope  $\varepsilon_0 = \varepsilon_L + \sum_\ell \delta_\ell$  (exact); the dual residual form of the rung with the certified per-rung bound (8R)  $\delta_\ell \geq \hat{s}^T U^{-1} \hat{s}$ , needing only  $A_c \succeq 0$  (Haynsworth) and the Parseval majorant  $U \succ 0$ ; and the coarse-to-fine induction whose hypothesis at rung  $\ell$  is the conclusion of rung  $\ell - 1$ , with base case  $\varepsilon_L \geq 0$ .

*The defect list — the counter reads differently now:*

- (F1) *one sign of the corner increments* — unchanged, quoted from T120/T121; with (F2) the *Harnack pair*, untouched by telescope and band machinery alike;
- (F2) *a uniform bound on the pairing ratio  $\delta$*  — unchanged:  $0.0126\text{--}0.1331$  over 360 rows, any finite value suffices;
- (F3) *closed* — the T123 reduction to the band profile of  $Wy$  plus the uniform oscillation positivity (5R\*); the telescope no longer needs it;
- (F4) *collapsed* — no longer a quantitative estimate: (8R) needs only the sign of  $A_c$ , which the coarse-to-fine induction already carries;  $\lambda_{\min}(A_c)$ ,  $\text{cond}(A_c)$  and  $\gamma^2$  leave the chain entirely, and the base case is pure semidefiniteness;
- (F5) *new, a bookkeeping convention, honestly stated* — the certified rung carries the coarse solution along (the solution-free version provably kills to zero: the coupling is a cancellation, not a perturbation); this is the same standard the whole chain already uses, but it is a convention, not a theorem, and it is listed.

*Honest movement:* the two-level worst case is gone — the last quantitative obstruction of the certification arc is replaced by a sign that the induction itself supplies, at the price of one named convention. What remains is the Harnack pair, one constant-not-rate loss per rung, and the assembly of the full chain end to end.

## 7.24 The grand assembly (T125): the chain composes

T125 (`grand_assembly_probe.py`, contract `GRAND.ASSEMBLY`, 34/34, `ASSEMBLY-GREEN`) is the series finale, and it does exactly one thing: it mounts every certificate of the arc on one ladder, end to end, and then states what the resulting object is. Five stages per zone — [A] the base-case Cholesky, [B] the telescope ascent with the (8R) rungs, [C] the  $\varepsilon$  lower bound on the target level, [D] the  $\kappa_{\text{end}}$ /parity chain, [E] the margin-free Albert handover — complete on 52 of 52 ladder zones, plus ten deep zones on stage [E] alone up to  $n = 1331$  ( $h_{\text{new}}$  up to 1495).

*The composing run.* The construction that makes this an assembly rather than a list is ladder 1: 30 *consecutive* prime-power atoms  $n = 2 \dots 79$  on *one* common resolution  $D = \min_k g_k / (2\nu) = 3.472446 \times 10^{-3}$  (with  $\nu_{\text{eff}} = 4.0\text{--}58.4 \geq 4$ , so every zone of the run is admissible in the T105 sense). Because the resolution is common, the new window of zone  $k$  is the old window of zone  $k + 1$ :

$M_n(k) = M_o(k+1)$  on all 29 consecutive pairs as an *integer* identity, the incoming atom's triangle restricted to the old window is the exact zero matrix (bit-exact,  $\max|\text{entry}| = 0$ ), so  $X = Q_{\text{old}}$  with no truncation, and the leading  $h_o \times h_o$  block of the next zone's balanced form equals that  $X$  to  $2.9 \times 10^{-14}$ . The output of every step *is* the input of the next. This is a constructive way *around* T114's [P2] seam rather than through it: 0 of 84 consecutive log-gap ratios are dyadic — per-zone frames have no common refinement — so the run chooses one frame for all of it, and pays for that with its length (30 zones,  $n \leq 79$ , because  $D$  is set by the smallest gap inside the run). Ladder 2 keeps the per-zone frames and reaches 22 further zones to  $\alpha = 2.885$ , with the inter-zone seam declared rather than closed.

*The load-bearing finding.* Reading the assembly stage by stage, the weakest stage of the *whole* mounting is [D] at WINDOW-CERT on every zone — but [D] is the second, independent route. The *load-bearing spine* [A][B][C][E] is at CHOLESKY-CERT on all 52 zones, its weakest stage being [A]. In other words: **the Harnack pair (F1)/(F2) no longer carries anything**. It survives only inside conclusion (iv) of the theorem below, which (i)–(iii) and (v) do not use.

*The numbers, stage by stage.* The four identities the assembly stands on hold on all 88 rungs of all 52 zones (nesting to  $3.0 \times 10^{-14}$  and  $6.2 \times 10^{-16}$ ; the saturation/dual-residual/Galerkin trio to  $4.7 \times 10^{-7}$ ; the exact piecewise-constant lag formula to  $2.2 \times 10^{-14}$ ; the telescope  $\sum_\ell \delta_\ell = \varepsilon_0 - \varepsilon_L$  to  $3.3 \times 10^{-7}$ ). (8R) is certified on 88/88 rungs by its three completed Choleskys, at retention  $\text{cb}/\delta = 0.9071\text{--}0.9971$  — T124's band 0.4739–0.9938, reproduced on a *different* ladder (frame windows, not the T124  $M$ -chain) — and the direction control holds: the primal minimum form returns values *above*  $\delta$  on 88/88 rungs, i.e. it is an upper bound and can never be a supply. The certified  $\varepsilon$  lower bound lands on 52/52 zones, retaining 0.0013–0.9694 of the measured  $\varepsilon_0$  but 0.9140–0.9929 of the  $L$ -level *ceiling*  $1 - \varepsilon_L/\varepsilon_0 = 0.001295\text{--}0.993502$  that an  $L$ -level chain can supply at all, and it exceeds the *declared* floating-point bound  $c_h u \text{cond}(A_0)$  by a factor  $3.2 \times 10^1\text{--}8.7 \times 10^{10}$ . The base case is an *equivalence*, not a bridge:  $\text{sign } \lambda_{\min}(Q^{(L)}) = \text{sign } \varepsilon$  on 52/52 zones at  $h = 416\text{--}1384$ , with  $\lambda_{\min}(Q^{(L)})$  sitting  $4.4 \times 10^5\text{--}5.6 \times 10^{10}$  times its own floating-point floor. The margin-free handover certifies 62/62 steps ( $n = 2 \dots 1331$ ) with the exact  $g_c \times g_c$  Schur complement at  $\lambda_{\min} = 1.484 \times 10^{-3}\text{--}1.253 \times 10^{-1}$ , a factor  $3.0 \times 10^{10}\text{--}4.6 \times 10^{13}$  above its own Cholesky floor — while the norm-perturbation surrogate of the *same* Schur complement, the one that divides by  $\lambda_{\min}(X)$  and *was* the T113 wall, is negative on 62/62 steps. In total: 430 completed Cholesky certificates and 440 verified identities.

*Three gates corrected in the open.* (a) The  $\kappa$  band: the unconditional pairing certificate (C1) dominates the measurement on 52/52 rows, but its worst case over the ladder is  $C_1 = 0.20927$  — and that worst case is a *single* zone, the shallowest one ( $n = 2$ ,  $\alpha = 0.3507$ , only  $n_C = 12$  corner cells, too short for the pairing to average); drop that row and  $C_1 \leq 0.06721$  on the remaining 51, against T120's 0.04745 measured on deep zones only. (b) The head-room minimum is 31.6, at the thinnest zone ( $n = 243$ ,  $\alpha = 2.7492$ ,  $h_0 = 679$ ), and it is thin because that zone's retention 0.6054 is itself pinned to its small  $L$ -level ceiling 0.629042 — not because the arithmetic is marginal; the factor stands over a *rigorous* bound, so any value above 1 certifies the sign. (c) The balance is fitted at *rung* level rather than on the collinear zone fit:  $\text{cb}/\delta \sim D^{-0.010} \alpha^{+0.001}$  (bands  $\pm 0.004$ ,  $\pm 0.005$ , rms 0.012) — flat in *both* parameters, the strongest form of T124's quoted  $\alpha^{-0.100}$  drift — and at zone level on the composing run, where  $D$  is common and only  $\alpha$  moves,  $\sum \text{cb} / \sum \delta \sim \alpha^{-0.012}$ : the certificate does not thin out along the run. The ratio reading  $\text{cb}/\varepsilon_c$  comes out at  $D^{+0.345} \alpha^{+1.566}$  here and is explicitly *not* comparable with T124's  $D^{-0.090} \alpha^{-0.100}$  (a different denominator: a 2–5-level frame chain against  $M = 32\text{--}2560$  pairs) — reported, not claimed.

*Five seams, named and measured.* (S1) the inter-zone frame seam — *closed* on ladder 1 by the common-frame construction; (S2) the direction of the  $\varepsilon$  bound — structural and declared: the certified bound lands on the *coarse* level, so the target window must be the coarsest level of its chain and the base case the finest (a bound at the fine level would need certified *upper* bounds on the rungs, i.e. the CBS constant T123 proved resists); (S3) the matrix cap — cost, not a defect; (S4) the (F5) accounting standard — open, declared inside the theorem, with the solution-free variant measured to return exactly 0 of the rung ( $|r_{\text{corr}}| = 0.3888\text{--}14.8018$ : a cancellation, not a

perturbation); (S5) finiteness — open, *and it is the gap*.

### Theorem V-final — the certified multilevel margin chain (T125)

*Setting.* Fix  $\nu \geq 4$ ; let  $n_i < \dots < n_j$  be consecutive prime-power atoms with  $u_k = \log n_k$  and log-gaps  $g_k$ , and put  $D = \min_k g_k / (2\nu)$  (one common resolution),  $M_k$  the smallest even integer with  $M_k D > u_k + D$ ,  $\alpha_k = M_k D / 2$ ,  $A^{(k)}$  the reflection-odd window form on  $h_k = M_k / 2$  coordinates built from the exact Weil window lags,  $b^{(k)}$  the odd pole vector,  $\varepsilon_k = 1 - b^\top A^{-1} b$  and  $Q^{(k)} = A^{(k)} - b^{(k)} b^{(k)\top}$ ; for each  $k$  let  $D = D_0 > \dots > D_{L_k}$  be halvings at fixed  $\alpha_k$ .

*Hypotheses.* **(H-base)**  $\varepsilon_{L_k}^{(k)} \geq 0$  at the finest level, equivalently  $Q^{(k, L_k)} \succeq 0$  (Albert 1969) — pure semidefiniteness, no rate, no margin. **(C-F5)** an accounting convention, *declared open*: the rung numerator  $\hat{s}_\ell = Z^\top b^{(k, \ell+1)} - B_x^\top u^{(k, \ell)}$  is a function of the *carried* coarse solution, the standard T122/T123 already used; a solution-free rung is *not* claimed — it is measured to fail. **(H-F1)/(H-F2)** the Harnack pair, needed *only* for conclusion (iv).

*Conclusions.* (i) [IDENTITY]  $P^\top A^{(k, \ell+1)} P = A^{(k, \ell)}$ ,  $P^\top b^{(k, \ell+1)} = b^{(k, \ell)}$ , hence  $\varepsilon_0^{(k)} = \varepsilon_{L_k}^{(k)} + \sum_\ell \delta_\ell$  exactly. (ii) [CERTIFIED, given  $A^{(k, \ell)} \succeq 0$ ] (8R)  $\delta_\ell \geq \hat{s}_\ell^\top U_\ell^{-1} \hat{s}_\ell =: \text{cb}_\ell > 0$ ; the rung is a *maximum*, so its denominator wants the form from above, and  $v^* = U_\ell^{-1} \hat{s}_\ell$  is optimal in closed form — no  $A_c^{-1}$ , no  $\lambda_{\min}(A_c)$ , no CBS constant. (iii) [CERTIFIED, given (H-base) and (C-F5)]  $\varepsilon_0^{(k)} \geq \sum_\ell \text{cb}_\ell > 0$ , i.e.  $Q^{(k)} \succ 0$ . (iv) [WINDOW-CERT, second and *independent* route] the  $\kappa_{\text{end}}$  chain, with  $|R - 1| \leq 0.20927$ , hence  $R \leq 1.20927$  and  $\kappa_{\text{end}} \geq 0.45264$  on this run. (v) [CERTIFIED, margin-free]  $X = Q_{\text{old}}$  exactly, and  $Q^{(k+1)} \succeq 0$  follows from Albert 1969 plus the Douglas range condition — nothing refers to the *size* of  $\varepsilon_k$ , only to its *sign*, which is (iii). (vi) [THE CHAIN] therefore (i)–(iii) and (v) compose:  $Q^{(i)} \succeq 0 \Rightarrow \dots \Rightarrow Q^{(j)} \succeq 0$ , each step a completed Cholesky above its declared floating-point floor; (iv) is not used.

*Attribution* (25 lines). 10 IDENTITY, 9 CHOLESKY-CERT, 3 WINDOW-CERT (all three inside (iv)), 1 HYPOTHESIS ((C-F5)); every eigenvalue, condition number and retention is labelled MEASURED, every exponent law a FIT.

*Verified instance.*  $D = 3.472446 \times 10^{-3}$ , 30 consecutive zones  $n = 2 \dots 79$ ,  $\alpha = 0.3507\text{--}2.1494$ ,  $h_k = 101\text{--}619$ , 88 telescope rungs, plus 22 non-composed reach zones to  $\alpha = 2.885$  and 10 handover-only zones to  $n = 1331$ ; all 430 Cholesky certificates completed.

*What is not claimed.* (a) No statement about all zones — the run is *finite* and its length is limited by  $\min_k g_k = 0.027780$ . (b) (H-base) is a Cholesky fact about *one* finite matrix, not a theorem. (c) (H-F1)/(H-F2) are window certificates, measured and bounded on the measured windows, with no uniform proof. (d) (C-F5) is a convention, not a solution-free bound. (e) No zeta function, no zero data, no explicit-formula positivity, *in either direction*: this is a statement about a given quadratic form at a given resolution.

## 8 The kill list

A program that only reports what survived is not reporting. These are the readings and routes that were refuted *by the program's own probes*, each with the measurement that killed it. They are listed as a section rather than as footnotes because they are the more transferable output: each one tells a later attempt where not to go.

- K1. The essential-singularity reading of the onset** (T96, refuted T102). T96 read the counterfactual onset as super-polynomial, compatible with  $\exp(-c/\delta')$ , i.e. an essential singularity. T102 showed the onset is *anchored* at the crossing  $\delta_c = 2w_k$  of two finite quantities ( $R^2 = 0.968$ ), that exponential versus power behaviour is not separable at this resolution, and that the scatter of  $c$  (CV 103%) collapses to  $c' = 1.52$  in the scale-free variable. The essential-singularity reading is *compatible but not distinguished* — and it is no longer used.
- K2. The  $C/g$  law as a causal law** (T101, refuted T102). Triple negative:  $g_k$  is causally impossible ( $Q_{k-1}$  is blind to atom  $k + 1$ ), statistically dispensable (causal law  $C \cdot u$ : CV 23.3% versus 27.6%), and arithmetically only a ceiling ( $C_k \leq g_k \mu_k$  exactly; the 16-zone extrapolation violates its own ceiling from  $k = 69$ ; checked on 18 120 prime-power atoms to  $n = 200\,000$ ). *It was a proxy.*

- K3. The naive margin route** (T104).  $M \geq m \cdot \text{Id}$  holds in 0/16 zones; the coupling-mass to margin ratio is  $10^3\text{--}10^6$ ;  $m \sim M^{-1.7}$  vanishes in the continuum (fit). The genuine form is nearly singular on the induction window and no margin rescues it. (The prolate wing basis recommended by T96/T103 was tested in both arms and also rejected.)
- K4. The density chain** (T106). It would need a hard gap  $w_0 \geq 0.55\text{--}2.44$ ; measured  $5.3 \times 10^{-6}\text{--}1.1 \times 10^{-3}$  — short by a factor  $5 \times 10^2$  to  $3.6 \times 10^5$ . Dead for good.
- K5. The Landau–Widom anchor** (T106). The wrong classical anchor for this problem (rms 0.50–1.55). The right one is Grenander–Szegő on the Planck-coarsened symbol — averaging over the mode-spacing cell  $\pi/M$  — which hits 0.92–2.43 and is better on 67/72 pairs. This also explains why a negative symbol (dip  $-21.3 \dots -3.4$ ) still gives  $Q \geq 0$ : the atoms oscillate exactly at the window’s Planck scale.
- K6. The accumulated invariant with self-propagation** (T106). A no-go *with a measured mechanism*. The wing demand gives  $\beta_{\text{wing}} = 2.21\text{--}14.43 > 1$  on 16/16, but inherited demands give  $\beta_{\text{int}} \sim 10^{-2}$  on 15/15: an accumulated invariant with absolute anchoring is false. The reason is transport — in a larger window an old wing pair becomes an interior pair and draws 99.7–99.95% of its demand from the 64 softest modes (coupling 751–2653× stronger). No self-propagation either ( $\beta_{\text{on}} < 1$  on 15/15, loss 13–320×,  $\theta_k = 0.9866\text{--}0.9998$ ). *The wing margin is an edge property of the window, not a property of the pair*. What survives is the window-wise family  $\sigma_k(\delta_0) \geq 1.31(\mu_k/2)$ , uniform in  $M$  with drift  $\leq 2.01\times$ .
- K7. The symbol route for the odd channel** (T107). Structurally dead, not merely weak: the symbol lower bound for  $T_{\text{odd}}$  is certified but *negative* ( $\lambda_{\text{cert}} = -2.54 \dots -0.81$ ) against a measured  $\lambda_{\text{min}}(T_{\text{odd}}) = 4.7 \times 10^{-5} \dots 6.4 \times 10^{-2}$  and a symbol infimum of  $-22 \dots -6$ : five orders of magnitude and a sign change. The soft edge of  $T_{\text{odd}}$  is a pure finite-section effect that no symbol can see. Coarse certified chains (Cauchy–Schwarz against  $\lambda_{\text{min}}$ ) close 0/16 — and fail even with the *measured*  $\lambda_{\text{min}}$ , so it is a category error, not a precision problem.
- K8. Three certificate candidates for the avoidance norm** (T108). Killed *quantitatively*, which is the useful form: the cancellation route ( $\sum |k| = 401$  against  $\sum k = 1$ ), the Bessel route (off by  $10^7$ ) and the pole Cauchy–Schwarz route (off by  $10^6$ ).  $\omega_{\text{cert}}$  is blocked by the compression Schur distance alone (0.18–0.52). What replaced them is not a fourth bound of the same kind but a different object class: a boundary-decay statement about one explicit vector. T109 then killed the natural fourth candidate as well: **the Combes–Thomas resolvent bound at the boundary**, evaluated with the *exact* Green row so that no constant can rescue it, closes 1/16 — the mechanism at the edge is cancellation ( $44\text{--}4.6 \times 10^4$  on 15/16 zones), which no absolute-value bound survives. What works instead is the residual certificate that carries the cancellation (Section 7).
- K9. The scalar step law for the margin** (T110). Killed structurally, in three pieces: bordered Weyl closes 0/15 handovers; the Schur/Friedrichs scalar cap is vacuous on 15/15 (its scalar-only bound degrades back to Weyl); and the strictly scalar graded minorant ( $n_{\text{soft}} = \text{dim}$ ) closes 0/15, with the strictly scalar *induction* compounding to a Schur factor of  $2.7 \times 10^{-36}$  by  $k = 3$ . The reason is a boundary layer: the new cells of a grown window attach at the edge where the pole source sits, so a one-number summary of the bordering is blind to the deciding direction. What survives is the graded ( $n_{\text{soft}} = 1$ ) Loewner minorant — one soft direction is exactly the amount of geometry the step needs.
- K10. Two earlier structural kills**, recorded here because they still bound the design space: the cross-block  $\sqrt{\lambda_0}$  law is Douglas range inclusion and therefore circular (T98), and the cone-library route saturates at 5/24 coverage and is dead as a covering strategy (T72, promoted into v540 as a named limit). Likewise the seam route to the archimedean place (T20–T25) and the Kohnen-plus route to the central values (T44, promoted into v537 as a scope fence).

The finale re-counted this list against the whole certification arc and made it quantitative: 24

**refuted routes across 18 parts**, every one of them structural — an identity (the T123 ceiling, the T124 primal direction), an exact sign ( $\lambda_{\min}(E_c) < 0$ , the negative norm surrogate, the non-dyadic gap ratios), or a measurement that tuning cannot move (the vacuous band split, the empty energy route, the cancelling solution-free rung). The counting matters because *four* of the 24 died of the *same* cause — a bound that needs the coarse form from *below* — which is why the exit was a direction reversal and not a sharper estimate. T125 added one entry of its own: a certified  $\varepsilon$  bound at the *fine* level of a chain is the CBS wall again, which is why V-final mounts the target window as the *coarsest* level.

## 9 The series balance (T102–T125)

The finale’s second half is bookkeeping about the series itself: what the hard core was at each stage, how much of the final chain is a certificate rather than a hypothesis, and how far that chain is from any infinite statement. All three are measurements against the same source, and they are reproduced here because they are the part a later reader can check without re-running anything.

### 9.1 The reduction cascade, seven stations

Each arrow is a strict reduction in dimension or in kind, and every one of them was paid for by a refutation in the list above:

1. **Matrix wall** (T104–T106) — a Loewner statement  $M \succeq m \text{Id}$  on  $M/2$  dimensions, killed as a margin route and localised by the parity split to one statement (R) in the odd channel.
2. **Scalar** (T107–T108) — Sherman–Morrison/Woodbury reduce (R) exactly to one scalar ratio  $r = \kappa/\varepsilon \leq 1$ , and  $\varepsilon$  itself becomes an identity (the square of the last Cholesky pivot).
3. **Ratio** (T108–T109) — the ratio splits into  $\omega < 1$ , certified unconditionally by a graded matrix cap, and one statement at an explicit vector.
4. **Boundary value** (T109) — what is left is literally one boundary value, carried by a residue certificate that *keeps* the cancellation instead of estimating it away.
5. **One inequality** (T110–T117) — the frame removes the ladder and  $\omega$  walls, the margin wall turns out to measure the discretisation, the margin-free Albert step removes the division that was the wall; what remains is one lower bound on  $\varepsilon$ .
6. **Harnack pair + one sign** (T118–T124) — the  $\varepsilon$  bound becomes an identity plus a supply, and the supply needs either the corner route ((F1)+(F2)) or the telescope route ((8R), whose only requirement is the *sign* of  $A_c$ ).
7. **One sign, one convention** (T125) — assembled: the spine  $[A][B][C][E]$  needs (H-base), a sign on *one* finite matrix, plus the (C-F5) accounting convention. **The Harnack pair survives as the second, independent route (iv) and is no longer a bottleneck of the spine** — this last station is what T125 added.

### 9.2 The certification degree

Every link of V-final counted once, classified by what actually establishes it:

| Class         | Links | %     | Spine | %     |
|---------------|-------|-------|-------|-------|
| IDENTITY      | 14    | 45.2  | 11    | 42.3  |
| CHOLESKY-CERT | 14    | 45.2  | 14    | 53.8  |
| WINDOW-CERT   | 2     | 6.5   | 0     | 0.0   |
| HYPOTHESIS    | 1     | 3.2   | 1     | 3.8   |
| <b>Total</b>  | 31    | 100.0 | 26    | 100.0 |

90.3% of the whole chain is identity or completed Cholesky; on the load-bearing spine it is **96.2%**, with *zero* window certificates and exactly one hypothesis — the (C-F5) accounting convention. The three window certificates are confined to conclusion (iv).

### 9.3 The honest RH distance, in six checkable statements

- (R1) **No zeta input.** The chain consumes exactly two arithmetic facts, both classical and both only for *admissibility* of the frame:  $g_k \leq \log 2$  (Bertrand–Chebyshev 1852) and  $g_k \geq \log(1 + 1/n)$ . No zeta function, no zero data, no zero-derived quantity — the probe’s own AST firewall rejects the tokens, the imports and write-mode file access in its source.
- (R2) **RH is not used, in either direction.** “RH  $\Rightarrow$  window Weil positivity” appears nowhere; (H-base) is a completed Cholesky of one explicit matrix of size  $h = 416\text{--}1384$ .
- (R3) **No infinite statement.** V-final quantifies over a *finite* run of consecutive atoms whose length is bounded by  $\min_k g_k = 0.027780$ : here 30 zones,  $n = 2 \dots 79$ .
- (R4) **A measurement ladder is not all zones.** 62 of 85 admissible handovers in the table, and 48 of 34 034 prime-power atoms up to 400 000 touched at all. Everything outside is untested, not implied.
- (R5) **The open items are uniformity, not size.** What is missing for any infinite statement is a uniform-in- $k$  version of stage [E]’s  $g_c \times g_c$  Schur complement, a uniform lower bound on  $\text{cb}_\ell$ , a proof rather than a measurement of (H-F1)/(H-F2) — classical addresses named (Widom 1974 / Fisher–Hartwig for the corner regularity of the inverse, Ostrowski 1937 / Varga 1962 for the sign) — and removal of the (C-F5) convention, which T124 measured to be a cancellation.
- (R6) **What it is.** A machine-certified, finite, multilevel *lower* bound on the Galerkin prediction defect of a log-singular Toeplitz form with a prime-power comb, together with a margin-free bordering induction that transports its sign along a nested window chain. That is a statement in numerical analysis. It is not a statement about zeros.

The defect counter closes accordingly: (F1)/(F2) window-certified and now *outside* the spine, (F3) closed in T123, (F4) collapsed to a sign in T124, (F5) declared as (C-F5) — and one new entry, **(F6) uniformity**, the only thing between V-final and an infinite statement.

## 10 Phase 2: toward uniformity (T126–T128)

The series closed at T125 with a mandate fulfilled and a distance named: what separates V-final from any infinite statement is *uniformity in the zone index*, not size. The first post-series part — T126 (`uniformity_seams_probe.py`, contract UNIFORMITY.SEAMS, 31/31, verdict SEAMS-CERTIFIED) — opens a second phase of the diary by attacking exactly that gap in its three components: the *segment seams* (consecutive runs carry different common resolutions, and the seam between them is one re-gridding), the *uniformity of the spine constants* (the (F6) entry of V-final, audited one constant at a time), and the *continuum argument* (finite-section positivity at every resolution  $\Rightarrow$  positivity of the continuum window form on a test class). All three move; what remains at the end is stated in Section 10.4 and is deliberately short.

### 10.1 The seam architecture: the partition question dissolves

The natural first question is arithmetic: can the prime-power axis be cut into segments such that each segment carries one common resolution  $D_{\text{seg}} \propto \text{min-gap}$  and neighbouring segments have a resolution ratio the certified transport reaches? Since  $D_{\text{seg}}$  is set by the segment’s minimum log-gap, this is a question about *minima of windows of the log-gap sequence* and nothing else — so it is decided *exactly*, by a dynamic programme over cut positions on the real gap table (9 700 zones,  $n \leq 10^5$ ; 33 DP runs in 2.71 s). The answer is a refutation with a replacement:

- **At T115’s synthetic threshold  $\rho^* = 1.83$  no partition exists.** The DP dies at zone  $n = 127$ , at the local gap pattern  $[0.0684, 0.0325, 0.0159, 0.0078, 0.0232, 0.0448]$  — a small gap directly behind a large one with no intermediate gap in the admissible window: the twin-after-a-large-gap mechanism, named exactly. The true requirement is  $\rho_{\text{crit}} = 2.0258$ , identical on  $n \leq 10^4$  and  $n \leq 10^5$  — not an aggregate but essentially *one* quotient of neighbouring gaps,  $g(125)/g(127) = 2.0238$ , reproducing  $\rho_{\text{crit}}$  to  $2.0 \times 10^{-3}$ . A partition at  $\rho_{\text{crit}}$  is *constructed* (306 segments; the T112 frame lemmas hold on all 9 700 zone steps,  $\min \nu_{\text{eff}} = 4.00$ , zero violations).
- **A refinement-only partition exists at no ratio  $\leq 64$ ,** so the tight-cap partition forces *coarsening* seams — and coarsening fails for real: on real pairs the transported certificate fails at ratios 0.309 and 0.531 while the target floor itself stays positive on 9/9 (what fails is the certificate, not the positivity), and ladders of intermediate resolutions make it *worse* ( $k = 1, 2, 3$  at  $n = 47$ :  $-2.54 \times 10^{-2} \rightarrow -2.52 \times 10^{-2} \rightarrow -4.63 \times 10^{-2}$ ) — the obstruction is the *destination grid* on which every ladder must land; a finer subdivision cannot cure a bad target.
- **The decisive new measurement:** holding the worst-case ratio 2.0256 *fixed* and sweeping the zone, the transported certificate survives on only 12 of 22 zones. The seam dies at the *zone*, not at the ratio — so no ratio threshold can ever close the programme, and the lemma the seam architecture needs is not “a larger  $\rho$ ” but a *zone-uniform*  $\eta/\tau$  bound at ratio  $\leq 2.026$  (U5 below, sharpened).

**What carries is the free-resolution route.** Releasing the resolution from the per-segment cap and taking the running minimum  $D_k = \min(\text{cap}_k, D_{k-1})$  makes monotonicity hold *by construction* — no DP, no gap theorem, not one coarsening seam anywhere. The largest single refinement drop over  $n \leq 10^5$  is 2.0256 (zone  $n = 8191$  — the same binding quotient, now read in the good direction), and re-gridding happens at only **764** of 9 700 boundaries (7.88%): a drop needs a *record-small* gap, and every other boundary is a no-op at ratio 1. All 12 affordable actual seams of that route ( $n = 3 \dots 125$ ,  $h \leq 1222$ , largest real ratio 1.943) are built and transported: 11 carry a positive certified floor in *one* re-gridding, the one at  $n = 31$  (transported mass  $\tau = 0.363$ , unusually poor) closes through a three-stage ladder with composed floor  $3.56 \times 10^{-2} > 0$ ; median retention 0.578, all 12 brackets valid. The seam device itself is sound throughout: the Haynsworth partial minimisation reproduces  $\lambda_{\min}(S)$  to  $2.8 \times 10^{-14}$  on both grids and the transport bracket contains the fine value on 30/30 real pairs; and a *full* seam is certified end to end at the ratio the arithmetic demands — spine on grid A, one re-gridding (lower bracket end  $8.671 \times 10^{-3} > 0$ , retention 0.103), spine on grid B — at  $\rho = 2.251 \geq 2.0256$  on 3 of 4 attempts (the failures are the zone dependence above, printed and not hidden). The price of the route is honest and structural: the resolution overhead  $D_{\text{free}}/\text{cap}$  has median 0.201, i.e. the *running minimum* of the gaps sets the window cost instead of the local gap — window size, not an assumption; certified seam by seam the route covers  $n \leq 125$ , an evidence limit rather than a gap in the argument.

## 10.2 The uniformity audit (U1–U6)

Every constant of the load-bearing spine is measured over zones *and* telescope depth (35 rungs, 14 zones, (8R) valid on 35/35 while audited; every fit jackknifed; the status rule declared before the numbers, and “measured-flat” is never upgraded to “uniform”):

| Constant                    | Finding                                                                                                                                                                                                                 | Status              |
|-----------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|
| U1 envelope margin          | forced by the construction's own doubling loop                                                                                                                                                                          | self-certifying     |
| U2 oversampling depth $L/M$ | needs one $\alpha$ -uniform Bernstein/Markov bound                                                                                                                                                                      | classical in form   |
| U3 rung quality $cb/\delta$ | measured-flat, but the reason is <i>direction</i> , not spectrum: the worst pencil direction of $(S, U)$ does not explain the flatness (alignment 0.866–0.987, median 0.975; $\mu_{\min}$ itself falls with the window) | genuinely new       |
| U4 base-case head-room      | a compute wall over the declared fp floor                                                                                                                                                                               | no proof obligation |
| U5 Albert–Schur floor       | $\lambda_{\min}(S)/\ A\ $ <i>falls</i> in the window direction — the one load-bearing gap; sharpened by the seam scan to a zone-uniform $\eta/\tau$ bound at ratio $\leq 2.026$                                         | genuinely new       |
| U6 resolution monotonicity  | the running minimum is monotone by construction                                                                                                                                                                         | resolved            |

The audit's own finding is that (F6) is not one lemma but three — U2 (classical in form), U3 (a direction lemma on  $\hat{s}$ ), U5 (a zone-uniform Schur floor) — of which exactly two are genuinely new mathematics.

### 10.3 The continuum step is a fractional Dirichlet energy

The chain one would write down first — density of piecewise-constant functions plus a continuity constant — is *refuted*: the archimedean kernel is Cauchy-type  $\sim 1/(2|s|)$ , not log-type, hence not locally integrable; its  $L^1$  mass over the window and  $\lambda_{\max}$  both diverge like  $\log(1/D)$ , so neither an  $L^2$  nor a sup-norm continuity constant exists (measured). What makes the divergence harmless is an *identity*: with row sums  $w_r$ ,

$$v^\top A v = \sum_r w_r v_r^2 + \sum_{r < s} (-A_{rs}) (v_r - v_s)^2,$$

with 86.1% of the off-diagonal entries measurably negative and the row sums bounded — the window form *is* a fractional Dirichlet energy, and the divergent diagonal is the  $\frac{1}{2}$ -energy of a step function, a property of the test class rather than a defect of the kernel. The limit argument is therefore quantitative, and it is measured: on the dyadic tower (5 levels,  $h = 63 \dots 1008$ , a  $16 \times$  spread in  $D$ ) the Galerkin values converge at order  $D^{1.798 \pm 0.057}$  for a  $C^\infty$  bump and  $D^{1.802 \pm 0.065}$  for a Lipschitz hat — superlinear in both cases, exactly as the energy identity predicts — with every level strictly positive and Richardson limit  $2.1256 \times 10^{-4} > 0$ . What is classical here is the consistency estimate for a Cauchy-kernel Galerkin scheme (Calderón–Zygmund; Bramble–Hilbert 1970); what is measured is its rate on this kernel; the abstract density step is not needed and not used. The statement that would stand at the end is **positivity of the Weil window form on Lipschitz test functions supported in the reached window** ( $\alpha = 1.5046$  on the tower, 1.6882 along the seam chain) — a finite-window statement whose bridge to RH is Weil's criterion (Weil 1952; Bombieri 2000), cited here and used in neither direction.

### 10.4 The full-proof map: sixteen points, two genuinely new

The definitive fail list after the three attacks has 16 points: 5 closed or classical, 9 conditional (constructed here, classical in form, or a machine limit), and exactly *two* genuinely new pieces of mathematics:

#### The two remaining inequalities (T126)

**(U3) The direction lemma.** Why the rung's data residual  $\hat{s}$  sits in the good subspace of the pencil  $(S, U)$ : the measured flatness of  $cb/\delta$  is provably *not* the pencil's worst-case behaviour (alignment 0.975 while  $\mu_{\min}$  falls), so a proof must control the *direction* of  $\hat{s}$ , not the spectrum.

**(U5) The zone-uniform seam floor.** A zone-uniform  $\eta/\tau$  bound for the re-gridding transport at ratio  $\leq 2.026$  — the constant every seam consumes. Ladders provably cannot repair a bad destination grid, so the reach must be won per *zone*, not per ratio.

S4, the coarsening bracket — which this part opened and which would have been a third genuinely new statement — is *dissolved* by the free-resolution route rather than proved. Still open and unchanged: the (C-F5) accounting convention, and the all- $\alpha$  step — the one thing no computation can close (at the tight cap, 6 boundaries in 9700 would need a double seam; for all  $n$  one needs a gap theorem, or the free-resolution relaxation, which is unconditional in the arithmetic). T127 (`two_inequalities_probe.py`, contract TWO.INEQUALITIES, 28/28, verdict BOTH-SHAPED) then took both inequalities head-on. The result — both come back proof-shaped, and *neither is the statement T126 named* — is the remainder of this section.

## 10.5 The separating variable is geometric, not arithmetic (T127)

The fixed-ratio sweep reproduces exactly: at the worst-case ratio 2.0256 the transported certificate survives on 12 of 22 zones, T126's list — and what dies is the *certificate*, not the positivity: the true target eigenvalue  $\lambda_{\min}(S_f)$  is positive on all 22 zones. Of 29 candidate separating variables, exactly 5 separate the surviving from the dying zones perfectly (22/22; the rank statistic is a description of a 22-point sample and is never called a test), and all five measure the same *geometric* quantity: the right-hand protrusion  $s_{l,r}$  — the two grids quantise the same window independently, their border blocks have equal width but sit offset by one rounding, and the transported mass  $\tau_{\text{dn}}$  loses exactly what protrudes past the coarse block's inner edge (correlation +0.997 with  $\tau_{\text{dn}}$ ). Every arithmetic candidate fails (gap,  $\alpha$ , atom weight, alias phase; best separation 0.727) — the separation is *not* the atom and *not* the gap. Deep arithmetic ( $n \leq 300\,000$ , 1082 record seams): the fine border block has 4 cells everywhere, 97.60% of real seams have ratio  $\leq 1.05$  with median 1.00140 — real seams are *near-identity re-griddings* — and the 11 seams with ratio  $> 1.2$  *all* have next-atom distance  $\leq 2$ : Mersenne (3, 7, 31, 127, 8191), Fermat-like (4, 16, 256, 65536, 131071), one twin (101, 103). Mihăilescu 2004 (Catalan) is cited as the classical address of that list and not used.

## 10.6 U5 as posed is refuted; the U5 band replaces it

As a zone-uniform certificate at ratio  $\leq 2.026$ , U5 is *dead as posed*. The replacement is **U5-band**: all 22 zones survive for every ratio  $\rho \leq 1.49531$ , and at least one dies for every  $\rho \geq 1.49766$  (bisected to  $2.3 \times 10^{-3}$ ; binding zone  $n = 19$ ) — a band that covers 99.26% of the route's real seams, the remainder being the 8 enumerated pairs above. The split of the form defect is an *identity* ( $\eta = \eta_{\text{cons}} + \eta_{\text{proj}}$ , exact), and  $\eta$  does *not* become small as  $\rho \rightarrow 1$  (max  $|\eta|/\lambda = 1.146$  at the *tightest* ratio): the band buys  $\tau_{\text{dn}} \approx 1$ , not small  $\eta$ . The floor is measured at  $c \geq 0.0090$  over 307 in-band transports with a positive single-step floor up to  $n = 1331$ ,  $h = 1496$ . **The sharpest new finding**: the near-identity census (43 deep zones, all at ratio 1.00140) fails *once* —  $n = 479$ ,  $|\eta|/\lambda = 1.472$ , within 0.14% of the identity — so the ratio band is only a *proxy* for the border geometry; the failure is repaired by a three-stage ladder (retention 0.5027), exactly as T126 repaired  $n = 31$  — the fallback is part of the statement. And the risk is recognisable *before* any factorisation: the indicator  $\tau_{\text{dn}}/\tau_{\text{flat}}$  ranks the margin at Spearman +0.749, and the one failure is its argmax ( $\tau_{\text{flat}}$  is a predictor, not a bound).

## 10.7 U3 defuses to a coarse floor: the harmonic-mean identity

Three identities restructure U3. The pole part of the rung data has the closed form  $(Z^{\text{T}}b)_j = -\frac{16}{\sqrt{2D}} \sinh^2(D/4) \cosh(\tilde{x}_j^c/2)$  (to  $3.2 \times 10^{-13}$ ); the rung quality is a *harmonic mean*,  $\text{cb}/\delta = 1/\sum_i w_i/\mu_i$  (to  $1.5 \times 10^{-7}$ ); and the weights split by Parseval (to  $6.1 \times 10^{-16}$ ). The residual's mass on the good half of the pencil spectrum  $\{\mu \geq \frac{1}{2}\}$  is 0.9665...0.99998 over 42 rungs (flat), so  $\text{cb}/\delta \geq 1/(2m + (1-m)/\mu_{\min})$  — a floor as *crude* as  $\mu_{\min} \geq 1.62 \times 10^{-2}$  (one part in 62) already gives  $\text{cb}/\delta \geq \frac{1}{4}$ , where T126 had to measure 0.192–0.307 sharply. The reading “ $\hat{s}$  avoids the bad direction” is *refuted*: the flatness is generic — every smooth, flat, alternating or quasi-random direction has good-band mass above 0.50 — and exactly *one* direction type fails, a *boundary layer* at the window edge (concentration  $3.1\times$  there, against  $0.03\times$  for  $\hat{s}$  and  $0.001\times$  for its closed-form

pole part). The bad pencil directions are *not* edge-localised in the  $\ell_2$  sense ( $1.23\times$ ) — only the  $U$ -metric concentration separates.

## 10.8 The map V2: three inequalities and one enumeration

The map V2 has 16 points: 10 carried, 2 conditional-classical (one Bernstein/Markov bound, the declared fp floor), and the 2 genuinely new ones of T126 *reformulated*:

### The remaining mathematics (T127, map V2)

**(T) The border-retention bound.** A lower bound for the transported mass  $\tau_{\text{dn}}$  under the edge-geometry hypothesis — from predictor to bound, via the support of the  $S_f$  eigenvector; the ratio band  $\rho \leq 1.50$  is only a proxy for it.

**(E) The  $|\eta|$  upper bound.** Both parts of the form defect  $\eta = \eta_{\text{cons}} + \eta_{\text{proj}}$ , to be bounded *jointly* with (T): the two rank together at  $+0.749$ .

**(M) The pencil-mass bound.** The residual's good-band mass plus the crude floor  $\mu_{\text{min}} \geq 1.62 \times 10^{-2}$ ; the one direction type to exclude is named (a boundary layer at the window edge), and the closed-form pole part provably is none.

**(L) The enumeration.** The 8 out-of-band seams up to  $n = 300\,000$  (Mersenne/Fermat pairs and one twin), a finite list; where a single step fails, a three-stage ladder composes — the fallback is part of the statement.

The craft judgement, labelled as one: the three remaining inequalities are of the *same craft* as the proven links of the chain — finite-dimensional inequalities between constructed matrices — and the residual risk is concentrated in the word *uniform*, which no amount of this probe's arithmetic can discharge. T128 (`teml_probe.py`, contract `TEML`, 27/27, verdict `THREE-OF-FOUR`) then worked that four-point list in order of cost, cheapest first. Three of the four points stand at their preregistered bars; the fourth — the protrusion concentration inside (T) — missed its own bar by 3.6%, systematically in the re-gridding ratio, which is exactly the kind of near-miss that distinguishes a bound from a fit. The remainder of this section is the four points in order.

## 10.9 (L) is derived and closed: the out-of-band table (T128)

The out-of-band class is no longer quoted but *derived*: applying the U5-band threshold  $\rho > 1.49531$  (T127's number, quoted and never re-bisected) to the free-resolution record schedule yields *exactly eight* seams up to  $n = 300\,000$ , all at next-atom distance 1 and all with identical border geometry ( $gc_c = 2$  coarse border cells,  $gc_f = 4$  fine) — each is a pair  $(2^k - 1, 2^k)$  or  $(2^k, 2^k + 1)$  whose log-gap halves the previous record. Every seam carries a status and no residue:

| seam $n$ | ratio $\rho$ | class       | status                     | retention / required $h_f$ |
|----------|--------------|-------------|----------------------------|----------------------------|
| 7        | 1.671        | Mersenne    | CERTIFIED-1STEP            | 0.179                      |
| 16       | 1.943        | Fermat-like | CERTIFIED-1STEP            | 0.156                      |
| 31       | 1.910        | Mersenne    | CERTIFIED-LADDER (2 rungs) | 0.091                      |
| 127      | 2.024        | Mersenne    | BUDGET-OPEN                | $h_f = 2\,476$             |
| 256      | 2.012        | Fermat-like | BUDGET-OPEN                | $h_f = 5\,694$             |
| 8191     | 2.026        | Mersenne    | BUDGET-OPEN                | $h_f = 295\,253$           |
| 65536    | 2.001        | Fermat-like | BUDGET-OPEN                | $h_f = 2\,907\,297$        |
| 131071   | 2.000        | Mersenne    | BUDGET-OPEN                | $h_f = 6\,177\,926$        |

The five open seams exceed the hard factorisation cap  $h \leq 1500$  by  $1.7\times$  to  $4119\times$ , and each is open for *one* reason only: a single Cholesky of size  $h_f - gc_f$  is missing — nothing conceptual. Each is characterised arithmetically and geometrically without any factorisation (ratio, both border cell counts, both window roundings, the flat transport factor), so the characterisation is available at any depth. “Closed list” means every item is enumerated and labelled, not that every item carries a Cholesky — that distinction is the point of the block. Mihăilescu 2004 (Catalan) remains the cited, unused address of the class; over a finite range a list is enumerated, not conjectured.

### 10.10 (T) becomes a bound, and its bar is missed honestly

The step from predictor to bound is an exact identity plus two explicit majorants. With  $y$  the transported  $S_f$  eigenvector,  $m_{\text{prot}}$  its mass on the protrusion,  $m_{\text{fill}} \geq 0$  the interior mass on the sliver the coarse block adds, and  $V_{\text{bord}} \geq 0$  the intra-coarse-cell variance of the transported function,

$$\tau_{\text{dn}} = \|y\|^2 - m_{\text{prot}} + m_{\text{fill}} - V_{\text{bord}}, \quad \text{hence} \quad \tau_{\text{dn}} \geq 1 - m_{\text{prot}} - V_{\text{bord}},$$

verified to  $9.0 \times 10^{-13}$  against an independent full projection on all 126 transports. The bound is valid on 126/126 (slack =  $m_{\text{fill}}$ , median 0.000); with the two explicit majorants — the outer-cell majorant for  $m_{\text{prot}}$  and the discrete oscillation majorant  $\rho(\lceil \rho \rceil + 1)/2 \cdot \sum (\Delta y)^2$  for  $V_{\text{bord}}$  — it stays valid on 126/126 and positive on 125/126 (the exact form is positive on 126/126); total slack median 0.275, and the loose part is the oscillation majorant (median  $7\times$  overshoot), so what (T) still needs is a Bernstein/Markov-type statement about the discrete oscillation of one  $\text{gc}_f \times \text{gc}_f$  eigenvector. **The honest miss:** the protrusion concentration  $\kappa = m_{\text{prot}} \text{gc}_f / (t_l + t_r)$  — eigenvector mass on the protrusion against what a flat border vector would put there — measures 0.086–2.071 over the 120 transports with a nonzero protrusion, and *misses the preregistered bar*  $\kappa \leq 2$  on 1/120 by 3.6%. The excess is systematic in the ratio, not accidental: the group maxima rise monotonically, 1.889/1.965/2.071 at  $\rho = 1.001/1.250/1.495$ , with the argmax at the band edge. The bar is *not* moved — a bar moved after the numbers would make everything else here worthless; the shape the next preregistration should declare is  $\kappa \leq 2 + C(\rho - 1)$ , the shape the U5 band already has. A side finding with a side:  $\kappa$  is 1.61–2.07 when the protrusion sits at the *inner* edge of the border block and 0.09–0.25 at the window edge — (T) is an *eigenvector* statement, not a rounding statement, contrary to what the +0.997 correlation invites one to think.

### 10.11 (E) and (T) jointly: the $\ell_2$ route dies, the energy route carries

The  $\ell_2$  route to  $\eta_{\text{proj}}$  is structurally lossy where it matters: the norm bound overshoots the true defect by a median factor of 55, and since the available margin relative to  $\lambda_c$  is smaller still, it certifies 0 of 12 real seams. What carries is the *energy route*, an exact Cauchy–Schwarz chain in the  $Q$ -inner product: with  $Q_f$  certified PSD and  $d = w_f - Pv$  the projection residual,

$$\sqrt{\lambda_f} \geq \sqrt{\lambda_c \tau^* - |\eta_{\text{cons}}|} - \sqrt{e_d}, \quad e_d = d^T Q_f d,$$

with  $e_d/\lambda_c = 0.041$ –0.154 on the real seams. The joint chain certifies a positive floor on 9/9 **in-band seams** (floor  $\geq 0.0850$ ) and 11/12 overall. The one tear —  $n = 31$ , already on the (L) list — is  $\eta$ -limited, not geometric:  $|\eta_{\text{cons}}|/\lambda_c = 0.392$  exceeds the whole geometric floor  $\tau^* = 0.363$ , so no sharper (T) would have helped; and shrinking the ratio is no alternative either, since the consistency defect does not vanish as  $\rho \rightarrow 1$  (median 0.089 at  $\rho = 1.0014$  over 43 zones). The tear is repaired by the declared two-rung ladder in which *every rung closes the chain by itself* — the fallback is part of the statement, not a patch.

### 10.12 (M) is earned: the exclusion is a proof, the floor is constructed

Three parts. (a) The boundary-layer exclusion for the pole part of  $\hat{s}$  is now a *proof* from the closed form  $(Z^T b)_j = -\frac{16}{\sqrt{2D}} \sinh^2(D/4) \cosh(\bar{x}_j^c/2)$ : the adjacent-component ratio lies in  $[e^{-D/2}, 1]$  and the outer- $k$ -cell share is bounded by  $kD(1 + \cosh \alpha)/\sinh \alpha$  — pole edge mass  $\Theta(D)$  against a layer's  $\Theta(1)$ , separation at least  $84\times$  on all 32  $(\alpha, D)$  pairs, by elementary inequalities rather than a finite sample. (b) The crude floor is *earned*:  $\mu_{\min}(S, A_z) = 1 - \gamma^2$  exactly (to  $5.9 \times 10^{-15}$ ; the strengthened Cauchy–Schwarz constant of the two-grid splitting, Eijkhout–Vassilevski 1991), and  $U \preceq (1 + \varepsilon_{\text{env}})A_z$  with  $\varepsilon_{\text{env}} \leq 0.038$  Cholesky-certified, so the constructed floor is a true lower bound on all 31 rungs: **0.188** against the  $1.62 \times 10^{-2}$  the harmonic-mean argument demands — an  $11.6\times$  margin, where T127 could only measure. (c) The comb residual stays measured, but it is measurably *no* layer: half of its mass needs 12–39% of the half-window against one cell for a

genuine boundary layer, and it vanishes with the cell width as  $D^{+1.54 \pm 0.09}$  (a fit with a jackknife band).

### 10.13 The map V3: twenty-one points, five genuinely open

#### The map V3 (T128): what is left, by kind

21 points: **10 done** (T127's map still had the (M) floor, the (T) identity and the (L) list among its open items; all three are now closed), **3 certified on real objects** — valid on every seam, transport and rung the hard cap admits, wanting only the word “uniformly” — and **3 conditional-classical** (Bernstein/Markov for the border oscillation, Bramble–Hilbert for the consistency defect and the  $Q$ -invisible residual, and the continuum leg). **5 genuinely open**, and no two of the same kind:

**M9** the  $\kappa$  bar — missed by 3.6%, systematically in  $\rho$ ; next hypothesis  $\kappa \leq 2 + C(\rho - 1)$ .

**M17** the pencil mass of  $\hat{s}$  on  $\{\mu \geq \frac{1}{2}\}$  — one measured scalar per rung, no derivation.

**M18** the comb residual is not a cell-scale layer — measured, not yet bounded.

**M19** uniformity over *all* zones — the word “for all”, which a finite probe structurally cannot supply.

**M21** the RH address itself — Weil positivity as the address of the limit statement, cited and never used.

The honest distance, in the probe's own three sentences, quoted:

*(1) After this block the finite-window chain has no step left that is merely plausible: every inequality it uses is either a completed certificate, an identity verified to machine precision, or an explicit inequality that has been evaluated — and found valid — on every seam, transport and rung the hard cap  $h \leq 1500$  admits, so the remaining mathematical content is three named classical estimates (Bernstein/Markov on one small Schur complement, Bramble–Hilbert twice) plus two measured scalars (the pencil mass, the comb's delocalisation). (2) The distance is nevertheless not small, and it is not small in a way no larger computation can shrink: 5 of the 21 map items are open because a finite list of objects cannot produce the word “for all zones”, and one of them — the protrusion concentration  $\kappa$  — has just been measured outside its own preregistered bar by 3.6 per cent, systematically in the re-gridding ratio rather than accidentally, which is exactly the kind of near-miss that distinguishes a bound from a fit. (3) And even with all four TEMPL points closed, what would stand is a positivity statement about a finite window at the  $\alpha$  actually reached: the continuum leg (M20) and the passage to Weil's criterion (M21) are untouched by anything in this probe, so the honest distance to RH is not “one uniformity lemma” but “one uniformity lemma, three classical estimates, a continuum limit, and a theorem nobody has”.*

T129 (`kappa_deep_seams_probe.py`, contract KAPPA.DEEP.SEAMS) is running: the  $\kappa$  law  $\kappa \leq 2 + C(\rho - 1)$  on a broad front, the two affordable deep seams (127, 256) via the T115 compression, an M18 bound, and a map V4 with a promotion assessment.

## 11 Status and outlook

### 11.1 What is certified, what is measured, what is hypothesis input

#### The three-way split, stated once

**Certified** (exact identities or closed-form bounds, machine-checked): the seven ledger modules of Section 4; and, in the sandbox arc, the bare bound  $b_0$  in closed form, the parity superselection  $JQJ = Q$ ,  $JB = -BR$ , the pole splitting  $ab^T + ba^T = ss^T - tt^T$ , the channel inertia, the avoidance law  $Q_{\text{full}} \geq 0 \Rightarrow \|B^T v_i\|^2 \leq w_i \Lambda_-$ , the growth inequality, the accumulation bounds, the Sherman–Morrison/Woodbury reduction of (R) to  $r \leq 1$ , the three T108 identities ( $\varepsilon$  as a  $Q|_{\text{odd}}$  energy, the rank-one overlap  $hh^T/\varepsilon$ , and  $\varepsilon$  as the square of the last Cholesky pivot), the two T109 certificates (the graded matrix cap  $\omega_{\text{cert}} = 0.2561\text{--}0.7569$ , no hypothesis input, and the residual certificate for the boundary value, sharpness  $1.0000\text{--}2.4937$ , 16/16 in budget), the three T110 pieces of the propagation circle

(the exact splitting of the growth step, embedding error  $\leq 2.6 \times 10^{-14}$ , with the new atom's restriction to the old window the exact zero matrix; the graded Loewner minorant step, valid exactly under the induction hypothesis, holding 15/15 with retention 1.0000; and the base-case Cholesky certificate at  $n = 2$ ,  $\lambda_{\min} \geq 6.93 \times 10^{-2}$  at three resolutions), and their T111 deep-ladder extension (117 handover certificates at retention 1.000000, largest single loss  $5.96 \times 10^{-8}$ ; atom entry the exact zero matrix on 117/117; and the arithmetic ladder-wall statement at frozen cell width, gap  $> 2D$  needed against the twin-prime gap 0.003831 at  $521 \rightarrow 523$ ), the T114 margin-free step certificates (Albert/pseudo-inverse Schur, 27/27 ladder steps to  $n = 1331$  including all seven torn zones, with the exact  $4 \times 4$  Schur complement at 42–67% of the block scale — no margin division anywhere), and the T115 compression certificates (the exact identity  $X_{\text{mixed}} = Q_{\text{old,mixed}}$ ; Albert on 66/66 (zone,  $q$ ) combinations; the one-sided second-order compression error; the deep margin-free step at  $n = 155\,921$  and the ten-step chains — certificates for a coarser Galerkin space, declared as such), and the T116 boundary-recursion identities (the partial-minimisation identity  $t^T T^{-1} t = r^T \Sigma^{-1} r + \sigma$  on six zones; the Haynsworth double-Albert equivalence 6/6 with sharp negative controls; the exact Riccati chaining, bit-exact on real windows; and the exact carriage of the global rank-one pole in the bordered state — the deep march itself certifies the band model's cost geometry, declared as such), and the T117 epsilon-theorem package (the four Galerkin identities — the cell functional of  $2 \sinh(x/2)$ , the exact nesting  $T_c = P^T T_f P$ , the two-level Pythagoras, the dual-norm residual form, all to  $10^{-10}$  or better; the two-level lower-bound chain, every link an identity or a completed Cholesky, positive on all 19 nested pairs; and the exact jump formulas — the Levinson bordering drop  $r_0^2/s_0$  and the rank- $(k_0+1)$  Woodbury corner update, both to  $10^{-10}$ ), and the T118 symbol package (the closed-form two-grid symbols —  $\sigma_S$  the weighted harmonic and  $\sigma_z$  the arithmetic mean of the aliasing pair; the saturation identity  $\varepsilon_c - \varepsilon_f = y^T S y$  to  $4.9 \times 10^{-11}$ ; the CBS identity  $\lambda_{\min}(L_z^{-1} S L_z^{-T}) = 1 - \gamma^2$  to  $2.1 \times 10^{-15}$ ; the isometry ordering onto  $T_h(\sigma_z)$ ;  $\inf \sigma_z > 0$  with explicit second-order margins on 7/14 windows; and the pole-mass identity in the oscillation space,  $\leq 2.1 \times 10^{-5}$  via the DTFT factor  $|\sin(\theta/2)|$ ), and the T119 arithmetic package (the closed-form comb symbol — one cosine per prime-power atom, linearly interpolated over the two nearest lags — with the parity rule for  $\sigma_z$ ; the atom-count bound  $|\sigma_z^{\text{comb}}| \leq \Xi(\alpha)$ , uniform in  $D$ ; the archimedean growth identity  $\inf \sigma_z^{\text{arch}} = \log(1/D) + B$  with universal  $B = -1.0474$ ; the resulting  $D_0$  theorem; the large-sieve section certificate on one window (floor  $+0.396$  where  $\inf \sigma_z = -0.349$ ); the Levinson corner identity and innovation mass budget; the identity  $\kappa_{\text{end}} = 1/(1+R)$  to  $1.1 \times 10^{-16}$ ; and the coarse-graining fence  $\lambda_{\min}(T_h(\sigma * K)) \geq \lambda_{\min}(T_h(\sigma))$  by unitary averaging), and the T120 Harnack package (the exact increment system  $Kv = s - u_0 G \mathbf{1}$  to  $2.4 \times 10^{-12}$ ; the exact symbol identity  $\sigma_c \sigma_z - \tau^2 = f_1 f_2$  to  $3.9 \times 10^{-16}$ ; the pairing certificate (C1) — a proved implication whose single hypothesis is one sign of the corner increments, giving  $|R - 1| \leq 0.04745$  and  $\kappa_{\text{end}} \geq 0.480827$  for  $n_C \geq 8$ ; the frame identity  $\alpha_o = u_k/2 + O(D)$  over all 732 handovers; and the universal- $B$  evaluation of the  $D_0$  criterion on all 1492 zones,  $\alpha^* = 0.693$ ), and the T121 wide-window package (the shifted-Cholesky section floors on the dense frame windows; the Christoffel/Fejér per-dip budget, positive on 8/16 frame windows up to  $\alpha = 3.446$  at sharpness 0.20–0.76; and the exact deficit identity  $y^T S y / \text{SUPPLY} = r_{\text{ray}} r_{\text{cbs}} / (1 - \gamma^2)$  to  $2.8 \times 10^{-14}$ ), and the T122 net-balance package (the closed form of the compressed Hankel term,  $(Z^T H Z)_{jk} = b_{j+k}$ , to  $7.0 \times 10^{-16}$  on 36/36 rows; the two-grid isometry identity  $A_z = W^T \text{Toeplitz}_M(c) W$  with  $W^T W = I$  exact; the certified band bound (5R) and its uniform  $y$ -free variant, non-vacuous on 36/36 rows to  $\alpha = 6.31$ ; the structural Rayleigh bound, positive and sound on 36/36 rows with residual slack 1.00–1.03; and the exact residual decomposition of the fitted balance, closing to  $2.9 \times 10^{-16}$ ), and the T123 structural-CBS package (the two-isometry completion  $V = B P$  with  $V^T V = I$  and  $V^T W = 0$  exact, the three-compression identities to  $3.4 \times 10^{-15}$ ; the closed forms of the coupling — half a first difference of the lag sequence to  $1.8 \times 10^{-15}$ , the aliasing-difference symbol to  $5.7 \times 10^{-16}$ ; the uniform oscillation positivity (5R\*),  $\lambda_{\min}(E_z) = +0.303 \dots + 5.017 > 0$  on every row at slack 1.0007–1.007; the exact dyadic depth-layer split of  $y^T G y$  to  $1.1 \times 10^{-15}$ ; and the two-term residual split of the certified balance, closing to  $2.8 \times 10^{-17}$ ), and the T125 assembly package (the integer window nesting  $M_n(k) = M_o(k+1)$  on all 29 consecutive pairs of the composing run with the incoming atom block bit-exactly zero and the composition residual  $2.9 \times 10^{-14}$ ; the four assembly identities on 88/88 rungs of 52/52 zones; (8R) certified on 88/88 rungs at  $\text{cb}/\delta = 0.9071$ – $0.9971$ ; the  $\varepsilon$  lower bound on 52/52 zones at  $3.2 \times 10^1$ – $8.7 \times 10^{10}$  over the declared floating-point bound; the base-case equivalence sign  $\lambda_{\min}(Q^{(L)}) = \text{sign } \varepsilon$  on 52/52; and the margin-free handover on 62/62 steps with the exact Schur complement  $3.0 \times 10^{10}$ – $4.6 \times 10^{13}$  above its floor while the norm surrogate is negative throughout — 430 completed Cholesky certificates in total), and the T126 seam package (the exact DP decision of the partition question on the real gap table —  $\rho_{\text{crit}} = 2.0258$ , no partition at 1.83, refinement-only infeasible at any ratio  $\leq 64$ ; the Haynsworth reproduction of  $\lambda_{\min}(S)$  to  $2.8 \times 10^{-14}$  with the transport bracket

containing the fine value on 30/30 real pairs; the twelve certified free-resolution seams, eleven in one re-gridding and the one at  $n = 31$  through a three-stage ladder; the end-to-end seam at  $\rho = 2.251$ ; and the exact fractional-Dirichlet identity of the window form).

**Measured** (holds on the resolved windows; no uniform constant):  $\rho_-$ ; the ratio  $r = 0.0053$ – $0.1814$  and  $\hat{r} < 1$  at all 92 resolution points;  $\omega = 0.2366$ – $0.7531$ ; the strict odd-channel margin  $\lambda_{\min}(Q|_{\text{odd}}) = 1.4 \times 10^{-5}$ – $5.1 \times 10^{-2}$ ; the margin map  $m_k \geq \text{need109}$  on 16/16 zones (ratio 2.09–179.6), extended by T111 to 199 zones with the crossing *measured* between  $n = 461$  and  $n = 463$  ( $n^* \approx 462$ , an upper bound; the T110 extrapolation  $n \approx 170$  was a fit artefact); the reserve trend  $f_{\text{crit}} \sim n^{-0.39}$  (fit); the requirement wall  $n = 727$ ; the closure map 44/64; the handover margins; the two race slopes; the T114 ratio closure  $r = 2 \times 10^{-4}$ – $0.103$  on 27/27 steps and the  $D$ -powers of  $\varepsilon$  and  $\kappa$  (fits); the T115 transport split (the bracket's lower end positive on 7/11 regrid pairs, threshold  $\rho^* = 1.83$  against a median ratio 2.001) and the fall laws  $\lambda_{\min}(S) \sim \rho^{-1.69}$ ,  $\varepsilon \sim \rho^{-1.71}$  (fits); the T116 law  $\varepsilon \approx 8.34 D^{1.790} \alpha^{-6.04}$  (jackknife  $\leq 0.041$ ), the factor-120 jumps at prime-power entries, the comb truncation budget growing like  $n^{2.17}$ , and the criticality  $\tilde{t}^T T^{-1} \tilde{t} \rightarrow 0.999975745$  (fits/measurements); the T117 rates  $\theta'$ (sum) = 1.74 against  $\theta = 1.76$  (loss 0.04 powers of  $D$ ) with bound/ $\varepsilon \in [0.111, 0.185]$ , the growth of  $\lambda_{\min}(S)$  under refinement, the entry direction (all 23 prime-power entries *raise*  $\varepsilon$  — correcting the T116 factor-120 reading, a sweep artefact), and the saturation ratio in  $[0.675, 0.744]$ ; the T118 measurements: the corner-exponent drift  $+0.238$  per halving ( $3$ – $7.4\sigma$ ), the concentration growth  $\sim h^{0.99}$ , the CBS floor  $1 - \gamma^2 \geq 0.181$ , the computed saturation band  $\beta \in [0.252, 0.336]$  on 20 triples, the log-versus-power growth decision on the  $128 \times$  FFT lever (fits with jackknife bands), the reduction sharpness 32–76% on 14/14 windows, and the zero crossings of  $\inf \sigma_z$  on 3/5 zones; the T119 measurements: the mass distribution of  $\|y\|^2$  (corner 0.18–0.42, middle 0.19–0.38, inner half 0.36–0.48), the exponents  $\varepsilon \sim D^{1.761}$ ,  $\|y\|^2 \sim D^{1.961}$ ,  $\lambda_{\min}(S)\|y\|^2 \sim D^{1.757}$  (fits), the flat band  $\kappa_{\text{end}} = 0.4906$ – $0.5093$  with drift  $\leq 0.0019$  per level, the ratio  $R = 0.9634$ – $1.0349$ , and the wide-zone crossing depths  $M \sim 10^{13}$ – $10^{187}$  (extrapolations of the certified growth law); the T120 measurements: the corner sign structure (the  $T^{-1}$  corner block positive share 1.0000 against a TP2 minor share of only 0.51–0.88 and a  $K^{-1}$  positive share of 0.44–0.51 — the two refutations), the sign purity 1.0000 on all 3915 rows with  $R \in [0.9098, 1.1011]$ , the per-cell ratio  $|a_j/w_j|$  up to  $4.8 \times 10^3$  (the per-cell inequality is false), the (C1) decay  $n_C^{-0.58 \pm 0.05}$  (fit), the comb-supremum incoherence 0.038–0.603 of  $\Xi$  (a grid supremum), the large-sieve floor usable on 2/6 zones, and the  $\gamma$  drift  $1 - \gamma^2 = 0.079$ – $0.283$  with fitted slope  $-0.297 \pm 0.013$  in  $\alpha$ ; the T121 measurements: section positivity on all 16 frame windows to  $\alpha = 6.277$  (the ten large rows Lanczos measurements, upper bounds by construction), the net balance  $\text{SUPPLY}/\varepsilon_c \sim D^{-0.177 \pm 0.096} \alpha^{-1.565 \pm 0.135}$  with per-resolution refits  $-1.88 \dots -1.62$  (fits), the link-(5) ratio 0.357–0.916 on 42/42 rows (the Hankel refutation), the Rayleigh and CBS slacks 1.29–13.19 and 3.42–16.67, the one-sign run 24/24 at  $n_C = h/16$  against 21/24 at  $h/8$  (run-share slope  $-0.0258 \pm 0.0042$ , a fit), the Fisher–Hartwig drift  $-0.008 \dots + 2.934$  against the allowed  $-0.120$  (inconclusive, declared as the result), the no-cancellation ratio  $4.6 \times 10^{-8}$ – $7.1 \times 10^{-5}$ , and the pairing quotient  $\delta = 0.0126$ – $0.1331$  over 360 rows ( $\kappa_{\text{end}} \geq 0.4688$ ,  $\delta \sim n_C^{+0.29}$ ); the T122 measurements: the Hankel/Toeplitz norm ratio 0.123–0.781 (the Weyl floor positive on only 12/36 rows), the band-bound recovery 0.429–0.995, the high-band mass of  $Wy$  (99.3–100.0% against 0.1–1.7% for  $y$  itself), the actual CBS coupling  $\kappa_y = 0.0067$ – $0.5311$  against  $\gamma^2 = 0.7173$ – $0.9647$ , the certified balance  $S_2/\varepsilon_c \sim D^{-0.003 \pm 0.075} \alpha^{-0.729 \pm 0.090}$  with the measured-CBS balance  $\alpha^{-0.113 \pm 0.063}$  against the ceiling  $\alpha^{-0.116 \pm 0.062}$  (fits), and the band margin  $0.024$ – $0.363 \sim D^{-0.042} \alpha^{-0.525}$  (a fit); the T123 measurements: the band mass of  $Wy$  (99.64–100.00% at  $d = |\theta - \pi| < \pi/2$  against  $8.9 \times 10^{-6}$ – $1.5 \times 10^{-3}$  in the  $\theta = 0$  cell) and of the worst-case  $y^*$  (81.22–100.00% middle-band), the coupling range  $\kappa_y = 0.0067$ – $0.4622$  against  $\gamma^2 = 0.7248$ – $0.9467$  (a factor 2–109), the coarse conditioning  $\text{cond}(A_c) = 2.1 \times 10^4$ – $3.1 \times 10^7$  with  $\lambda_{\min}(A_c) = 1.7 \times 10^{-5}$ – $2.9 \times 10^{-4}$  (envelope-margin ratio  $8.4 \times 10^2$ – $4.0 \times 10^5$ , rising like  $\alpha^{+4.202}$ , a fit), the two-factor overshoot 1.20–3.51 with the certified margin positive on 5/24 rows, the certified balance  $S_6/\varepsilon_c \sim D^{-0.069 \pm 0.099} \alpha^{-0.544 \pm 0.113}$  against the ceiling  $\alpha^{-0.044 \pm 0.076}$  (fits; gap  $\alpha^{+0.500}$ ), and the lid share  $\varepsilon_f/\varepsilon_c = 0.177$ – $0.546 \sim \alpha^{+0.030}$  (a fit); the T125 measurements: the retention against the measured  $\varepsilon_0$  (0.0013–0.9694) and against the  $L$ -level ceiling (0.9140–0.9929), the pairing band  $C_1 = 0.01063$ – $0.20927$  over the ladder (worst case a single shallowest zone;  $C_1 \leq 0.06721$  on the remaining 51),  $R = 0.8243$ – $1.0377$  and  $\kappa_{\text{end}} = 0.4907$ – $0.5481$ , the head-room minimum 31.6 at  $n = 243$ , the rung-level balance  $\text{cb}/\delta \sim D^{-0.010} \alpha^{+0.001}$  and the run-level  $\sum \text{cb}/\sum \delta \sim \alpha^{-0.012}$  (fits), and the (S4) cancellation ratio  $|r_{\text{corr}}| = 0.3888$ – $14.8018$  with the norm-only variant returning exactly 0; the T126 measurements: the transport reach (positive to  $\rho = 2.251$ , first failure at 2.061, an overlap band rather than a clean split), the fixed-ratio zone sweep (12/22 zones at 2.0256), the free-resolution overhead  $D_{\text{free}}/\text{cap} = 0.021$ – $1.000$  (median 0.201) with re-gridding at 764/9 700 boundaries, the seam retention (median 0.578), the pencil alignment 0.866–0.987 (median 0.975), the Galerkin consistency orders  $D^{1.798 \pm 0.057}$  and  $D^{1.802 \pm 0.065}$

with Richardson limit  $2.1256 \times 10^{-4}$  (fits with jackknife bands), and the diverging continuity constants of the refuted naive continuum route; every scaling exponent quoted in this document.

**Hypothesis input / classical** (used, named, not proved here): Kneser, Witt/Eichler, Shimura, Waldspurger/Baruch–Mao, Cohen, Weil–Guinand, Bombieri, Yoshida, Connes–Consani, Douglas, Slepian–Landau–Pollak, Gronwall/Robin (unconditional form only), Grenander–Szegő, Sherman–Morrison/Woodbury, Christoffel–Darboux, Szegő–Levinson, the Weyl bound, Combes–Thomas/Jaffard, Prager–Synge, first-order Kato perturbation, Albert (the pseudo-inverse Schur criterion), Haynsworth (the partial-minimisation form of the Schur complement), C  a/Strang’s first lemma, Bramble–Hilbert, Brandt/Yserentant two-scale spaces, Szeg  –Kolmogorov, Aubin–Nitsche duality, Payne–Weinberger, Levinson/Durbin, Kac–Murdock–Szeg  /Widom/Fisher–Hartwig, Fej  r–Riesz, Hartman/Peller (Hankel smoothing), Gershgorin, Parseval, the local Fourier / two-grid symbol (Brandt, Hackbusch, Fiorentino–Serra), Yserentant’s strengthened Cauchy–Schwarz constant, Axelsson/Axelsson–Vassilevski (the band-split strengthened-CBS mechanism), the Bank–Smith/D  rfler–Nochetto saturation assumption (with Bornemann–Erdmann–Kornhuber as the honest caveat that saturation can fail), the Montgomery–Vaughan/Selberg large sieve (upper direction only), Chebyshev/Mertens (the atom count), Trench/B  ttcher–Silbermann (explicit Toeplitz inverses), Delsarte–Genin (the split Levinson recursion), Wilkinson/Higham backward error.

## 11.2 The asymptotic mountain

Three separate reasons why the cascade of Figure 8 is not a route to a theorem as it stands, and why the document does not present it as one. Since T110 the mountain has a *precise* formulation — the three gaps of Section 7: no reserve, no scalar step law (structurally excluded by the boundary layer), and no uniformity in  $k$ . Since T111 it also has a *measured* one — the three walls: the margin wall at  $n^* \approx 462$  (measured, an upper bound; the earlier  $n \approx 170$  was a fit artefact), the arithmetic ladder wall at the twin-prime pair  $521 \rightarrow 523$ , and the requirement wall at  $n = 727$ . The mechanism itself never failed on the deep ladder — all 117 handovers certify at retention 1.000000, and the reserve even opens with depth ( $f_{\text{crit}} \sim n^{-0.39}$ , fit) — so the actionable content of the measurement is the reinterpretation: the operating variable is the window depth, not the zone index, and a proof along this route must beat the exponent gap  $-0.974 = -0.681 - 0.293$  and couple the depth  $\delta_0$  to the local prime gap. T112 (`adaptive_scaling_probe.py`) then performed that coupling: the ladder wall and the depth dependence of the requirement wall fall *structurally* in the gap-coupled frame, but the margin wall is frame-invariant at exponent  $-0.974$ , with the prime-gap dependence disclosed (Bertrand–Chebyshev suffices upward; Zhang/Maynard force  $\theta$ -uniformity downward). T113 (`limit_operator_probe.py`) then settled what that invariant number *is*: the wall carries the same exponent in every admissible currency — it is substance — but it measures the *discretisation*, not the spectrum: at a fixed window both leading eigenvalues carry the same cell-width power under refinement (the continuum window form has no gap), the positive floor is a cancellation of relative size  $10^{-7}$ – $10^{-4}$ , and the T109 requirement chain divides by an artifact margin. T114 (`margin_free_step_probe.py`) then discharged the leading part of that balance: the margin-free step certificate exists (Albert 1969, the exact Schur complement), runs eleven steps past the old wall to  $n = 1331$ , certifies all seven torn zones, and shows the wall was an  $O(1)$  numerator divided by an artifact floor — chains now stop only at the  $h \leq 1500$  window cost, never at a step, and the (R) demand itself is grid-bound and deleted ([P5]). T115 (`schur_transport_probe.py`) then split the two-object core: the transport of the exact Schur complement across a real regrid certifies only mild refinement ( $\rho \lesssim 1.8$ ) — and for a principled reason, since on nested ladders, where the transport error is exactly zero,  $\lambda_{\min}(S)$  itself falls like  $\rho^{-1.7}$  — while the two-scale compression breaks the compute cap: a certified margin-free step at  $n = 155\,921$  ( $117\times$  deeper), ten-step chains, the stopper always cost and never a step, with the contiguous reach moving from  $n \leq 125$  to  $n \leq 5437$  (the  $n \log n$  barrier divided, not broken). T116 (`boundary_formulation_probe.py`) then ran the boundary reformulation itself: the induction step *is* a boundary process — the global pole rides exactly in a  $12 \times 12 + 12 + 1$  state and one unbroken Riccati march ran 169 236 prepends to 1.35 million cells ( $903\times$  the old cap) at flat cost, declared as a cost-geometry demonstration and not a Weil certificate — but the state cannot be bounded, because the prime comb itself refuses

compression (every incoming cell reflects off every prime power in the window; the truncation budget grows like  $n^{2.17}$ ). T117 (`epsilon_theorem_probe.py`) then made that inequality theorem-shaped: the Galerkin reading is four exact identities, the lower-bound direction is a certified two-level chain that loses no power of  $D$  ( $\theta' = 1.74$  against  $\theta = 1.76$ ), the jumps have exact closed forms — and, correcting T116, the prime-power entries *raise*  $\varepsilon$ , so the factor-120 falls were a sweep artefact and the jump side-condition disappears. Since T117 the mountain is three named analytic lemmas about one symbol, each with a classical address — the Kac–Murdock–Szegő/Widom corner asymptotics of  $T^{-1}$ , a one-level lower bound for  $\lambda_{\min}(S)$  at the Nyquist frequency, and the Bank–Smith saturation constant, two of the three constants rather than rates — plus the comb as the honest support obstruction. T118 (`symbol_lemmas_probe.py`) then measured the three against their addresses, and two of the three stand: saturation is an *identity* here; the Schur floor is rescued by the strengthened Cauchy–Schwarz shift onto the oscillation Gram, certified on 7/14 windows with the remainder shown to be under-resolved rather than obstructed on the 15 680-point FFT lever; and the corner lemma corrects itself — a log singularity is the boundary of all powers, so the chain is repaired at its last step and the rate now lives on the oscillation mass  $\|y\|^2$ . Since T118 the mountain is *one* genuinely new analytic statement — (H2), a  $D^{1.75}$  lower bound for the oscillation mass — plus the standard CBS uniformity (H1), two arithmetic reductions for the comb dips, and the [SUPPORT] obstruction. T119 (`oscillation_mass_probe.py`) then closed the arithmetic half as a theorem —  $\inf \sigma_z > 0$  for all  $D < D_0(\alpha) = \exp(-(\Xi(\alpha) + B))$ , with  $\Xi$  the prime-power atom count and  $B = -1.0474$  universal — proved the energy route to (H2) empty (the hypothesis is genuinely new content), and reduced the oscillation-mass half to one discrete Harnack inequality through the exact identity  $\kappa_{\text{end}} = 1/(1 + R)$ . T120 (`harnack_probe.py`) then *explained* that inequality — the two increment families are one sequence offset by a single fine cell, so given one sign of the corner increments the summed bound  $|R - 1| \leq 0.047$  is a proved two-line implication, while the per-cell version is provably false and the two obvious structural routes to the sign (total positivity, the discrete maximum principle) are refuted by measurement — and at the same time sharpened the wide-window defects: the frame consumes  $\alpha = u_k/2$  by construction, so the unconditional  $D_0$  route covers not one handover of the real ladder ( $\alpha^* = 0.693$ , below the first admissible zone), and  $1 - \gamma^2$  decays with  $\alpha$ . Since T120 the mountain is four named defects — one sign of the corner increments (a corner-localised Fisher–Hartwig estimate), a uniform bound on the pairing ratio (a discrete gradient estimate),  $D_0$  for the whole ladder, and the  $\alpha$ -dependence of  $\gamma$  — plus the [SUPPORT] obstruction; the chain contains no zeta input anywhere. T121 (`wide_window_probe.py`) then ran the  $\sigma_z$  route against the real ladder and restructured two of the four: the route survives as a *section* statement ( $\lambda_{\min}(T_h(\sigma_z)) > 0$  on 16/16 frame windows to  $\alpha = 6.277$ , half of them where the symbol infimum is negative; the certified Christoffel budget reaches  $\alpha = 3.446$ ) and dies as a symbol statement, and the net  $\alpha$ -balance of the whole chain is one poly-log ( $\alpha^{-1.57}$ , uniform in  $D$ ), split exactly into the Rayleigh and CBS steps with nothing in  $\gamma$  — while one link recorded as classical was refuted (the odd sector’s Hankel term). Since T121 the mountain is the same four defects changed in kind — the corner sign (now a head-versus-tail inequality at  $n_C = h/16$ ), the pairing bound, the section lemma with the Hankel term, and one poly-log of  $\alpha$ -slack with a structural address — plus the [SUPPORT] obstruction. T122 (`net_balance_repair_probe.py`) then repaired the two worst-case steps that poly-log named: the Hankel term is the reflection half of an exact isometry ( $A_z = W^T \text{Toeplitz}_M(c)W$ ,  $W^T W = I$ ), so link (5) is retired by an identity — and the per-dip Nikolskii budget with it — while the certified band floor reaches  $\alpha = 6.31$  where the Christoffel budget tore at 3.45; the structural Rayleigh bound is sharp (residual slack 1.00–1.03, flat in  $\alpha$ ), the certified deficit halves to  $\alpha^{-0.73}$  (exactly uniform in  $D$ ), and with the measured CBS coupling the balance closes to within the chain’s own ceiling ( $\alpha^{-0.113}$  against  $\alpha^{-0.116}$ ). Since T122 the mountain is the two unchanged sign/pairing defects plus two scalar estimates of the same band-profile kind — the band margin (F3) and a structural CBS bound (F4) — plus the [SUPPORT] obstruction. T123 (`structural_cbs_probe.py`) then ran the same band machinery at both scalars, and they *split*: the margin closes almost for free (the reduction to the band profile of  $Wy$  costs 0.7%, and the oscillation block gains uniform certified positivity, (5R\*)), but the coupling

resists structurally — every certified route to  $\kappa_y$  needs the coarse form from below, whose near-null direction is smooth and lives on a cancellation no pointwise symbol minorant can see — and the remaining  $\alpha^{+0.500}$  gap to the ceiling is identified with the  $\varepsilon_f$  lid itself, the same object 4–8 recursion levels deep. Since T123 the mountain is the unchanged Harnack pair (F1)/(F2), one band-profile lemma (F3), and coarse-level positivity (F4) — no longer an estimate to sharpen but the recursion base case — plus the [SUPPORT] obstruction. T124 (`multilevel_telescope_probe.py`) then built that recursion, and the telescope carries: the ladder is one window form on nested spaces ( $P^\top A^{(\ell+1)} P = A^{(\ell)}$  exact — the two-level system is literally one rung), the rung contribution is a residual in the inverse norm — a maximum, not a minimum — so it is certified from *above*, where the envelope works ((8R), valid on 400/400 rungs at a drift seven times weaker than T123’s), and (F4) collapses from a quantitative estimate to a sign the coarse-to-fine induction already carries, with base case pure semidefiniteness —  $\lambda_{\min}(A_c)$ ,  $\text{cond}(A_c)$  and  $\gamma^2$  leave the chain entirely. T125 (`grand_assembly_probe.py`, the series finale) then mounted all of it on one ladder and changed the *shape* of the mountain rather than its height: on 52 of 52 zones — 30 of them a literally composing run of consecutive atoms on one common resolution — the load-bearing spine [A][B][C][E] is at CHOLSKY-CERT, so the Harnack pair (F1)/(F2) *leaves the spine* and survives only in the second, independent route (iv). Since T125 the mountain is therefore *one* declared accounting convention ((C-F5)) on the spine, the Harnack pair as an alternative route rather than a bottleneck, and **(F6) uniformity in the zone index** — no uniform-in- $k$  handover Schur complement, no uniform lower bound on  $\text{cb}_\ell$  — plus the [SUPPORT] obstruction. That uniformity gap, not any missing estimate, is the honest distance from the chain to any infinite statement (Section 9), and it is where the series closes. T126 (`uniformity_seams_probe.py`, the first part of the second phase, Section 10) then opened the attack on exactly that gap and finished the seam architecture: the segment-partition question dissolves into a monotone free-resolution construction (re-gridding only at record-small gaps, all twelve affordable real seams certified, a full seam end to end at the ratio the arithmetic demands), the continuum step closes through an exact fractional-Dirichlet identity onto Lipschitz test functions in the reached window, and the audit splits (F6) into three named lemmas of which exactly *two* are genuinely new — the U3 direction lemma and the U5 zone-uniform seam floor. T127 (`two_inequalities_probe.py`) then dissected both, and both changed: the U5 separating variable is pure interval geometry (the border protrusion  $s_{l,r}$ , not the atom and not the gap), U5 as posed is *refuted* and replaced by a band covering 99.26% of the real seams with the 8 exceptions enumerated (Mersenne/Fermat pairs and one twin), and U3 collapses to a coarse floor through the harmonic-mean identity, with exactly one direction type to exclude. T128 (`teml_probe.py`) then worked the four-point list cheapest first, and three of the four stand at their preregistered bars: (L) is derived and closed (the three affordable seams certified, the five open ones each lacking exactly one Cholesky beyond the cap), (T) becomes an exact identity plus majorants valid on every transport, (E) certifies through the  $Q$ -inner-product chain on all nine in-band seams, and (M)’s boundary-layer exclusion is now a proof with an  $11.6\times$  floor margin — while the fourth point, the protrusion concentration  $\kappa$ , missed its own bar by 3.6%, systematically in the ratio. Since T128 the mountain is *five* named open points — the  $\kappa$  bar (M9), two measured scalars (M17 the pencil mass, M18 the comb), the word “for all” (M19), and the RH address itself (M21) — plus the (C-F5) convention, the all- $\alpha$  step that no computation can close, and the [SUPPORT] obstruction; T129 (`kappa_deep_seams_probe.py`) is running at the  $\kappa$  law.

1. **Sixteen zones are sixteen zones.** Everything below the top line of the cascade is established on the resolved zones  $n = 2 \dots 29$  (T111’s deep ladder extends the measurements to  $n = 521$ , T114’s margin-free ladder to  $n = 1331$ , T115’s compressed run to  $n = 155\,921$ , T116’s boundary march to  $1.35 \times 10^6$  cells with atom reach  $n = 173$  — finitely many zones either way). Not one constant in the chain is certified uniformly in  $k$ . The T101 requirement list (A)–(D) is unchanged: the hardness would sit in (A), a sharp, non-perturbative arithmetic lower bound at a Weil-positivity edge.
2. **Finite resolution is a wall, and a measured one.** Numerics as a witness is exhausted

beyond  $\alpha \approx 0.55$  ( $O(M^{-1.8})$ , T96); the finite-block route was judged the structurally wrong instrument, not an under-resourced one (T92);  $\lambda_{\min}$  collapses geometrically by a factor 11.4 per mode, so double precision ends at  $N \approx 10$  while  $N = 24$  would need  $10^{-27}$  entries.

3. **The induction base is not free either.** The recursion terminates arithmetically (240/240 in  $\leq 4$  steps to the classical zone, T99; 960/960 with longest chain 13 steps, T101), but the zone peaks need percent-level bounds, not 0.1% (T99: extrapolated peak margins  $+2.7 \times 10^{-2}$ ,  $+1.4 \times 10^{-2}$ ,  $+1.7 \times 10^{-2}$ , i.e. 2.2–5.5% of  $\mu_k/2$ , with zone 5 saturating exactly).

### 11.3 Extraction candidates: the community-checkable path

The useful question at this point is not “how close is the program to RH” — it is not close, and no such claim is made — but “which pieces of it are unconditional theorems that someone else can check”. Three candidates are named in the diary and are repeated here as the outlook; the third one is the finale’s own recommendation and now leads:

- **The assembled chain, as a conditional lemma paper in numerical linear algebra** (the T125 recommendation). Stated without one word of this project’s vocabulary: *a certified multilevel lower bound for the Galerkin prediction error of a log-singular Toeplitz form with a prime-power comb*, with the (8R) direction reversal as the theorem and the margin-free bordering induction as its corollary — the comb is then merely the example that makes the symbol log-singular, and no arithmetic conjecture is mentioned anywhere. The reassessment scored it against the matching-lemma package on the five criteria of the July extraction analysis and found it *stronger* on three (prior-art independence — the ingredients are classical but the assembly is a corollary of none of them; self-containedness — decisive, since the statement needs no project vocabulary at all; referee surface — three Schur-complement facts, Parseval, Galerkin orthogonality and one floating-point convention) and *comparable* on the other two (machine-checkability: 430 Cholesky certificates and 440 identities against v541’s 33/33 with a  $10^6$  enumeration; honest boundary). Two conditions come with the recommendation: carry (C-F5) in the *abstract* rather than in a remark, together with the measured failure of the solution-free rung as a separate proposition — that negative result is publishable content and it is what keeps the gap from looking overlooked — and do *not* extract (F1)/(F2) with it, since they are window certificates outside the spine and their classical address belongs in a separate note whose whole content is a discrete Harnack inequality for the increments of the inverse.
- **Zone extension beyond  $\log 2$ .** Unconditional Weil positivity is classically known exactly up to support width  $\log 2$  (Yoshida 1992; Bombieri 2000). T91’s target inequality — for  $a \in (\log 2, a^*]$  and  $\|f\| = 1$ ,  $(P_{\text{pole}} + A_{\text{arch}})(f) \geq \sqrt{2} \log 2 \cdot h_f(\log 2)$  — is typed as *provable-shaped* and would, if proved, be a standalone classical result: a zone extension past Bombieri’s boundary. Its typing is explicit and unflattering:  $\text{RH} \Rightarrow (\text{T})$ , but  $(\text{T}) \not\Rightarrow \text{RH}$ . T92 showed the finite-block route to it is the wrong instrument; that narrows, rather than closes, the target.
- **The matching lemma.** Already proved exact-integer on  $[4, 10^6]$  with an explicit margin inside v541, and reduced by T77 to a restricted  $\sigma_{-1}$  divisor bound in the Gronwall regime — with the note that Robin’s *unconditional* 1983 bound suffices at the measured headroom, so the RH-equivalent sharpening is not needed. Closing the tail (the correlated-cancellation lemma on credit-rich non-coherent atoms) would be a self-contained classical statement.

#### Final fence

The prime line of TFPT has produced seven machine-checked modules about lattice-native Hecke structure, a finite relative-trace identity, and an exactly closing transport ledger. It has *not* produced evidence for the Riemann Hypothesis, and the one place where RH appears — I5 — appears as an equivalence typing that names the residual difficulty rather than reducing it. The corresponding open gates in the ledger (ZETA.CENSUS.TO.GL1, ZETA.HP.CARRIER) are untouched by everything in Sections 5–7, and no marker moves anywhere on account of this document.

## A Probe index, diary parts T11–T128

Every row is one exploratory probe in `experiments/tfpt-discovery/`; the file column omits the common suffix `_probe.py`. “Checks” is the probe’s own assertion count, all green. Verdict strings are the diary’s preregistered verdict enums. Rows marked “ $\rightarrow vN$ ” were later consolidated into the corresponding load-bearing module; the promotion, not the probe, is the load-bearing artefact. The index covers the 118 parts T11–T128 that the current diary file records (119 runs — T104 was a double run with two independent arms — and 3033 green checks); the ten earlier probes T1–T10 predate that file and are not indexed here. Series totals as recorded by the diary at T125, the final part of the series: 125 **probes**, 3139/3139 **sandbox checks**, plus **v535–v541** promoted; with T126–T128, the first three parts of the second phase, the running totals are 128 **probes and** 3225/3225 **checks**. T129 (`kappa_deep_seams_probe.py`, contract `KAPPA.DEEP.SEAMS`) is running and not yet indexed.

| Part | Probe                                           | Contract / finding                                                                          | Verdict                                   | Checks |
|------|-------------------------------------------------|---------------------------------------------------------------------------------------------|-------------------------------------------|--------|
| T11  | <code>e8_glue_chi4_signed_theta</code>          | the $\chi_4$ -signed glue theta welds parts 6 and 7                                         | NO PROMOTION                              | 19/19  |
| T12  | <code>e8_glue_lseries_selfdual</code>           | L-series of the signed glue count; Apéry/ $\zeta(3)$ bridge; Poisson self-dual width $2\pi$ | NO PROMOTION                              | 16/16  |
| T13  | <code>e8_glue_character_atlas</code>            | glue-character atlas of $E_8$                                                               | NO PROMOTION                              | 11/11  |
| T14  | <code>seam_selfdual_width</code>                | SEAM.SELFDUAL.WIDTH; unimodularity gate                                                     | DEFLATED                                  | 13/13  |
| T15  | <code>zeta_arith_completion</code>              | ZETA.ARITH.COMPLETION; shell Hamiltonian                                                    | PARTIAL                                   | 12/12  |
| T16  | <code>zeta_local_euler</code>                   | ZETA.LOCAL.EULER; one weight-4 Hecke polynomial                                             | PARTIAL                                   | 16/16  |
| T17  | <code>zeta_funceq_duality</code>                | ZETA.FUNCEQ.DUALITY; Poisson involution                                                     | PARTIAL                                   | 9/9    |
| T18  | <code>zeta_weil_recovery_trace</code>           | ZETA.TRACE.COMPLETENESS; Guinand–Weil balance                                               | BENCHMARK-STANDING                        | 17/17  |
| T19  | <code>prime_thinning_seam_truncation</code>     | PRIMON.SEAM.TRUNCATION; thinning vs. the finite tower                                       | MIXED                                     | 19/19  |
| T20  | <code>seam_archimedean_kernel</code>            | archimedean kernel from the seam density of states                                          | KILLED-AS-NAIVE                           | 25/25  |
| T21  | <code>prime_prediction_mechanics</code>         | four prediction channels quantified                                                         | MAPPED                                    | 19/19  |
| T22  | <code>seam_halfcut_dilation_arch</code>         | non-DOS successor of the T20 kill                                                           | PARTIAL (K1, K2 pass; K3 fails)           | 25/25  |
| T23  | <code>seam_boost_arch_dictionary</code>         | arch dictionary via the seam boost                                                          | DEAD                                      | 23/23  |
| T24  | <code>seam_truncated_trace_rvm</code>           | truncated trace / Riemann–von Mangoldt route                                                | DEAD                                      | 25/25  |
| T25  | <code>seam_tate_phase</code>                    | scattering phase as the last observable                                                     | DEAD (arch route closed)                  | 29/29  |
| T26  | <code>shell_multiplicity_one_quotient</code>    | shell quotient towards multiplicity one                                                     | PARTIAL                                   | 18/18  |
| T27  | <code>e8_pneighbor_hecke</code>                 | Kneser $p$ -neighbours carry the Eisenstein half                                            | PARTIAL $\rightarrow v535$                | 27/27  |
| T28  | <code>e8_kneser_transport_holonomy</code>       | cusp coefficient from mod-3 lattice geometry                                                | TRANSPORT-VISIBLE                         | 22/22  |
| T29  | <code>primon_transport_category</code>          | primon monoid relations                                                                     | RELATIONS-VERIFIED                        | 16/16  |
| T30  | <code>e8_transport_cusp_uniformity</code>       | uniformity of the T28 reading at $p = 5, 7$                                                 | BROKEN (reading); transport holds         | 32/32  |
| T31  | <code>e8_neighbor_operator_decomposition</code> | cusp extraction is a rule $\nu_p = a \text{Id} + b T_p$                                     | REPAIRED-BY-DICTIONARY $\rightarrow v535$ | 49/49  |
| T32  | <code>census_newform_recovery</code>            | census redundancy is 2-adic oldform structure                                               | REDUNDANCY-IS-OLDFORM $\rightarrow v535$  | 24/24  |
| T33  | <code>e8_nu_rule_many_primes</code>             | hardening of the $\nu$ -rule to $p \leq 100$ ; Eichler finding                              | HARDENED $\rightarrow v536$               | 26/26  |
| T34  | <code>compiler_functor_theta</code>             | project ZETA.COMPILER.FUNCTOR founded                                                       | FOUNDED-WITH-ANCHOR                       | 13/13  |

| Part | Probe                         | Contract / finding                                       | Verdict                               | Checks |
|------|-------------------------------|----------------------------------------------------------|---------------------------------------|--------|
| T35  | rankin_selberg_functor        | Rankin–Selberg translates the abelian shadow             | RS-TRANSLATES-EISENSTEIN              | 13/13  |
| T36  | e8_lambda_charsum_evaluator   | cusp remainder isolated geometrically                    | TWO-SIDED → v536                      | 29/29  |
| T37  | e8_salie_signed_cusp          | signed cusp channel, scalable                            | SCALABLE-AND-SIGNED                   | 37/37  |
| T38  | cuspidal_bridge               | Shimura preimage of $f_8$ inside the compiler monoid     | BRIDGE-FOUND → v537                   | 13/13  |
| T39  | zeta_center_atlas             | the abelian drop                                         | ABELIAN-DROP-CLOSED                   | 13/13  |
| T40  | weil_recovery_compiler_sector | Weil-recovery sector terrain                             | TERRAIN-MAPPED                        | 24/24  |
| T41  | functor_wiring_halfintegral   | the half-integral bridge is Hecke-equivariant            | FUNCTOR-WIRED → v537                  | 13/13  |
| T42  | e8_local_densities            | local densities closed                                   | CLOSED → v536                         | 25/25  |
| T43  | kohnen_plus_waldspurger       | Kohnen plus space                                        | PARTIAL                               | 12/12  |
| T44  | kohnen_operator_projection    | central-value route via Kohnen plus                      | OUTSIDE-KOHNEN-SCOPE → v537           | 17/17  |
| T45  | baruch_mao_metaplectic        | central-value family as $b(d)^2$ ; $R$ constant          | CENTRAL-VALUES-WIRED → v537           | 13/13  |
| T46  | stage4_prolate_prototype      | stage-4 prolate prototype                                | IMPL-OPEN                             | 16/16  |
| T47  | stage4_prolate_uv_dirac       | pointwise UV/Dirac prototype                             | UV-NO-SIGNAL                          | 17/17  |
| T48  | lambda_scale_uniqueness       | does the compiler fix the $\Lambda$ scale?               | SCALE-TORSOR (claim killed)           | 17/17  |
| T49  | rankin_positivity_miniature   | the Deligne/Rankin mechanism runs in-suite               | MINIATURE-RUNS                        | 28/28  |
| T50  | half_integral_selfconvolution | the central-value family aggregates exactly              | FAMILY-AGGREGATED                     | 17/17  |
| T51  | waldspurger_family_kernel     | first non-collapsing infinite carrier candidate          | INFINITE-CARRIER-CANDIDATE            | 22/22  |
| T52  | hecke_orbit_transfer          | path space of Hecke words; fibre determinants            | TRANSFER-REALIZED                     | 16/16  |
| T53  | relative_trace_compiler       | v535/v536/v537 are three projections of one RTF identity | ONE-FORMULA → v538                    | 16/16  |
| T54  | waldspurger_canonical_measure | the “canonical discriminant measure” claim               | MARGINAL-WEIGHT (killed)              | 26/26  |
| T55  | rtf_period_pairing            | canonical distributional limit of the finite RTF         | DISTRIBUTIONAL-RTF-TYPED → v539       | 36/36  |
| T56  | rtf_pairing_sharpening        | all T55 identifications to sub-percent                   | SHARPENED                             | 27/27  |
| T57  | infinite_rtf_packing          | both infinity directions in one double series $Z(s, w)$  | PACKED                                | 24/24  |
| T58  | zeta_rtf_xi_residue           | GL(1) shadows of $Z(s, w)$                               | GL1-SHADOW-CLASSICAL (route A closed) | 24/24  |
| T59  | weil_gns_identification       | residual factor in closed form                           | TRANSFORM-REQUIRED                    | 19/19  |
| T60  | plancherel_descent            | Sato–Tate moment contraction                             | PARTIAL                               | 21/21  |
| T61  | st_isotype_gl1_core           | the trivial Sato–Tate isotype is a GL(1) core            | GL1-CORE-EXACT → v539                 | 24/24  |
| T62  | core_isolation                | $d$ -twist mixing vanishes fibrewise                     | CORE-ISOLATED → v539                  | 24/24  |
| T63  | weil_core_completion          | Plancherel correction irreducibly non-automorphic        | EXTRA-TERMS → v539                    | 33/33  |
| T64  | rtf_stabilization             | $\det_2$ absorbs obstruction 2; the minus stays          | HALF-STABLE                           | 34/34  |
| T65  | pi_digits_prime_distribution  | prime distribution in the digits of $\pi$                | PI-NULL                               | 16/16  |
| T66  | pi_prime_four_level           | the $\pi$ /prime question on four levels                 | FOUR-LEVEL-NULL                       | 23/23  |
| T67  | amplitude_dirac_sqrt          | Dirac square root of the amplitude plane                 | DIRAC-SQRT-EXACT → v540               | 27/27  |
| T68  | geometric_polarization        | geometric polarisation $b = N_+ - N_-$                   | RESHUFFLE-ONLY → v540                 | 34/34  |

| Part | Probe                        | Contract / finding                                                       | Verdict                                    | Checks  |
|------|------------------------------|--------------------------------------------------------------------------|--------------------------------------------|---------|
| T69  | mixed_channel_full_weight    | even- $k$ deletion is a determinant invariant                            | DELETION-UNIVERSAL $\rightarrow$ v540      | 33/33   |
| T70  | linear_measure_weil          | positive linear $\Theta$ carrier; pure plus combination                  | LINEAR-PLUS-COMBINATION $\rightarrow$ v540 | 29/29   |
| T71  | transport_terrain            | FE of the linear family exactly anchored                                 | TERRAIN-MAPPED $\rightarrow$ v540          | 40/40   |
| T72  | cone_enlargement             | the cone library saturates at 5/24                                       | COMPLEMENTARY-PAIR $\rightarrow$ v540      | 34/34   |
| T73  | continuous_cone              | per-direction hybrid cones in the continuum                              | HYBRID-GAINS                               | 28/28   |
| T74  | door_a_spectral_polarization | the spectral side is provably sign-blind                                 | VACUUM-STRUCTURE                           | 32/32   |
| T75  | door_b_lambda_transport      | the gap functional $\lambda^*$ has its own structure                     | LAMBDA-STRUCTURED                          | 28/28   |
| T76  | hybrid_universality          | the recipe certifies 91/91 adversarial Weil directions                   | RECIPE-UNIVERSAL-ON-BATTERY                | 22/22   |
| T77  | matching_lemma_structure     | the matching lemma reduces to a restricted divisor bound                 | LEMMA-CLASSICAL-SHAPED                     | 21/21   |
| T78  | lemma_window_proof           | matching lemma exact-integer on $[4, 10^6]$                              | WINDOW-PROVED $\rightarrow$ v541           | 25/25   |
| T79  | transport_ledger             | the transport ledger closes; I5 named and typed                          | LEDGER-CLOSES $\rightarrow$ v541           | 22/22   |
| T80  | tail_correlation_lemma       | character-exact signed envelope; tail to $N^* \approx 10^{23}$           | RESERVE-PARTIAL $\rightarrow$ v541         | 29/29   |
| T81  | recipe_coherent_avoidance    | reachability theorem: coherent targets need coherent $m$                 | AVOIDANCE-FAILS $\rightarrow$ v541         | 31/31   |
| T82  | heat_arch                    | $\Delta_{\text{arch}}$ is the internal $\Gamma$ difference               | ARCH-INTERNAL $\rightarrow$ v541           | 22/22   |
| T83  | fe_symmetrization            | symmetrisation absorbed; the wall is FE-transversal                      | INVARIANT-NULL                             | 27/27   |
| T84  | gaussian_lift                | Größencharacter phases give $L(1, \lambda_k)$ convergence                | LIFT-WORKS-UNANCHORED                      | 29/29   |
| T85  | lambda_equivariant_design    | the $\lambda$ design closes the coherent class                           | LEMMA-CLOSES-LAMBDA $\rightarrow$ v541     | 24/24   |
| T86  | correlated_cancellation      | the $Q$ -pairing involution closes the non-coherent tail                 | LEMMA-FULLY-CLOSED                         | 29/29   |
| T87  | i5_connes_dictionary         | I5 one-family form vs. the semi-local Connes structure                   | DICTIONARY-EXACT-CORE                      | 22/22   |
| T88  | i5_tightness                 | eight near-vertical tight curves; $\Gamma$ -side spacing                 | TIGHT-SET-PARAMETRIZED                     | 27/27   |
| T89  | sonin_crossing               | window compression is Bombieri's classical object                        | CROSSING-MAPPED                            | 19/19   |
| T90  | residual_dissection          | the residual vectors are explicit Gaussian modes                         | CORE-DISSECTED                             | 17/17   |
| T91  | thin_band_analytic           | the band is the one-atom zone; atom rescue                               | BAND-PARTIAL (3/4 closed)                  | 19/19   |
| T92  | zone_extension_proof         | the finite-block route is the wrong instrument                           | T-SKELETON                                 | 36/36   |
| T93  | band_reconciliation          | third independent implementation of the T89/T91 band map                 | MIXED (map recalibrated)                   | 41/41   |
| T94  | prime_blind_demo             | 753/753 primes of an unseen window from lattice counting                 | BLIND-100                                  | 16/16   |
| T95  | continuum_extension          | C1 chain proved; band top behaves as a positivity edge                   | T-CONTINUUM-NUMERIC                        | 28/28   |
| T96  | relay_test                   | the T95 edge was a map artefact; the relay is real                       | EDGE-ARTIFACT + RELAY-CONFIRMED            | 21/21   |
| T97  | relay_induction              | the induction step has proof shape; one cross-block lemma                | ALIGNMENT-ONLY                             | 105/105 |
| T98  | cross_lemma                  | the $\sqrt{\lambda_0}$ law is Douglas range inclusion (circular)         | LAW-CONFIRMED-MECHANISM-OPEN               | 44/44   |
| T99  | dk_recursion                 | exact mirror-parity selection rule; recursive inequality                 | DECAY-LAW-FOUND                            | 23/23   |
| T100 | bulk_remainder               | closure 11/24 $\rightarrow$ 24/24 samples; the drift was a grid artefact | REMAINDER-CLOSES-ZONES                     | 27/27   |
| T101 | k_asymptotics                | 16 zones resolved; primitives flat, instrument loses the race            | CROWDING-TRENDS                            | 31/31   |
| T102 | arithmetic_bound             | the onset is manufactured by the Schur dressing; $C/g$ was a proxy       | MECHANISM-IDENTIFIED                       | 42/42   |

| Part | Probe                | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                         | Verdict             | Checks |
|------|----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|--------|
| T103 | instrument           | slope halves; closure map 7/64 $\rightarrow$ 44/64                                                                                                                                                                                                                                                                                                                                                                                                         | INSTRUMENT-IMPROVED | 29/29  |
| T104 | schur_profile_bound  | arm B: naive margin route dead, threshold split closes 16/16                                                                                                                                                                                                                                                                                                                                                                                               | CHAIN-PARTIAL       | 21/21  |
| T104 | schur_profile_chain  | arm A: independent twin of the same contract                                                                                                                                                                                                                                                                                                                                                                                                               | CHAIN-PARTIAL       | 47/47  |
| T105 | bare_wing_bound      | bare bound certified in closed form; avoidance is a theorem                                                                                                                                                                                                                                                                                                                                                                                                | ONE-OF-TWO          | 28/28  |
| T106 | friedrichs_angle     | parity split of the Weil pole; hardness in the odd channel                                                                                                                                                                                                                                                                                                                                                                                                 | DENSITY-MAPPED      | 32/32  |
| T107 | odd_channel_closure  | Sherman–Morrison/Woodbury reduction to $r = \kappa/\varepsilon$                                                                                                                                                                                                                                                                                                                                                                                            | SCALAR-TRACTABLE    | 30/30  |
| T108 | ratio_certificate    | $\varepsilon$ is an exact $Q _{\text{odd}}$ energy; (R) on two scalars; first non-vacuous margin chain                                                                                                                                                                                                                                                                                                                                                     | EPSILON-IDENTITY    | 44/44  |
| T109 | boundary_decay       | cancellation, not decay; both scalars certified (graded cap + residual certificate); one strict-margin input left                                                                                                                                                                                                                                                                                                                                          | BOUNDARY-CERTIFIED  | 29/29  |
| T110 | margin_propagation   | the induction circle closes on the measured zones (certified base case + graded Loewner minorant; atom entry free); three gaps named                                                                                                                                                                                                                                                                                                                       | MARGIN-PROPAGATES   | 28/28  |
| T111 | deep_zone_stress     | the crossing measured, not extrapolated ( $n^* \approx 462$ , upper bound); three walls (margin / twin-prime ladder / requirement); mechanism never fails, 117/117                                                                                                                                                                                                                                                                                         | CROSSING-CONFIRMED  | 23/23  |
| T112 | adaptive_scaling     | gap-coupled frame: ladder wall and $\omega$ depth-dependence fall structurally; margin wall frame-invariant ( $-0.974$ ); trade to [P1]/[P2], prime-gap inputs disclosed                                                                                                                                                                                                                                                                                   | SCALING-PARTIAL     | 20/20  |
| T113 | limit_operator       | the wall is currency-invariant ( $-1.168$ in all five currencies) but measures the discretisation: no continuum gap, floor a $10^{-7}$ – $10^{-4}$ cancellation, need109 an artifact margin; new balance [P1]–[P4]                                                                                                                                                                                                                                         | SUBSTANCE-CONFIRMED | 27/27  |
| T114 | margin_free_step     | the margin-free step: Albert/pseudo-inverse Schur certifies 27/27 steps (11 past the old wall, to $n = 1331$ ) and all 7 torn zones; the wall was an $O(1)$ numerator over an artifact floor; (R) grid-bound, deleted ([P5]); new two-object core                                                                                                                                                                                                          | WALL-DISSOLVES      | 22/22  |
| T115 | schur_transport      | transport certifies only mild refinement (split at $\rho^* \approx 1.8$ ; on nested ladders $\lambda_{\min}(S)$ itself falls like $\rho^{-1.7}$ ), but the two-scale compression breaks the cap: margin-free step at $n = 155\,921$ , ten-step chains, stopper always cost; remaining list three points, one inequality (Szegő–Levinson)                                                                                                                   | TRANSPORT-BLOCKED   | 26/26  |
| T116 | boundary_formulation | the induction as a boundary process: the $(\Sigma, r, \sigma)$ state carries the global pole exactly and marches 169 236 prepends to $1.35 \times 10^6$ cells at flat cost (cost-geometry demo, not a Weil certificate); the prime comb refuses compression (budget $\sim n^{2.17}$ ); $\varepsilon \approx 8.34 D^{1.790} \alpha^{-6.04}$ with factor-120 jumps — a Galerkin error under Aubin–Nitsche duality                                            | RICCATI-PARTIAL     | 33/33  |
| T117 | epsilon_theorem      | the one inequality theorem-shaped: four exact Galerkin identities ( $\varepsilon > 0 \iff \tilde{t} \notin T(V_c)$ ); a certified two-level lower bound on 19/19 pairs at $\theta' = 1.74$ against $\theta = 1.76$ (no power of $D$ lost); exact jump formulas (Levinson bordering, rank- $(k_0+1)$ Woodbury) — entries <i>raise</i> $\varepsilon$ , the factor-120 falls were a sweep artefact; three named analytic lemmas remain, two of them constants | THEOREM-SHAPED      | 23/23  |

| Part | Probe            | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Verdict           | Checks |
|------|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|--------|
| T118 | symbol_lemmas    | the three lemmas against their classical addresses: saturation is an <i>identity</i> here ( $\varepsilon_c - \varepsilon_f = y^\top S y$ , $\beta \in [0.252, 0.336]$ ); the $S$ -symbol is a weighted <i>harmonic</i> aliasing mean, vacuous on the comb — rescued by the CBS shift onto the oscillation Gram, whose <i>arithmetic</i> mean suppresses the comb quadratically ( $\inf \sigma_z > 0$ on 7/14 windows; on the 15 680-point FFT lever the floor rises logarithmically and crosses zero on 3/5 zones — under-resolved, not obstructed); the corner exponent drifts (log is the boundary of all powers) and the chain is repaired on $\ y\ ^2$ ; one analytic statement remains — (H2), $\ y\ ^2 \geq c D^{1.75}$                                                                                                                                                                                                                        | TWO-OF-THREE      | 36/36  |
| T119 | oscillation_mass | the arithmetic half closes as a theorem: the comb symbol in closed form (one cosine per prime-power atom; parity rule under $\sigma_z$ ), the atom-count bound $ \sigma_z^{\text{comb}}  \leq \Xi(\alpha)$ uniform in $D$ , the archimedean growth <i>identity</i> $\inf \sigma_z^{\text{arch}} = \log(1/D) + B$ with universal $B = -1.0474$ — hence $\inf \sigma_z > 0$ for all $D < D_0(\alpha) = e^{-(\Xi(\alpha)+B)}$ , deep-verified ( $M$ up to 65 536, matrix-free); the large sieve certifies one window's section <i>without</i> pointwise positivity (+0.396 where $\inf \sigma_z = -0.349$ ) and dies structurally for $\alpha \geq 2.5$ (a moment problem); the energy route to (H2) is proven empty (genuinely new content); and $\kappa_{\text{end}} = 1/(1+R)$ <i>exactly</i> ( $1.1 \times 10^{-16}$ ) — one discrete Harnack inequality remains                                                                                    | ARITHMETIC-DONE   | 27/27  |
| T120 | harnack          | the Harnack core is proof-shaped: $a_j$ and $w_j$ are the odd/even half of <i>one</i> increment sequence, offset by a single fine cell, so one sign of the corner increments gives $ R-1  \leq \sum  v_{2j+1} - v_{2j}  / \sum  v_{2j}  \leq 0.04745$ <i>unconditionally</i> (C1; the per-cell inequality is provably <i>false</i> , $ a_j/w_j $ up to $4.8 \times 10^3$ ); $R \in [0.9098, 1.1011]$ on 3915 rows, $\kappa_{\text{end}} \geq 0.480827$ for $n_C \geq 8$ ; two structural routes refuted ( $T^{-1}$ corner positive but not TP2; the discrete maximum principle false — $u$ oscillates in the bulk); the $\alpha$ -audit makes (D2) total ( $\alpha_o = u_k/2$ by construction, $D_0$ holds on 3/1492 zones, $\alpha^* = 0.693$ below the first admissible handover); $\gamma$ gains an exact symbol identity but $1 - \gamma^2$ falls with $\alpha$ (slope $-0.297 \pm 0.013$ ); theorem V4: 18 links, defect list $3 \rightarrow 4$ | HARNACK-EXPLAINED | 21/21  |

| Part | Probe              | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | Verdict               | Checks |
|------|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------|--------|
| T121 | wide_window        | <p>the <math>\sigma_z</math> route survives the real ladder as a <i>section</i> statement and dies as a symbol statement:<br/> <math>\lambda_{\min}(T_h(\sigma_z)) &gt; 0</math> on 16/16 frame windows<br/> <math>(n = 7 \dots 283\,303, \alpha \text{ up to } 6.277, h \text{ up to } 107\,790;</math><br/> 8 with <math>\inf \sigma_z &lt; 0</math>, the section at<br/> <math>0.64\text{--}3.23 \times  \inf \sigma_z </math> positive), the<br/> Christoffel/Fejér per-dip budget certified on 8/16<br/> up to <math>\alpha = 3.446</math> (each dip 0.156–0.720 of one<br/> Fejér width; the dip count grows like <math>e^{2.41\alpha}</math>, a fit);<br/> the net balance <math>\text{SUPPLY}/\varepsilon_c \sim D^{-0.18} \alpha^{-1.57}</math> —<br/> one poly-log, uniform in <math>D</math>, split exactly<br/> <math>(2.8 \times 10^{-14})</math> into the Rayleigh (<math>\alpha^{-1.02}</math>) and CBS<br/> (<math>\alpha^{-0.51}</math>) steps, nothing in <math>\gamma</math>; link (5) refuted on<br/> 42/42 (ratio 0.357–0.916 — the odd sector’s<br/> Hankel term); one-sign 24/24 at <math>n_C = h/16</math> (only<br/> 21/24 at <math>h/8</math>; certificates restated, free since<br/> <math>\delta \sim n_C^{+0.29}</math>); the FH comparison inconclusive at<br/> reachable <math>h</math> (drift <math>-0.008 \dots + 2.934</math> against an<br/> allowed <math>-0.120</math>); <math>\delta = 0.0126\text{--}0.1331</math> over 360<br/> rows, <math>\kappa_{\text{end}} \geq 0.4688</math>; theorem V5: 19 links, the<br/> defect count stays 4 but every entry changes kind</p>                                                                                                                | WIDE-<br>RESTRUCTURED | 21/21  |
| T122 | net_balance_repair | <p>the Hankel term is the reflection half of an<br/> isometry, not an error term:<br/> <math>A_z = W^T \text{Toeplitz}_M(c) W</math> with <math>W = BZ</math>,<br/> <math>W^T W = I</math> exact (the compressed Hankel is half a<br/> second difference of the lag sequence, closed form<br/> to <math>7.0 \times 10^{-16}</math>; norm ratio 0.123–0.781, <math>O(1)</math> —<br/> <math>Z</math> suppresses both halves equally; the Weyl floor<br/> positive on only 12/36); the certified band bound<br/> <math>\lambda_{\min}(A_z) \geq \text{thr} - \lambda_{\max}(G)</math> non-vacuous on 36/36<br/> rows to <math>\alpha = 6.31</math> (the T121 Christoffel budget<br/> tore at 3.45; recovery 0.429–0.995; caveat: the<br/> optimum sits at <math>\text{supp}(g) \approx 1</math>); the structural<br/> Rayleigh bound is <i>sharp</i> (<math>q_{\text{str}} = 1.00\text{--}1.03</math>, drift<br/> <math>\alpha^{-0.002}</math>), removing the whole <math>r_{\text{ray}} \sim \alpha^{+0.83}</math> slack;<br/> the certified balance halves the deficit,<br/> <math>S_2/\varepsilon_c \sim D^{-0.003} \alpha^{-0.729}</math>, exactly <math>D</math>-uniform, and<br/> with the measured CBS coupling<br/> (<math>\kappa_y = 0.0067\text{--}0.5311</math> against <math>\gamma^2</math>) it reads <math>\alpha^{-0.113}</math><br/> against the chain’s own ceiling <math>\alpha^{-0.116}</math> (residual<br/> decomposition closes to <math>2.9 \times 10^{-16}</math>); theorem V6:<br/> the count stays 4, link (5) retired by an identity,<br/> (F3) the band margin<br/> <math>(0.024\text{--}0.363 \sim D^{-0.042} \alpha^{-0.525})</math> and (F4) a<br/> structural CBS estimate absorb (E3*)/(E4*)</p> | NET-IMPROVED          | 20/20  |

| Part | Probe                | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     | Verdict           | Checks |
|------|----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------|--------|
| T123 | structural_cbs       | <p>the coupling resists structurally: the coarse block is the other half of the T122 construction (<math>V = BP</math> an isometry, <math>V^\top V = I</math>, <math>V^\top W = 0</math> exact — <math>A_c, B_x, A_z</math> three compressions of one window Toeplitz form, identities to <math>3.4 \times 10^{-15}</math>); the coupling is half a <i>first</i> difference of the lag sequence, in symbol form <math>-\frac{i}{2} \sin(\varphi/2)[\sigma(\varphi/2) - \sigma(\varphi/2 + \pi)]</math> — the aliasing difference, vanishing where the coarse space lives, which explains the factor 2–109 between <math>\kappa_y = 0.0067\text{--}0.4622</math> and <math>\gamma^2 = 0.7248\text{--}0.9467</math>; both certified routes to <math>\kappa_y</math> need the coarse form from below (band-split CS is correct 24/24 and non-vacuous 0/24, costing <math>\text{cond}(A_c)</math> up to <math>3.1 \times 10^7</math>; the Schur route yields (5R*) <math>\lambda_{\min}(E_z) = +0.303 \dots + 5.017 &gt; 0</math> on every row at slack 1.0007–1.007 but needs <math>E_c \succ 0</math>, true on 0/24) — the near-null direction of the window form is smooth, and its eigenvalue is a cancellation no pointwise symbol minorant sees (envelope margin <math>8.4 \times 10^2\text{--}4.0 \times 10^5 \times \lambda_{\min}(A_c)</math>, rising <math>\alpha^{+4.202}</math>); the band margin (F3) reduces to the band profile of <math>Wy</math> at a cost of 0.7% (dyadic split exact to <math>1.1 \times 10^{-15}</math>, overshoot 1.20–3.51); the certified balance <math>S_6/\varepsilon_c \sim D^{-0.069} \alpha^{-0.544}</math> against the ceiling <math>\alpha^{-0.044}</math> — gap <math>\alpha^{+0.500}</math>, all of it the coupling (two-term residual split closes to <math>2.8 \times 10^{-17}</math>) — and the gap is identified as the <math>\varepsilon_f</math> lid itself, 4–8 recursion levels deep; theorem V7: the count stays 4, restructured — (F1)/(F2) the Harnack pair, (F3) one band-profile lemma, (F4) coarse positivity = the recursion base case</p>                                                                                                                                                                                                                  | CBS-RESISTS       | 20/20  |
| T124 | multilevel_telescope | <p>the telescope carries: the ladder is one window form on nested spaces (<math>P^\top A^{(\ell+1)} P = A^{(\ell)}</math> exact to <math>2.4 \times 10^{-14}</math> — the T122/T123 two-level system is one rung), every rung is the saturation identity and the Galerkin energy of the level correction (<math>\delta_\ell = \ u_{\ell+1} - Pu_\ell\ _A^2</math> to <math>3.7 \times 10^{-8}</math>; telescope <math>\sum \delta_\ell = \varepsilon_0 - \varepsilon_L</math> to <math>9.2 \times 10^{-10}</math>; top rung 0.27–0.88, geometric decay, <math>L = 5</math> ceiling 0.9911–0.9987); the rung is a <i>maximum</i> (dual residual form <math>\delta_\ell = \hat{s}^\top S^{-1} \hat{s}</math>; <math>r_\ell = W \hat{s}</math> by Galerkin), so it wants <math>S</math> from <i>above</i>: <math>S \preceq A_z</math> needs only <math>A_c \succeq 0</math> (Haynsworth), <math>A_z \preceq U = Z^\top T_M(\text{up})Z</math> is Parseval against the certified majorant — (8R) <math>\delta_\ell \geq \hat{s}^\top U^{-1} \hat{s}</math> valid on 400/400 rungs (80 chains, 20 zones), <math>\text{cb}/\delta = 0.4739\text{--}0.9938</math> at drift <math>\alpha^{-0.080}</math>, 7× weaker than T123's <math>\alpha^{-0.544}</math> (<math>\lambda_{\min}(E_c) &lt; 0</math> on 399/400 is irrelevant — (8R) never forms <math>E_c</math>; the Levinson bordering route falls: refinement is not a bordering, <math>\text{tail}/\delta = 13\text{--}863</math>); the near-null direction stays smooth and self-similar (prolonged Rayleigh quotient exactly <math>\lambda_{\min}(A^{(\ell)})</math>, overlap 0.944–1.000), and (8R) needs only the <i>sign</i> of <math>A_c</math> — the previous rung's conclusion: the non-circular induction runs coarse <math>\rightarrow</math> fine, base case <math>\varepsilon_L \geq 0</math>; balance <math>\sum \text{cb}/\varepsilon_0 = 0.6404\text{--}0.9729</math> against the ceiling, two-level <math>\text{cb}/\varepsilon_c \sim D^{-0.090} \alpha^{-0.100}</math> against T123's <math>\alpha^{-0.544}</math> — <math>+0.444</math> of the <math>\alpha^{+0.500}</math> gap; theorem V8: (F1)/(F2) the Harnack pair, (F3) closed, (F4) <i>collapsed</i> to a sign, (F5) new, a bookkeeping convention — the solution-free rung provably fails</p> | TELESCOPE-CARRIES | 28/28  |

| Part | Probe          | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | Verdict        | Checks |
|------|----------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------|--------|
| T125 | grand_assembly | <p>the series finale — the chain composes: all five stages ([A] base Cholesky, [B] telescope with (8R) rungs, [C] <math>\varepsilon</math> lower bound, [D] <math>\kappa_{\text{end}}</math>/parity, [E] margin-free handover) complete on 52/52 ladder zones plus 10 handover-only zones to <math>n = 1331</math>, and ladder 1 is a <i>literally composing</i> run — 30 consecutive atoms <math>n = 2 \dots 79</math> on ONE common resolution <math>D = 3.472446 \times 10^{-3}</math>, <math>M_n(k) = M_o(k+1)</math> as an integer identity on 29/29 pairs, the incoming atom block bit-exactly zero, so <math>X = Q_{\text{old}}</math> and the next zone's leading block is that <math>X</math> to <math>2.9 \times 10^{-14}</math> (T114's [P2] seam constructively <i>circumvented</i>: 0/84 gap ratios dyadic, so one frame is chosen for the whole run, at the price of its length); the load-bearing finding: the weakest stage of the whole mounting is [D] at WINDOW-CERT, but the SPINE [A][B][C][E] is CHOLESKY-CERT on 52/52 — <i>the Harnack pair (F1)/(F2) no longer carries</i>, it lives only in the independent second route (iv); (8R) on 88/88 rungs at <math>\text{cb}/\delta = 0.9071\text{--}0.9971</math> (primal form above <math>\delta</math> on 88/88 — an upper bound, never a supply), base-case equivalence <math>\text{sign } \lambda_{\min}(Q^{(L)}) = \text{sign } \varepsilon</math> on 52/52 at head-room <math>4.4 \times 10^5\text{--}5.6 \times 10^{10}</math>, <math>\varepsilon</math> lower bound on 52/52 at 0.9140–0.9929 of the <math>L</math>-level ceiling and <math>3.2 \times 10^1\text{--}8.7 \times 10^{10}</math> over the declared fp bound, Albert Schur <math>3.0 \times 10^{10}\text{--}4.6 \times 10^{13}</math> above its floor on 62/62 steps while the norm surrogate is negative throughout; 430 completed Choleskys, 440 identities; theorem V-final printed with line-by-line attribution (25 lines: 10 identity, 9 Cholesky, 3 window — all in (iv) —, 1 hypothesis = (C-F5)) and a five-point “what is not claimed”; series balance: kill list 24 routes over 18 parts (4 of them one cause), cascade 7 stations, certification degree 90.3% chain / <b>96.2%</b> spine identity-or-Cholesky, RH distance in six checkable points; new defect (F6) uniformity; extraction reassessment: stronger on 3/5 criteria than the matching-lemma candidate</p> | ASSEMBLY-GREEN | 34/34  |

| Part | Probe            | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | Verdict         | Checks |
|------|------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|--------|
| T126 | uniformity_seams | <p>the first part of phase 2 (“the full proof”) — the seam architecture is finished: the partition question is decided <i>exactly</i> by DP on the real gap table (9 700 zones, <math>n \leq 10^5</math>) and refuted at T115’s <math>\rho^* = 1.83</math> (dies at <math>n = 127</math>, the twin-after-a-large-gap pattern; true requirement <math>\rho_{\text{crit}} = 2.0258</math>, essentially the single quotient <math>g(125)/g(127)</math>; refinement-only infeasible at any ratio <math>\leq 64</math>; coarsening fails at the destination grid, and ladders make it worse); the decisive new measurement: at fixed ratio 2.0256 the certificate survives on only 12/22 zones — the seam dies at the <i>zone</i>, not the ratio, so no threshold closes the programme; what carries is the <i>free-resolution route</i>: the running minimum <math>D_k = \min(\text{cap}_k, D_{k-1})</math> is monotone by construction, re-gridding at only 764/9 700 boundaries (7.88%, record-small gaps), all 12 affordable real seams certified (11 in one re-gridding, <math>n = 31</math> via a three-stage ladder, median retention 0.578), a full seam end to end at <math>\rho = 2.251 \geq 2.0256</math>; the uniformity audit (35 rungs, 14 zones) gives every spine constant a status — U1 self-certifying, U2 one Bernstein bound, U3 measured-flat but <i>directional</i> (alignment 0.975), U4 a compute wall, U5 <i>falls</i> (the load-bearing gap), U6 resolved; the continuum step closes via the exact fractional-Dirichlet identity (86.1% negative off-diagonals; consistency <math>D^{1.798}/D^{1.802}</math>, Richardson limit <math>2.1256 \times 10^{-4} &gt; 0</math>); the full-proof map has 16 points, 5 done/classical, 9 conditional, exactly 2 genuinely new — the (U3) direction lemma and the (U5) zone-uniform seam floor at ratio <math>\leq 2.026</math></p> | SEAMS-CERTIFIED | 31/31  |

| Part | Probe            | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               | Verdict     | Checks |
|------|------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------|--------|
| T127 | two_inequalities | <p>both remaining inequalities are proof-shaped, and neither is the statement T126 named: the fixed-ratio sweep reproduces (12/22, T126's list — the <i>certificate</i> dies, not the positivity: <math>\lambda_{\min}(S_f) &gt; 0</math> on all 22); of 29 candidate separating variables 5 separate perfectly (22/22), all the same <i>geometric</i> quantity — the border protrusion <math>s_{l,r}</math>, i.e. what the fine border block protrudes past the coarse block's inner edge (corr +0.997 with <math>\tau_{\text{dn}}</math>; every arithmetic candidate fails, best 0.727); U5 as posed <i>refuted</i>, replaced by <b>U5-band</b>: zone-uniform for every <math>\rho \leq 1.49531</math>, at least one zone dies for <math>\rho \geq 1.49766</math> (bisected to <math>2.3 \times 10^{-3}</math>, binding <math>n = 19</math>), covering 99.26% of real seams — the remaining 8 are an enumerated list of next-atom-distance <math>\leq 2</math> pairs (Mersenne 3, 7, 31, 127, 8191; Fermat-like 4, 16, 256, 65536, 131071; twin 101, 103; Mihăilescu 2004 cited, not used); the band buys <math>\tau_{\text{dn}} \approx 1</math>, not small <math>\eta</math> (<math>\eta = \eta_{\text{cons}} + \eta_{\text{proj}}</math> exact; max <math> \eta /\lambda = 1.146</math> at the tightest ratio; floor <math>c \geq 0.0090</math> over 307 in-band transports); the sharpest finding: the near-identity census (43 deep zones, <math>\rho = 1.00140</math>) fails <i>once</i> (<math>n = 479</math>, <math> \eta /\lambda = 1.472</math>) — the band is a <i>proxy</i> for the border geometry, repaired by a three-stage ladder (retention 0.5027), and the risk indicator <math>\tau_{\text{dn}}/\tau_{\text{flat}}</math> (Spearman +0.749, the failure its argmax) sees it before any factorisation; U3 defuses via the harmonic-mean identity <math>\text{cb}/\delta = 1/\sum w_i/\mu_i</math> (<math>\hat{s}</math>-mass on <math>\{\mu \geq \frac{1}{2}\} \geq 0.9665</math> over 42 rungs, so a <i>crude</i> floor <math>\mu_{\min} \geq 1.62 \times 10^{-2}</math> gives <math>\text{cb}/\delta \geq \frac{1}{4}</math>); “<math>\hat{s}</math> avoids the bad direction” refuted — flatness is generic, exactly one direction type fails (a boundary layer at the window edge, concentration <math>3.1\times</math> vs <math>0.03\times</math> for <math>\hat{s}</math>); map V2: 16 points, 10 carried, 2 conditional-classical, the 2 genuinely new reformulated as (T)/(E)/(M) plus the enumeration (L) — three inequalities of the same craft as the proven links, residual risk in the word <i>uniform</i></p> | BOTH-SHAPED | 28/28  |

| Part | Probe | Contract / finding                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | Verdict       | Checks |
|------|-------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------|--------|
| T128 | teml  | <p>the four-point list of map V2, worked cheapest first, and three of the four stand at their preregistered bars: <b>(L)</b> is <i>derived</i> and closed — the out-of-band class <math>\rho &gt; 1.49531</math> has exactly 8 seams up to <math>n = 300\,000</math>, all at next-atom distance 1 with identical border geometry (<math>gc_c = 2</math>, <math>gc_f = 4</math>); <math>n = 7, 16</math> certified in one step (retention 0.179/0.156), <math>n = 31</math> through a two-rung ladder (0.091), and 127/256/8191/65536/131071 BUDGET-OPEN at <math>h_f = 2\,476 \dots 6.2 \times 10^6</math> (<math>1.7 \times -4119 \times</math> over the cap; each lacks exactly <i>one</i> Cholesky, nothing conceptual; Mihăilescu 2004 cited as the address, not used); <b>(T)</b> from predictor to bound — the exact identity <math>\tau_{dn} = \ y\ ^2 - m_{prot} + m_{fill} - V_{bord}</math> (<math>9.0 \times 10^{-13}</math> against an independent full projection, 126 transports), <math>\tau_{dn} \geq 1 - m_{prot} - V_{bord}</math> with explicit majorants valid 126/126 (exact bound positive 126/126, fully explicit 125/126; slack median 0.275, the oscillation majorant the loose part at <math>7 \times</math>) — <i>but</i> the protrusion concentration <math>\kappa = m_{prot} gc_f / (t_l + t_r) = 0.086\text{--}2.071</math> misses the preregistered bar <math>\kappa \leq 2</math> on 1/120 by 3.6%, <i>systematically</i> in <math>\rho</math> (group maxima 1.889/1.965/2.071 at <math>\rho = 1.001/1.250/1.495</math>); the bar is not moved, the next hypothesis is <math>\kappa \leq 2 + C(\rho - 1)</math>; and <math>\kappa</math> splits by side (1.61–2.07 right against 0.09–0.25 left) — <b>(T)</b> is an <i>eigenvector</i> statement; <b>(E)+<b>(T)</b></b> jointly: the <math>\ell_2</math> route certifies 0/12 (median <math>55 \times</math> overshoot), the <i>Q</i>-inner-product chain <math>\sqrt{\lambda_f} \geq \sqrt{\lambda_c \tau^* -  \eta_{cons} } - \sqrt{e_d}</math> (<math>e_d/\lambda_c = 0.041\text{--}0.154</math>) certifies 9/9 in-band seams at floor <math>\geq 0.0850</math>, 11/12 overall; the one tear (<math>n = 31</math>, already on the <b>(L)</b> list) is <math>\eta</math>-limited (<math> \eta_{cons} /\lambda_c = 0.392 &gt; \tau^* = 0.363</math> — no sharper <b>(T)</b> would have helped, nor a smaller ratio), repaired by the two-rung ladder, every rung closing the chain by itself; <b>(M)</b> earned: the boundary-layer exclusion is a <i>proof</i> from the closed form (adjacent-component ratio in <math>[e^{-D/2}, 1]</math>, outer-<math>k</math>-cell share <math>\leq kD(1 + \cosh \alpha) / \sinh \alpha</math>, separation <math>\geq 84 \times</math>), <math>\mu_{min}(S, A_z) = 1 - \gamma^2</math> exact (<math>5.9 \times 10^{-15}</math>), <math>U \preceq (1 + \varepsilon_{env})A_z</math> with <math>\varepsilon_{env} \leq 0.038</math> Cholesky-certified — constructed floor 0.188 against the <math>1.62 \times 10^{-2}</math> demanded, an <math>11.6 \times</math> margin on all 31 rungs; the comb residual stays measured but is measurably no layer (half its mass needs 12–39% of the half-window); map V3: 21 points — 10 done, 3 certified on real objects, 3 conditional-classical, 5 genuinely open (M9 the <math>\kappa</math> bar, M17 the pencil mass, M18 the comb, M19 “for all”, M21 the RH address)</p> | THREE-OF-FOUR | 27/27  |

## B Symbols used in Sections 5–7

All of these are sandbox objects; they are listed so the figures and formulas above can be read without the diary.

| Symbol                                                             | Meaning                                                                                                                                                                                                                                                                                 |
|--------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $Q_{\text{Weil}}, Q_{\text{cert}}, \Delta_2, \Delta_{\text{arch}}$ | the four terms of the transport ledger; $Q_{\text{Weil}} = Q_{\text{cert}} + \Delta_{\text{arch}} + \Delta_2$ closes at $4.4 \times 10^{-16}$ (v541)                                                                                                                                    |
| $A_{\text{fam}}, A_{\text{shift}}$                                 | the two archimedean halves of the heat family; $\Delta_{\text{arch}} = A_{\text{fam}} - A_{\text{shift}}$ (T82)                                                                                                                                                                         |
| I5                                                                 | the one remaining coupled inequality, $Q_{\text{cert}} + \Delta_2 + A_{\text{fam}} - A_{\text{shift}} \geq 0$ (T79)                                                                                                                                                                     |
| $\mu_k, w_k, g_k$                                                  | atom strength, handover window width, atom spacing at zone $k$                                                                                                                                                                                                                          |
| $D_k(\alpha)$                                                      | the Schur-complement defect of the induction step; target $D_k \leq \mu_k/2$ (T98)                                                                                                                                                                                                      |
| $\sigma_k(\delta)$                                                 | the $E_-$ Schur profile of the genuine Weil form just above atom entry (T102)                                                                                                                                                                                                           |
| $E_-, E_0, E_+$                                                    | the three atom eigenspaces; the $k$ -th atom acts as $\text{diag}(-\frac{1}{2}, 0, +\frac{1}{2})$                                                                                                                                                                                       |
| $\text{bare}_k, b_0$                                               | $\lambda_{\min}(Q_{\text{full}} _{E_-})$ and its certified excess over $\mu_k/2$ (T105)                                                                                                                                                                                                 |
| $L, \Lambda_-$                                                     | the soft dressing scalar and its free cap (5.63–7.52, divergent) (T105)                                                                                                                                                                                                                 |
| $J$                                                                | the mirror-parity involution; $JQJ = Q, JB = -BR$ (T105)                                                                                                                                                                                                                                |
| $\rho_{\pm}$                                                       | the Friedrichs-angle cosines of the even / odd channel (T106)                                                                                                                                                                                                                           |
| (R)                                                                | $Q_{\text{full}}(\alpha_k) _{J=-1} \succeq (\mu_k/2)P_- _{J=-1}$ , the remaining Loewner statement (T106)                                                                                                                                                                               |
| $s^*, \tau, \kappa, \varepsilon$                                   | Sherman–Morrison scalar, its Woodbury split $s^* = \tau + \kappa$ , and $\varepsilon = 1 - \tau$ (T107)                                                                                                                                                                                 |
| $r$                                                                | the residual ratio $\kappa/\varepsilon$ ; (R) $\iff r \leq 1$ ; measured 0.0053–0.1814 (T107)                                                                                                                                                                                           |
| $x, h, x_0$                                                        | the pole response $x = T_{\text{odd}}^{-1}\tilde{t}$ , its demand compression $h = V^{\top}x$ , and the single edge entry that carries $\ h\ $ on 8 zones (T108)                                                                                                                        |
| $\omega, \omega_{\text{cert}}$                                     | the wing scalar $(\mu/2)\lambda_{\max}(\Gamma)$ with $\Gamma = V^{\top}T_{\text{odd}}^{-1}V$ , and its graded-matrix-cap certificate (T108/T109)                                                                                                                                        |
| $m_k, \text{need109}$                                              | the strict odd-channel margin $\lambda_{\min}(Q _{\text{odd}})$ at zone $k$ , and the explicit demand $(\mu/2)H_{\text{cert}}^2/((1 - \omega_{\text{cert}})\kappa)$ of the T109 chain; $m_k \geq \text{need109}$ on 16/16 measured zones (T110)                                         |
| $(\Sigma, r, \sigma); \theta, \varphi$                             | the boundary state of the Riccati recursion (partial minimisation of the pole-free form over the interior; carries the pole exactly), and the two exponents of the measured law $\varepsilon \approx 8.34 D^{\theta} \alpha^{\varphi}$ , $\theta = 1.790$ , $\varphi = -6.04$ (T116)    |
| $\delta_{\ell}, \hat{s}_{\ell}, U$                                 | the rung contribution of the level telescope ( $\varepsilon_{\ell} = \varepsilon_{\ell+1} + \delta_{\ell}$ ), its oscillation-space residual data ( $r_{\ell} = W\hat{s}_{\ell}$ ), and the certified Parseval majorant $U = Z^{\top}T_M(\text{up})Z$ of the per-rung bound (8R) (T124) |